University of London Computer Centre

DIMFILM

Preliminary User Guide

John C. Gilbert

Table of Contents

`NOS`

Table of Contents

Table of Contents

University of London Computer Centre

DIMFILM

Preliminary User Guide

John C. Gilbert

Part 1 – INTRODUCTION

Chapter 1 – Preamble

 DIMFILM was first released at ULCC in May 1973 and the original User Guide
published one year later. In its initial form DIMFILM was described as a
library of Descriptive Instructions for MicroFILM and was designed for use
with Fortran on the CDC computers then operated at ULCC. In that format
it enjoyed considerable popularity for ten years. Additionally, it was also
implemented on a number of other computers driving different (non-microfilm)
graphics devices.   The advent of the Cray and Amdahl computers at ULCC required
a major conversion exercise if DIMFILM was to continue. At the same time
it seemed sensible to redesign the internal logic of DIMFILM to enable support
of any graphic device through the provision of appropriate driver software
and to ensure a high degree of machine independence so implementation on
different machine ranges would be considerably simplified and – of overriding
importance – the user interface would be constant. Thus the decision was
taken that the new release of DIMFILM would provide a user interface in accord
with the Fortran 77 standard. The major difference that the user of the
first release of the new DIMFILM will encounter is in the change of many
routine names. The opportunity was taken to not only reduce lengths of names
to six characters but to rationalise certain of the names. Further, to ensure
sensible device independence certain parameter ranges were changed – primarily
to remove limitations imposed by the then operating CalComp 1670 Microfilm
Recorder. These changes have been kept to a minimum and most users will have
to do little more than edit references to DIMFILM library routines. A later
chapter in this guide is devoted to changes of this nature that will be required
during conversion of programs. This chapter will also serve as an addendum
to the original User Guide until a full descriptive document is available.
 The new DIMFILM has been designed for the future incorporation of many
new facilities and enhancements. These will be described in this preliminary
guide as they become available.   DIMFILM development has always benefited
from the suggestions and requests made by users. We look forward to this
continuing and would restate here our commitment to consider all written
comments etc. forwarded to us. At the same time, this release is made on
the same basis as those previously – DIMFILM is a developing system and
as such any corrections/improvements will be incorporated at the earliest
possible time. There can be no guarantee that all routines will function
according to specification and your cooperation in reporting any failings
(however seemingly trivial) will be appreciated. The user interface specification
as detailed in this guide is intended to be definitive and we shall endeavour
to support it. Although DIMFILM is a developing product it is supported
in the same manner as other software at ULCC

Chapter 2 – General Notes

```
  In this User Guide a number of conventions have been adopted either  in
the interest of clarity or to circumvent limitations of the  typeface.  These
should be borne in mind when reading  the text.   In the following, actual
Fortran references are not given (CALL being  omitted) and, unless explicitly
stated otherwise, parameters are  typed implicitly in accordance with standard
Fortran default  (i.e. real, unless symbolic name commences I,J,K,L,M,N when
integer).  Thus, for example SUB(X,Y,I) would be referenced with the Fortran
statement CALL SUB(X,Y,I) with parameters X and Y being real and  I integer.
Unless stated otherwise all references are as subroutines.   There is frequent
reference to ranges of values; in particular, in  reference to acceptable
parameter values.  A standard notation  has been adopted, whereby square brackets
[ and ] denote inclusive  lower and upper limits respectively, while round
parentheses ( and )  denote exclusive lower and upper limits.  Thus, for
example,  (0.0,1.0] indicates the range greater than zero and less than or
equal  to unity.  Occasionally it has been necessary to use the notation
<= for less than or equal and >= for greater than or equal; where  underlining
would result in ambiguity equivalent notations <= and >=  have been used.
```

Chapter 3 – DIMFILM – an overall view

```
  DIMFILM is a wide-ranging and comprehensive library of subprograms  for
the generation of graphical output by computer.  Through its use  of descriptive
names and short argument lists, combined with the  sensible use of defaults,
DIMFILM offers a straightforward method  of plotting suited both to the simplest
application and the  experienced user, for whom many advanced facilities are
provided.   The user is able to define, and plot within, his own coordinate
space, while facility to clip and blank at the boundaries of  specified sub-areas
is possible.  Available are several levels  of error checking, with which
programs may be speedily  debugged.   Elaborate plotting is possible with
minimal experience.  Features  such as alternate line styles (e.g. dashed
or dotted lines),  rotation and translation of axes, specification of colour/intensity
etc. are included.  For textual output, a variety of fonts is  included offering
a range of type styles and complexity.  The user  may readily modify his output
script to contain underlined,  italicised or emphasised passages, while super-
and sub-scripts  can be simply incorporated for the output of complex mathematical
expressions.  A variety of geometric functions is included.   Complementary
to the range of vector plotting facilities,  DIMFILM provides comprehensive
device independent raster  functions.  These may be used in conjunction with
all other  features including the wide range of supported colour models.
In addition to these facilities for straightforward plotting, DIMFILM  is
capable of generating other varied types of graphical output.  Many of these
are primarily aimed toward the scientific plot  suitable for publication,
thus various axis scaling (e.g. linear or  logarithmic) are supported together
with polar plotting.  The user  has full control over frame and axis annotation.
Blanking is  operable for the inclusion of legends or other such captions.
Where discrete data coordinates are given interpolation may optionally  be
performed.  Histograms, month-labelled axes and pie-charts are  among the
other features.  The advanced user is able to uniquely  tailor his output
style by redefinition of the various defaults.  These defaults enable the
novice user to immediately produce  acceptable output.  A number of plots
```

Part 1 – INTRODUCTION

```
may be contained in a single  graph, while the user has full control over
the number and  positioning of graphs in a single "frame".  Graphical, diagrammatic
and textual items may be freely intermixed in any plot.   The production
of contour plots is possible through DIMFILM.   Other facilities are under
review and will be introduced as they  are developed. User suggestions/wishes
are valued, and all such  offerings will be given serious consideration.
This document constitutes a concise guide to using the new  DIMFILM. Descriptions
of routines are not intended to be definitive  but are accurate for basic
access to most features.  Definitive  documentation will be prepared at a
later date.  It will be found  that when new, advanced facilities are available
access by certain basic  routines involves predefinition of a range of sophisticated
options.  However, the actions/functions described here will ensure compatibility
with the previous release of DIMFILM when used in conjunction with the  other
basic routines.
```

Part 2 – GENERAL PLOTTING

Chapter 1 – Basic Plotting

 This section introduces the basic concepts of diagrammatic plotting (or
drawing), although restricting discussion to calligraphic (that is, line
or vector) graphical output. The topic of shaded (that is, area filled or
raster) output is deferred to a following dedicated section. This demarcation
enables the whole subject of raster plotting (at both basic and refined levels)
to be presented in a unified way, while users with interest only in line
drawing may skip that section. However, some aspects of basic plotting <u>will</u>
be of essential relevance to users with a principal interest in shaded images
(in particular, device control, plotting areas and colour models).

1.1 – Plotter Control

1.1.1 – Device initiation

 At the simplest level, DIMFILM may be initialised by directly referencing
one of the device nomination routines. One of these <u>must</u> be the first
DIMFILM reference. Each device has its own characteristics and capabilities.
An Appendix is included to describe each device in detail. This contains
details of specific device parameters the user may need to know and also,
where relevant, will detail any device specific routines that may access
functions/features that are not sufficiently widespread as to form a part
of the general DIMFILM package. There is also an Appendix with general
notes on differences between different classes of device (for example, in
how overlapping vectors are generated and their appearance).

D35   selects the Dicomed in 35mm black and white aperture card mode (on
unperforated film), with landscape aspect ratio. The default coordinate ranges
will be horizontal: 0.,32236. and vertical: 0.,24708., corresponding to a
nominal image size of 36.41x27.94 millimetres on film.

D35P selects
the Dicomed in 35mm black and white (on unperforated film) in portrait aspect
ratio. The default coordinate ranges will be horizontal: 0.,24708. and vertical:
0.,32236., corresponding to a nominal image size of 27.94x36.41 millimetres
on film. This mode is particularly suited to documentation applications
for traditional page layouts.

D16 selects the Dicomed in 16mm black
and white cine mode (on double perforated film), with landscape aspect ratio.
The default coordinate ranges will be horizontal: 0.,32767. and vertical:
0.,23920., corresponding to a nominal image size of 10.26x7.49 millimetres
on film.

D35C selects the Dicomed in 35mm colour slide mode, with landscape
aspect ratio. The default coordinate ranges will be horizontal: 0.,32148.
and vertical: 0.,21698., corresponding to a nominal image size of 36.3x24.5
millimetres on film.

D16C selects the Dicomed in 16mm colour cine mode
(on double perforated film), with landscape aspect ratio. The default coordinate
ranges will be horizontal: 0.,32767. and vertical: 0.,23920., corresponding
to a nominal image size of 10.26x7.49 millimetres on film.

   Other device types will be considered where there is sufficient user demand
(to this end, please advise us of your requirements).

1.1.2 – Control

 Generally, a single job will consist of a number of discrete pictures and
it is required that each of these be separate and the display surface cleared
and readied for the next image.

FRAME is used to achieve this, the effect being appropriate to the
type of device. Thus, film will be advanced on a film recorder, paper advanced
or a new page readied on a pen plotter, or the screen cleared on a display
terminal.

1.1.3 – Termination

 To ensure the plot data is flushed/disposed (as relevant to the device/system)
each DIMFILM execution <u>must</u> be terminated by

DIMEND which will correctly conclude the current DIMFILM execution.

 It is possible to combine more than one DIMFILM execution in a job-step
provided each is correctly initiated and terminated.

1.2 – Plotting Areas

 DIMFILM is designed to operate in the user's chosen coordinate range, with
checks to avoid plotting beyond this area. Each device when used for DIMFILM
initiation will cause its default coordinate ranges to be used for plot bounds.
Generally, it is more convenient for the user to specify his own ranges.
Such user coordinates may be referred to as **world coordinate space** .
Additionally, the user may specify sub-areas to which plotting will be confined
(by clipping at the boundaries) or in which plotting will be prohibited (by
blanking). All these areas are defined by their ranges in each of the x-
and y- directions. Thus, parameters (XLEFT,XRIGHT,YLOW,YUP) define a rectangle
(aligned with the frame edges) running from XLEFT to XRIGHT in the x-direction
and from YLOW to YUP in the y-direction.

1.2.1 – Overall Bounds

BOUNDS(XLEFT,XRIGHT,YLOW,YUP)   is used to set the user's overall plotting
bounds as the rectangle with x-range [XLEFT,XRIGHT] and y-range [YLOW,YUP],
thereby defining world coordinates. This area is mapped to the largest rectangle
of the same aspect ratio that can be fitted on the device display surface.
Clipping will be performed at the boundaries of this area, which may be
redefined at any time. (Redefinition may rescale symbol/marker heights
and broken line patterns.) The current plot position will be set to the
lower left corner of the area, viz. (XLEFT,YLOW). Any currently active preclipping/blanking
areas will be cancelled.

 While the user's plotting is performed within the defined overall bounds
it is possible to select sub-areas to which plotting is restricted or in
which plotting is suppressed.

1.2.2 – Clipping

PANE(XLEFT,XRIGHT,YLOW,YUP)   will define an area to which plotting
is restricted, i.e. all plot data will be pre-clipped at the boundaries of
this rectangle (defined in user's overall coordinates as the rectangle with
x-range [XLEFT,XRIGHT] and y-range [YLOW,YUP]).

ENPANE   cancels the
current pre-clip area, which effectively reverts to the overall bounds.

1.2.3 – Blanking

BLANK(XLEFT,XRIGHT,YLOW,YUP)   will define an area in which plotting
is suppressed, i.e. pre- blanking will be performed at the boundaries of
this rectangle (defined in users's overall coordinates as the rectangle with
x-range {XLEFT,XRIGHT} and y-range {YLOW,YUP}).

ENBLNK   terminates
pre-blanking of the user's plot data.

1.3 – Categories of Drawing

 The components of any plot may be categorised by function. First, there
are the line drawing operations through which the user constructs the picture:
this may be the specific drawing of each graphic element or the application
of functional routines (including graphing and contouring operations). Second,
each plot may include a textual content: this may be specific text defined
by the user or text produced automatically by, for example, axis valuing
in graph plotting. Third, plots may incorporate non-textual annotation:
this might be the outlining of boundaries (for example, the current clipping
pane may be drawn) or axes and their ticking during graph generation.   The
user may wish to visually discriminate between these categories through using
lines of different appearance for each class. The visual appearance of any
line is governed by both the <u>colour</u> in which it is drawn and the <u>style</u>
(being such properties as solid/broken and thickness). For this reason,
DIMFILM categorises the drawing functions into three classes or types – referred
to as **C**olour/**S**tyle Types (CSTypes). Each CSType function will be
drawn with the same line, both in colour and style – these may be set independently
by the user. These groupings of line properties   Similarly, DIMFILM supports
three latent groups of **C**olour/**S**tyle which are accessed according
to the particular plot function. These CSGroups each define a set of line-drawing
properties, including colour/intensity and style, values for which may be
assigned to each group independently. The three CSGroups are associated
with the three CSTypes.   When a drawing operation is performed DIMFILM first
determines the CSType of the function and then ensures plotting is performed
with the corresponding CSGroup colour and style.   This may be summarised:

```
     CSType 1   –  general line drawing
                   most functional plots
```

```
CSType 2  -  all text

CSType 3  -  non-textual annotation (graphs,etc.)
             most boundary plotting
```

When plotting, the corresponding CSGroup is utilised (e.g. for text, a
CSType 2 operation, CSGroup 2 is used). The user may set different colours
and styles for each of the three groups and the relevant one will be used
according to the actual type of the plot operation. (It is possible to set
all three groups to the same colour/style.) It is possible to reassign the
mapping between plot operation and colour/intensity group – a more rigorous
description is given in a later section.

1.4 – Basic Line Drawing

All basic plotting is performed in the user's coordinate space, subject
to any transformations in effect (see later), having regard to any active
pre-clipping/blanking. Lines are drawn at the current colour/intensity and
in the current style – that is, CSType 1 as described at 1.3.2.

1.4.1 – Drawing

OFF2(X,Y)   is used to reposition the current plot position to (X,Y)
without drawing any line.

ON2(X,Y)   causes a line to be drawn from
the current plot position to (X,Y) at the current colour/intensity and in
the current style (i.e. CSType 1).

POINT(X,Y)   will cause a point
to be plotted at coordinates (X,Y), subject to any clipping/blanking (CSType
1).

1.4.2 – Line Style

 The user may select the style in which he wishes lines to be drawn. Style comprises all properties controlling the appearance of the line other than its colour/intensity – i.e. it is the combination attributes such as the continuity and thickness of the line.   It is possible to specify solid (i.e. unbroken or continuous) lines (being the default) or dashed, dotted or dash-dotted lines. A user defined pattern may also be defined (described in a later section).   The following all set the style for <u>all</u> CSGroups:

DASH specifies the style to be dashed lines.

DOT specifies the style
to be dotted lines.

DSHDOT specifies the style to be dash-dotted lines.

DSHOFF cancels any broken line pattern; the style reverts to solid (i.e. continuous) lines (the default).

 It should be noted that any broken line style is continuous, in the sense that the pattern remains in phase over adjacent drawn line segments (thereby enabling curve approximations to be drawn in broken patterns) and across clipped/blanked regions. A pattern will be restarted following any non-drawing relocation of the current position (e.g. by references to OFF2, BOUNDS, FRAME etc.).   The thickness of lines may be specified for all or any of the CSGroups. The user specifies a scale factor by which the default line on any device is to be scaled in thickness. It should be noted that thickness is device related: not all devices can support the same range of line widths. In particular, the default on many devices will be for the finest line width so a factor less than unity will be ignored in such cases. Equally, most devices can only support a discrete range of line thicknesses between a set minimum and maximum. Each device will select an appropriate line thickness depending on its default and the specified scale factor (which by default is unity) – the Appendix on supported devices should be consulted for details of each device's capability.

LINSF(FACTOR)   sets the line thickness scale factor for all CSGroupings (1-3) to FACTOR ($\geq$ 0.0).

LINSFn(FACTOR)   – n = 1,2,3   sets the
line thickness scale factor for CSGroup n (1,2 or 3) to be FACTOR ($\geq$ 0.0).
  In both the foregoing references, the case of FACTOR = 0.0 is treated as a special case whereby the finest line of which any device is capable is selected.

1.4.3 – Point/Spot Style

 By default, DIMFILM generates points, or spots, as a software drawn mark.
However, most devices have a facility to plot a point directly; such a point
is known as a hardware generated point, and the user may specify whether
hardware or software points are required.

HWSPOT will cause all subsequent points to be produced, where supported,
on the device.

SWSPOT will cause all subsequent points to be generated
by DIMFILM software. [This is the default.]

   When software points are produced, the basic spot size will be (approximately)
one ten-thousandth of the maximum device dimension. As with lines, the
size of the basic spot may be scaled using scale factors, and this may be
specified for all or any of the CSGroups. The user specifies a scale factor
by which the default spot on any device is to be scaled. Where hardware spots
are selected, it should be noted that the size is device related: not all
devices can support the same range of spots. In particular, the default
on many devices will be for the smallest spot so a factor less than unity
will be ignored in such cases. Equally, most devices can only support a
discrete range of spot sizes between a set minimumn and maximum. Each device
will select an appropriate spot depending on its default and the specified
scale factor (which by default is unity) – the Appendix on supported devices
should be consulted for details of each device's capability.

SPOTSF(FACTOR)   sets the spot size scale factor for all CSGroupings
(1-3) to FACTOR ($\geq$ 0.0).

SPOTSn(FACTOR)   **– n = 1,2,3**   sets the
spot size scale factor for CSGroup n (1,2 or 3) to be FACTOR ($\geq$ 0.0).
   In both the foregoing references, the case of FACTOR = 0.0 is treated
as a special case for hardware points whereby the smallest spot of which
any device is capable is selected. For software spots, the basic spot is
scaled to create a spot of the required size. For factors above unity the
spot will be (partially) infilled. Factors below one will result in an
appropriately scaled spot; a zero factor, for software spots will result
in no discernible spot being plotted.

1.5 – Basic Text

 The DIMFILM user has available to him a very sophisticated text plotting
"subsystem". Much of the power and flexibility of this will be described
in later sections. At the elementary level the user may control the size,
positioning and orientation of his text string. However, within the actual
text string the user may incorporate certain escape sequences to change the

output characteristics (e.g. underlining, italicisation,super- or sub- scripting, type density etc.) or even to switch to or load an alternate font. DIMFILM offers a range of fonts (initially being a modification and extension of the Hershey occidental data); this range will be expanded as time and manpower permits – again your requirements are sought. The mathematical capabilities of the original DIMFILM are now embodied in the single text output routine. Details of these and the other escape sequences will be described later as they become available. For the purposes of basic text output the user should refrain from the inclusion of characters * (star) and (dollar) in text strings.

1.5.1 – Text Production

SYMTXT(STRING) will output the character string passed in the parameter STRING, which must be of type CHARACTER. The string will be plotted from the current symbol string position – which by default is the current position (normally located via reference to ON2 or OFF2. The string will be produced at the current symbol height and orientation, and (by default) will be plotted in the normal left to right mode. The current symbol colour/style (the CSGroup currently associated with CSType 2) will be used.

The passing of character strings to SYMTXT (and, indeed, to other DIMFILM routines) may be accomplished in a number of ways. The user may pass the string in a <u>simple</u> CHARACTER variable, or as a reference to <u>one</u> element of a CHARACTER array, or by directly defining the string argument in the reference to DIMFILM. When using variable references it is possible to pass substrings, rather than the whole string stored in the simple variable or array element. Consider the following FORTRAN statements:

```
        CHARACTER*26 STRING
        DATA STRING/'abcdefghijklmnopqrstuvwxyz'/

        CALL SYMTXT(STRING)        (i)

        CALL SYMTXT(STRING(4:6)) (ii)
```

The declaration of STRING is of a simple CHARACTER variable of length 26. The first reference to SYMTXT (i) will cause the <u>whole</u> of STRING to be output, i.e. all twenty-six letters of the alphabet will be plotted. The second reference (ii), however, will pass only the **substring** of STRING consisting of characters 4 to 6 (inclusive) to SYMTXT, and the plot will be def only. {FORTRAN 77 will assume an omitted first substring pointer to indicate the beginning of the string, and an omitted second pointer to

indicate the end of the declared string.}   When a CHARACTER array is used
the concept is similar; only a single array element will be processed by
SYMTXT. In this case:

```
        CHARACTER*10 STRING(3)
        DATA STRING(1)/'abcdefghij'/
        DATA STRING(2)/'qrstuvwxyz'/
        DATA STRING(3)/'1234567890'/

        CALL SYMTXT(STRING(3))    (iii)

        CALL SYMTXT(STRING(1))    (iv)

        CALL SYMTXT(STRING(2)(:6)) (v)

        CALL SYMTXT(STRING)       (vi)
```

 Here a CHARACTER array of three elements has been declared, with each element
having a length of ten characters. Reference (iii) to SYMTXT specifically
passes the third element of the array STRING and the plot will comprise the
digits 1234567890. Reference (iv) is to the first element and the plot will
be abcdefghij. The next reference (v) is to the substring consisting of
the first (defaulted) to sixth character in the second array element and
the plot will be qrstuv. The last of these references (vi) is an unsubscripted
reference to the array and by default the first element will be passed to
SYMTXT, and the effect will be identical to that achieved by the specific
element reference at (iv).   Finally, the user is free to specify a constant
string as the argument:

```
        CALL SYMTXT('xyz')
```

 In this instance the whole of the CHARACTER constant will be passed and
the plot will be of xyz.   In general the string will be plotted in the current
default font – this topic is dealt with later.

1.5.2 – Numerics

 Whilst it is usual for numeric values to be passed to SYMTXT in the character
string, there are cases when the user may wish to output the current value
of a variable in his computer program. DIMFILM supports this requirement

for both REAL and INTEGER simple variables (or single array elements) and
permits the user to supply the Fortran FORMAT whereby the output will be
controlled.

RNUM(RNUMB,FMT)   will output a REAL value, according to the supplied
format, at the current symbol string position (see descriptions pertaining
to SYMTXT) at the current symbol height and orientation. This is a CSType
2 function.

   RNUMB holds the REAL value (it may be a reference to an
array element) to be plotted

   FMT is a string containing the Fortran
FORMAT by which the value is to be output, and must be of type CHARACTER.
The FORMAT <u>must</u> commence with an opening parenthesis and contain a matching
closing parenthesis to terminate output. Text within the FORMAT will be
plotted as if it were directly passed to SYMTXT. The FORMAT <u>must</u> be a
valid Fortran specification, comprising only one output record of not more
than 150 characters; the only format specifier should be compatible with
type REAL (i.e. E, F, G, etc.) – scale factors etcetera are supported. However,
only <u>one</u> variable may be output (that is, output lists of several entities
cannot be handled – the user should use continuation strings {see SYMCON/TXTCON
later}).

INUM(INUMB,FMT)   will output an INTEGER value, according to
the supplied format, at the current symbol string position (see descriptions
pertaining to SYMTXT) at the current symbol height and orientation. This
is a CSType 2 function.

   INUMB holds the INTEGER value (it may be
a reference to an array element) to be plotted

   FMT is a string containing
the Fortran FORMAT by which the value is to be output, and must be of type
CHARACTER. The FORMAT <u>must</u> commence with an opening parenthesis and contain
a matching closing parenthesis to terminate output. Text within the FORMAT
will be plotted as if it were directly passed to SYMTXT. The FORMAT <u>must</u>
be a valid Fortran specification, comprising only one output record of not
more than 150 characters; the only format specifier should be compatible
with type INTEGER (e.g. I). However, only <u>one</u> variable may be output (that
is, output lists of several entities cannot be handled – the user should
use continuation strings {see SYMCON/TXTCON later}).

   The following examples of FORTRAN usage should help to clarify the use
of the format parameter.


```
     CALL RNUM(A,'(F10.3)')

        will output the current value of A (REAL) as
        it would be output by the Fortran format
        specification F10.3

     IJK = 2*35 + 3
          .
```

```
        .
    CALL INUM(IJK,'(''THE VALUE OF IJK IS '',I3)')

      will output the current value of IJK (INTEGER)
      according to the Fortran format, giving -

        THE VALUE OF IJK IS  73
```

 In each case the format could have been supplied in a variable of type
CHARACTER. Note also, when single quotes are to be used in a character string
it is necessary to insert two such quotes – the whole string being commenced
and terminated by single quotes.

1.5.3 – Text Size/Orientation

  Text output via SYMTXT will be plotted at the previously specified height
and orientation. This latter may be either relative to the frame or to the
current axes. Once set values hold until redefined.

SYMHT(HEIGHT)   sets the character height for subsequent strings to
HEIGHT, in the dimensions of the users overall coordinates. At first entry
to DIMFILM the symbol height is set to one-fiftieth of the maximum bounds
range.

SYMANG(ANGLE)   specifies that text strings are to be produced
at ANGLE (AGroup1, default degrees – see 1.9 for discussion of angular measure)
to the reference X-direction (which may correspond to either axis or frame)
in the counter-clockwise direction.

RELAX indicates that subsequent text
strings will be oriented with respect to the current positive X-axis.

RELFR indicates
that subsequent text strings will be oriented with respect to the X-boundary
of the frame. (This is the default action.)

1.6 – Markers

  It is possible to plot various marker symbols centred at the current position.

MARKER(NMARK)   cause the marker number NMARK from the current marker
set to be plotted, centred on the current position and utilising the current
colour/intensity (i.e. CSType 1 – basic line plotting rather than text).
This will be plotted at the current marker height. Within any marker font
there is a maximum of 48 symbols, consequently NMARK should be in the range
[1,48].

MARKHT(HEIGHT)   sets the current marker height to be HEIGHT
in units of the user's current overall bounds. At DIMFILM initialisation
this will be set to one-fiftieth of the Y-range.

1.7 – Colour/Intensity

 To support a diverse range of graphic output devices DIMFILM provides considerable
flexibility in the specification of colour and intensity. By adopting a consistent
approach to the interrelationship between these parameters output can be
directed to monochrome or achromatic devices and to true or pseudo colour
devices.   As described elsewhere, DIMFILM supports three latent groups of
colour/style which are accessed according to the particular plot function.
These CSGroups each define a set of line-drawing properties, including
colour/intensity and style. The user may set up different colours/intensities
for each of the three groups and the relevant one will be used dependent
on the nature of the actual plot operation. Alternatively, the user may
set all three groups to the same colour/intensity value.

1.7.1 – Intensity

 For any CSGroup, the relation between colour and intensity is maintained.
Where a colour is specifically set the corresponding intensity will be computed
according to the predefined relationship (covered later). However, where
a specification is made for intensity alone, the colour will be recomputed
to maintain the same hue whilst yielding the desired intensity. Thus for
monochromatic or achromatic use the intensity only need be specified. (It
should be noted that due to the methods of colour mapping not all hues utilise
the whole range of intensity. Appropriate diagnostics are issued and hue
maintained where possible.)

INTSY(ZINT)   sets the intensity for all CSGroupings (1-3) to value
ZINT, which should be in the range [0.,1.]. (For each device this unit range
is mapped onto the available intensities.)

INTSYn(ZINT)   – n = 1,2,3
  sets the intensity for CSGroup n (1,2 or 3) to value ZINT in range [0.,1.].

At DIMFILM initialisation the intensity is set to a default value of 0.8.

1.8 – Colour Models

 DIMFILM provides the user with the means to specify colour either through
reference to a look up table or specifically by definition of a triplet of
colour defining values in any of the range of supported colour models. The
methods may be freely intermixed and the user provided with the maximum flexibility
in attaining the desired result. Each method is described in the following
section on colour selection. The remainder of this section is concerned
with a description of the individual colour models. An appendix will later
be provided defining the various look up table definitions in terms of selected
models; this should serve as a useful guide to the user wishing to specify
his own colour palette.

1.8.1 – General Terms

 Light may be seen as achromatic or chromatic: **achromatic**   light has
no discernable colour – it is a composite of frequencies in the spectrum
yielding "white" light, perceived as a shade in the "grey" scale running
from black–grey–white; **chromatic**   light has a dominant colour and is
(generally) a mixture of frequencies. **Monochromatic** light is a specific
case of chromatic light: it is a single frequency and may be seen as varying
brightness of the one colour.   Achromatic and monochromatic light have a
single attribute – intensity – by which the observer can differentiate output
from different sources. This is the single parameter that is important when
outputting to a black and white device or, indeed, a monochromatic device
(such as a green screen storage tube). Chromatic light – or colour – has,
however a range of sensations experienced by an observer. In order to uniquely
define a colour these attributes must be identified and specified. Colour
is to a large extent subjective. Indeed, the physiological response to "colour"
is still under investigation. However, it is generally accepted that there
are three psycho–physiological characteristics of colour that may be quantified,
and which together can define the observed colour; these are typically **hue**
 (essentially the "colour"), **saturation** (the "purity"), and **brightness**
 ("luminance" or intensity). These will be considered in turn. **hue**
 – hue is used to distinguish the perceived "colour". This could be a single
waveband in the visible spectrum, but would more normally be a group of different
wavelengths that are emitted and result in the (subjectively) perceived colour.
In practice, a dominant (i.e. single) wavelength may be derived which (possibly
mixed with pure white) results in the same physical interpretation by an
observer. Thus, hue may be defined as light of a single frequency corresponding
to the dominant wavelength.   For convenience, the spectrum of visible light
may be arranged as a circle onto which the hue may be mapped. This angular
measure defaults to degrees. A convention has been adopted whereby zero
corresponds to red, with hue increasing counterclockwise through yellow.
Note that complementary colours (see CMY discussion below) are diametrically

opposed; i.e. they lie 180 degrees apart.

This may be portrayed diagrammatically:

```
                           *
                       -       -
        green  120.0  -            -  60.0  yellow
                    *                    *

                -                       -

               -                         -

              -                           -
  cyan  180.0 *         hue spectrum         * 0.0  red
              -                           -

               -                         -

                -                       -

                   *                 *
         blue  240.0  -            -  300.0  magenta
                       -       -
                           *
```
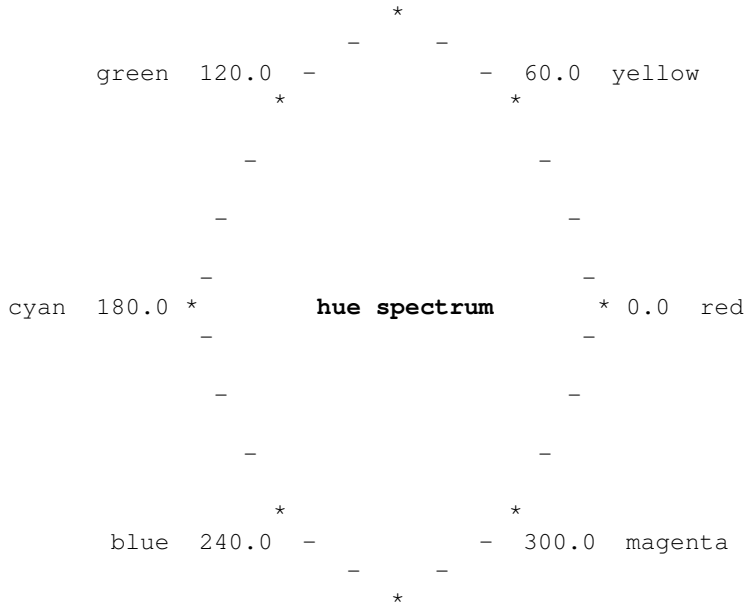
It should be noted that white, gray (any shade), and black are **not**
hues but may be considered as colours.   **saturation** – saturation
is the degree to which a colour is undiluted by white (i.e. the absence of
an amount of its complementary colour, which would effectively add "white").
More strictly this is the purity of a colour. It is expressed on a normalised
scale [0.0,1.0]; for example, pure red is more saturated than pink. It must
be noted that it is independent of brightness; at any brightness a hue may
attain any saturation in the range [0.0,1.0] (excepting the cases of black
and white – nominally corresponding to minimum, or zero, and maximum brightness).
In essence, saturation controls the appearance of a colour running from
the pastel shades to strong (and pure) colours. [In some colour models saturation
is not truly equivalent to purity, as its ratio relates only to the particular
colour subspace modelled.]   **brightness** – brightness (also variously
referred to as lightness, intensity etc.) is equivalent to the intensity
attribute of achromatic light. It is notionally independent of hue and
saturation. Strictly it would be an expression of the luminance of a colour
surface. However, the term brightness (alternatively lightness, intensity)
has become accepted in the very subjective area of colour definition. This
is normally quantified as a normalised value in the range [0.0,1.0]. While
discussing the general definition of colour it is worth considering the terms
used by artists. They are more concerned with describing colours by reference
to the attributes **tints** , **shades** , and **tones** of pure colours
(i.e. the available pigments). From a pure pigment (e.g. one that is fully

saturated), the artist may choose to add white pigment (thereby creating
a **tint** – and decreasing the saturation), or adding black pigment (creating
a **shade** – decreasing the brightness). Addition of both white and black
pigments results in a **tone**. All these changes are variations of a <u>single</u>
hue.

    This may be shown diagrammatically:

```
                    Tints
         White    *--------------------* pure colour
                  |                    /   (i.e. fully
                  |                   /     saturated
                  |                  /          hue)
                  |                 /
                  |                /
                  |               /
                  |              /
                  |             /
                  |    Tones   /
                  |           /
         Grays    |          /   Shades
                  |         /
                  |        /
                  |       /
                  |      /
                  |     /
                  |    /
                  |   /
                  |  /
                  | /
                  |/
         Black    *
```

1.8.2 – Look Up Tables

 DIMFILM contains a 256 entry colour **L**ook **U**p **T**able (**LUT**)
that the user may reference. A simple integer value in the range 0–255 will
access a colour (internally defined by an appropriate value triplet) from
the LUT. The user may redefine individual entries in the table, or the complete
table, either through his own specification of triplet(s) or by reference
to one of the standard LUTs supported by DIMFILM (reference may be to the
complete table or to an individual member). In defining triplets for insertion

into the LUT the user may work in <u>any</u> of the colour models supported in
DIMFILM. Although the default model is the RGB system, the user may well
prefer to work in one of the other models (possibly due to familiarity or
the particular colour properties required). Switching between models is
possible at any time, but will <u>not</u> affect any colours already set. The
use of look-up-tables is dependent on the output device capabilities; the
DIMFILM table of 256 may be excessive for many devices and the user should
ensure familiarity with his device. For many purposes the user will find
need of relatively few colours, but will want this to be discrete and well
defined. With this in mind, all the look up tables offered by DIMFILM will
(unless specifically stated otherwise) have a common set of colours for the
first 16 entries. The remaining entries will be dependent on the nature
of the selected model.

For all look up tables these common entries are:

entry (R ,G ,B)		colour
0		background (device specific)
1		foreground (device specific)
2	(1.,0.,0.)	red
3	(0.,1.,0.)	green
4	(0.,0.,1.)	blue
5	(1.,1.,0.)	yellow
6	(0.,1.,1.)	cyan (blue-green)
7	(1.,0.,1.)	magenta (red-blue)
8	(1.,.5,0.)	red-yellow (orange)
9	(.5,1.,0.)	yellow-green
10	(0.,1.,.5)	green-cyan
11	(0.,.5,1.)	cyan-blue
12	(.5,0.,1.)	blue-magenta
13	(1.,0.,.5)	magenta-red
14	(1.,1.,1.)	white
15	(0.,0.,0.)	black

 Note the foreground/background colours are device defaults (refer
to Appendix on devices). They are usually black and white, although which
is used as foreground and which background will depend on device characteristics.
For example, a pen plotter is likely to have white background (paper) and
black foreground (pen's ink). Conversely, a raster display or film recorder
may have a black background (blank screen/unexposed film) on which the default
drawing action, the foreground colour, will be white. The inclusion of white
and black at 14 and 15 does not imply that all devices are capable of accessing
such colours. The result of drawing in the background colour is also device
specific. Unless a device is capable of doing so, picture elements cannot
be erased dynamically. (E.g. an active film recorder will already have
exposed the vector to be erased and no action, barring total redrawing of
frame, can eliminate that vector. However, on most raster devices the background
colour may be used to erase data and restore the background.) The Appendix
on device types should clarify this situation.

1.8.3 – RGB Colour Model

 The RGB colour model is used in colour television monitors and in most
colour raster displays. The **R**ed, **G**reen and **B**lue primary colours
may be used additively to create any other colour (that is, as primaries
they cannot be decomposed into any combination of the other two chosen primaries).
Note that red, green and blue are commonly referred to as the **additive
primaries.** For the purpose of RGB colour definition, a unit cube defined
by three orthogonal axes is used; with the cartesian axes for the red,
green and blue components each normalised to be in the range [0.0,1.0].
[This is a subspace of three-dimensional colour space.]   Any colour in
this space is uniquely defined by the relative amounts of the three primaries
– i.e. by the triplet value (**red, green, blue**). Of particular
interest is the principal diagonal – this is the achromatic, or grey, scale
– being the set of values with equal components of each of the three primaries;
this runs from black (at the origin) to pure white (when red=green=blue=unity).

The RGB colour cube may be shown diagrammatically:

```
      Green                          Yellow
         *--------------------------*
        /                          /|
       /  |                       / |
      /                          /  |
     /    |                     /   |
    /                          /    |
   /      |                   /     |
  Cyan*------------------------*White |
   |      |                 . |    | |
   |      |                   |    | |
   |      |                   |    | |
   |      |            .      |    | |
   |      |                   |    | |
   |      |                   |    | |
   |      |                   |    | |
   |      |        .          |    | |
   |      |                   |    | |
   |      |      .            |    | |
   |      |                   |    | |
   |      |  .                |    | |
   |      |                   |    | |
   | Black* – – – – – – –| – – *Red
   |      |                   |   /
   |    /                     |  /
   |      |                   | /
   |  /                       |/
   |                          | /
   |/                         |/
    *--------------------------*
     Blue                      Magenta
```
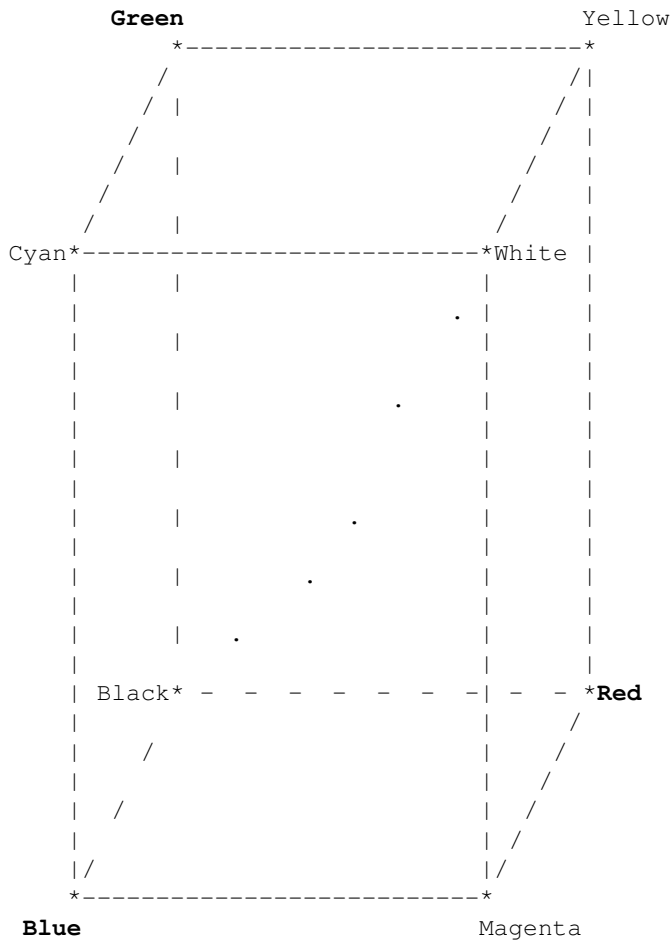
{The orientation of the tri-orthogonal axes R, G and B - as depicted here
- is chosen for two reasons: (i) the R-G-B axes correspond to the normal
three-dimensional cartesian axis convention for X-Y-Z, and (ii) when viewed
as a projection along the achromatic diagonal (from white toward black) there
is a good visual relationship with the hue disc for which red lies at zero
degrees (see Section 1.7.1).}

1.8.4 – CMY Colour Model

 Whereas red, green and blue are known as the additive primaries, their
complements – cyan, magenta and yellow, respectively – are referred to as
the **subtractive primaries**. This is because – on transmitted or reflected
light – they remove or **subtract** their complementary colour. Thus:

```
     -red   = blue  + green = cyan
     -blue  = green + red   = yellow
     -green = red   + blue  = magenta
```

 Indeed, the colour subspace of the CMY system is the identical subspace
in cartesian space as defined by the RGB colour cube. However, an essential
difference is that white light is located at the origin, with black at the
opposite vertex. The additive primaries add to black, while the subtractives
subtract from white light.   The subtractive primaries are of especial importance
in photographic emulsions and in print applications, where the pigments
in print dyes operate in a subtractive mode. Thus, for example, green is
achieved through the combined effects of yellow and cyan inks. Yellow absorbs
blue – leaving green and red – while cyan absorbs red, leaving only green
to be reflected from the surface. Similarly, cyan, magenta and yellow inks
together will reflect no light with a resultant black. The coloured pigment
employed in most hard copy devices (e.g. ink jet, thermal ink transfer and
electrographic techniques) functions in this manner – CMY colour space is
therefore useful in defining these processes.

 The conversion between RGB and CMY triplet values
 is given by:

$$
\begin{vmatrix} C \\ M \end{vmatrix} = \begin{vmatrix} 1 \\ 1 \end{vmatrix} - \begin{vmatrix} R \\ G \end{vmatrix} \qquad \begin{vmatrix} R \\ G \end{vmatrix} = \begin{vmatrix} 1 \\ 1 \end{vmatrix} - \begin{vmatrix} C \\ M \end{vmatrix}
$$

The CMY colour cube may be shown diagrammatically:

```
        Magenta                           Blue
              *─────────────────────────*
             /                         /|
            / |                       / |
           /                         /  |
          /    |                     /   |
         /                          /    |
        /      |                   /     |
   Red*─────────────────────────*Black  |
    |          |                 |       |
    |                       .    |       |
    |          |                 |       |
    |                            |       |
    |          |          .      |       |
    |                            |       |
    |          |                 |       |
    |                            |       |
    |          |          .      |       |
    |                            |       |
    |          |        .        |       |
    |                            |       |
    |          |    .            |       |
    |                            |       |
    |  White* ─ ─ ─ ─ ─ ─ ─|─ ─ *Cyan
    |                        |  /
    |     /                  | /
    |                        |/
    |   /                    | /
    |                        |/
    |/                       |/
    *─────────────────────────*
   Yellow                      Green
```
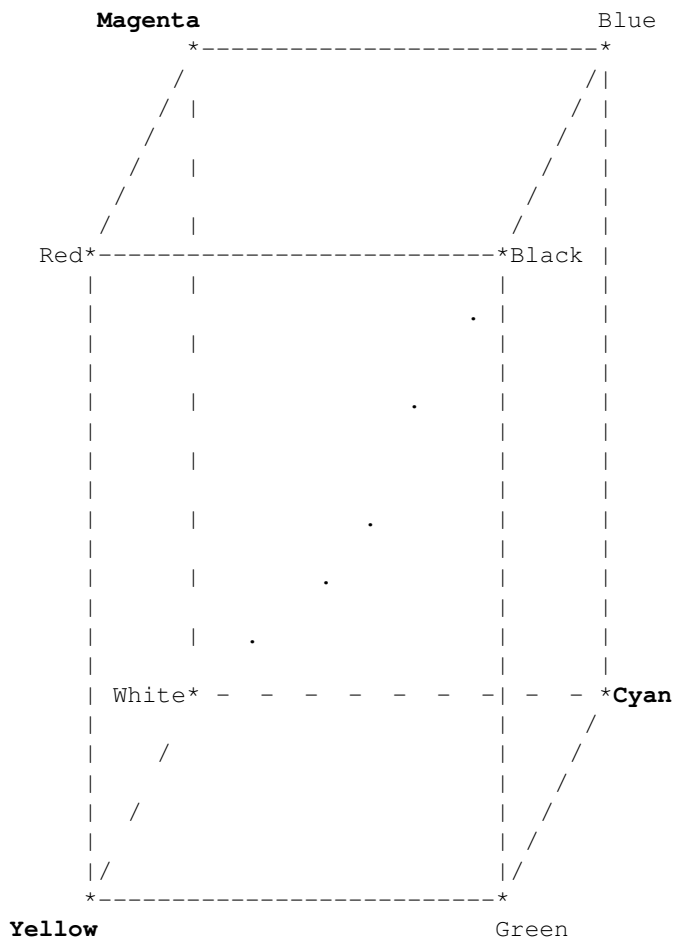
1.8.5 – Broadcast TV Colour Models

 Broadcast television has developed colour models to serve its own particular
requirements. Indeed, the need for a colour methodology arose after a number
of years of monochrome television transmissions. This dictated the overriding
concern for compatibility of a colour signal when received by a monochrome
receiver. The outcome was a value triplet one member of which was a measure
of **brightness** . For computer graphics this has the fortunate result
that by assigning different levels of brightness to different colours one
can be sure that they may be discriminated when displayed on a monochrome
device. The brightness axis runs along the achromatic scale (suited to the
particular display device characteristics); the two remaining axes retain
orthogonality by their definition in the plane perpendicular to brightness.
The direction of one of these axes must be chosen, thereby leaving the plane
containing the remaining axis defined. Essentially, these two axes are
together a measure of the chroma component when all brightness (i.e. the
"white" component) is removed. As indicated above in the discussion of general
terms hue and saturation describe this colour – and, indeed, the two remaining
axes may be related to hue and saturation trigonometrically. However, by
maintaining all three axes in three dimensional cartesian space simple transformation
exists between this space and RGB (as also to CMY). The importance of this
should not be forgotten. Further, a straight line in all these spaces transforms
into a straight line in each of the others – thus, interpolations between
colours may be accomplished simply. [This is **not** true in general for
transforming interpolations between HSV and HLS spaces and RGB, CMY or "television"
colour spaces.]   The definition of "television" colour space depends on the
chosen brightness axis and the orientation of the remaining axes. The first
is selected physically on the basis of what RGB value yields white (not the
ideal model of the diagonal in the colour cube). Here there arise differences
between adopted national standards. PAL (Phase Alternation Line) as used
in Great Britain, Germany and most non-American Western nations has defined
brightness   (**Y**) as:

$$Y = 0.299R + 0.587G + 0.114B$$

 The other two components are known as **U** and **V** , and are chosen
so U is "near-blue" (to the magenta side of blue. V is then at right angles
to U (on the magenta side of red).

   The conversions between **YUV** and **RGB** are then defined as:

$$\begin{vmatrix} R \\ G \\ B \end{vmatrix} = \begin{vmatrix} 1.000 & 0.000 & 0.140 \\ 1.000 & -0.394 & -0.581 \\ 1.000 & 2.020 & 0.000 \end{vmatrix} \begin{vmatrix} Y \\ U \\ V \end{vmatrix}$$

   and

$$\begin{vmatrix} Y \end{vmatrix} \quad \begin{vmatrix} 0.299 & 0.587 & 0.114 \end{vmatrix} \begin{vmatrix} R \end{vmatrix}$$

```
|U| = |-0.148  -0.291   0.439| |G|
|V|   | 5.007  -4.193  -0.814| |B|
```

It is instructive to look at the YUV axes against the RGB
axes. [Note: for this representation the RGB axes have been
rotated from the previously shown orientation in order to
better illustrate the YUV axis selection.]

```
                          |  B
                          |
                          |
                          |
                          |
                          |
                          |
                          |
                          |
                         -|                    /  Y
              U  –   |  –                    /
               \–       |       –           /
                –  \     |         –      /
   sat.            \    |             /
      –      –    .\    |          /   –
         –     .    \   |        /
       –  – . hue \  |     /          –
            –      \ |  /
        –        –   \|/           –
          –        – *-----------------------
          –        –/                    –      G
              – /
        –     –  /               –
            –   /
           –   /              –
         –    /
      V  –       /–                    –
            /   –
           /       –     –
          /           –
         /
        /
       /
      /  R
```
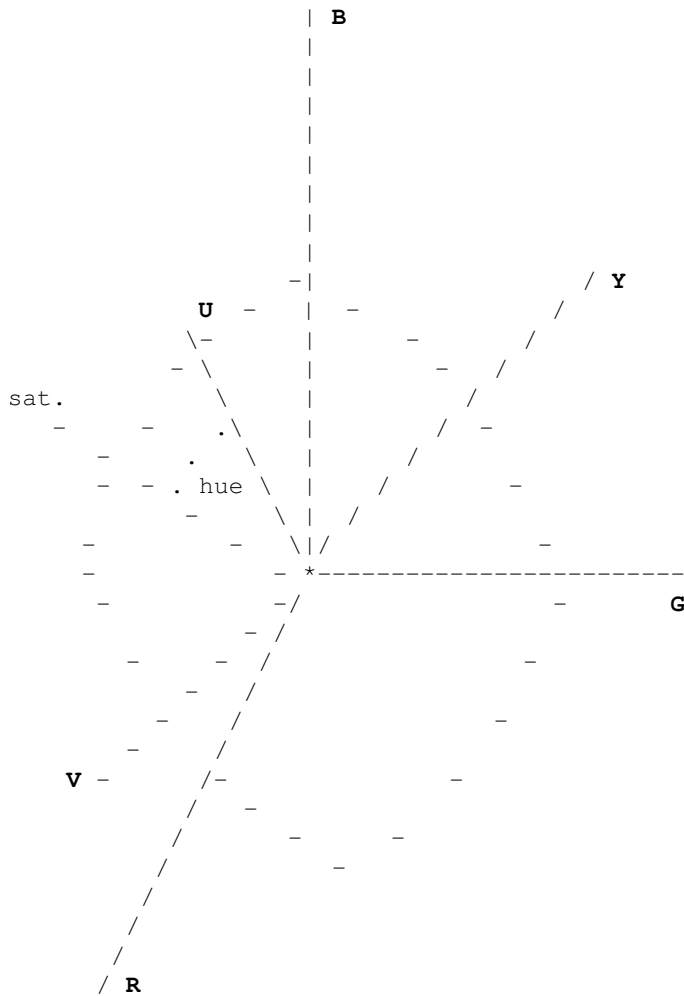
[Note: hue angularly from U, saturation radially outwards.]

 NTSC (National Television System Committee) is the system used in the United
States, South America and Japan. The NTSC definition of **brightness**
(**Y**) is:

     Y = 0.30R + 0.59G + 0.11B

 The other two components are known as **I** and **Q,** and were chosen
with the I axis being aligned with orange-cyan and the Q axis as green-magenta.

   The conversions between **YIQ** and **RGB** are then defined as:

         |R| | 1.00   0.95   0.62| |Y|
         |G| = | 1.00 -0.28 -0.64| |I|
         |B| | 1.00 -1.11   1.73| |Q|

     and

         |Y| | 0.30   0.59   0.11| |R|
         |I| = | 0.60 -0.28 -0.32| |G|
         |Q| | 0.21 -0.52   0.31| |B|

 The differences between the NTSC and PAL models were occasioned by the
criteria used in the location of the I-Q and U-V axes. The NTSC system was
based on the degree of differentiability between colours of the same brightness
for different hues: the eye being particularly perceptive for orange-cyan.
PAL, on the other hand, was designed to minimise errors in hue reproduction
(a disadvantage of NTSC – "Never Twice the Same Colour" – is its liability
to hue distortion, particularly flesh tones) and to simplify the design of
decoders.

1.8.6 − HSV Colour Model

 The above described colour models are primarily hardware related. Although,
for instance, RGB may be used to define colours rationally it does not lend
itself to the artists' concept of colour. Indeed, a perceptual model requires
reference to tints, tones and shades rather than the more physically precise
RGB rectangular coordinates. Compromise models exist: HSV is one of these.
In actual fact, such models are defined in terms related to the display
hardware, but in a way that − at least in practice − is sympathetic to the
subjective visual system. The **HSV** model defines the appearance of a
colour by a value triplet measuring its **H**ue, **S**aturation and (notional)
Value.   If one considers the RGB colour cube it is possible to visualise
the solid from a viewpoint along the extended achromatic diagonal. From
this vantage point a hexagonal projection is seen. (Three planes are visible,
while three others are obscured.)

```
    Viewed along the principal diagonal the RGB colour cube
    will appear thus:




               green  *------------* yellow
                   /.            /\
                  /  .          /  \
                 /    .        /    \
                /      .      /      \
               /        .    /        \
              /          .  /          \
             /            ./white        \
       cyan  *------------ . . . . . * red
             \            .\          /
              \          .  \        /
               \        .    \      /
                \      .      \    /
                 \    .        \  /
                  \.            \/
             blue  *-----------* magenta


   [Note: Visible edges are shown solid, invisible
          edges dotted.  The black vertex is obscured.]
```
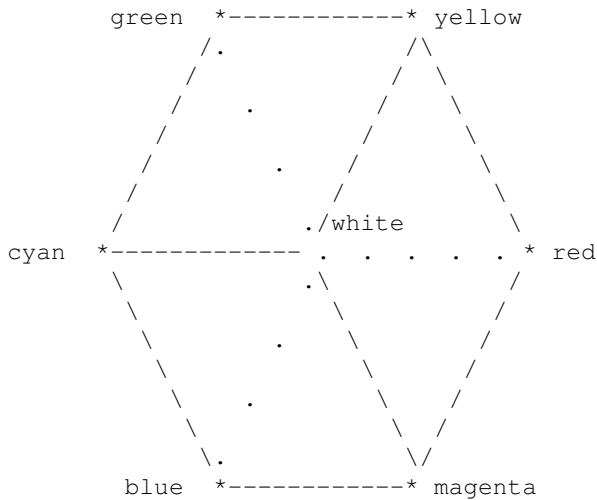
   If the RGB colour cube is considered as a projection onto the plane orthogonal
to the achromatic axis (i.e. perpendicular to the grey diagonal) a hexagonal
disc results. For each **value** of grey there is defined a subcube with
its own projection. As the value is increased from zero (black) the hexagonal
disc will increase in size, until it is at a maximum for the cube defined
by the full (white) diagonal. [Notes: 1 − at black the disc will be a point,

2 – the projection is scaled so the length of a side in the projected disc
equals the length of a side in the solid.] Each disc is the projection of
the three "brightest" faces of the associated subcube; where the defining
diagonal **value** (= v) runs from 0.0 to 1.0 the subcube will be comprised
of sides of that same length; I.e. each disc comprises colours on the surfaces
of the cube defined by r=v, g=v, b=v. So by specifying v at least one of
r,g,b will equal v; conversely, v = max(r,g,b). Thus the **value** is defined
by the grey value, in the range [0.0,1.0]. The **hue** is an angular measure
of the hue, or colour component, in degrees – nominally in the range [0.0,360.0].
The convention adopted here is for red to equate to 0.0, with counter clockwise
rotation (increasing angle) through yellow, green, etc. The remaining member
of the defining triplet is the **saturation**, which is measured radially
outward from the value axis of the hexcone. The saturation is essentially
a ratio in that for any hexagonal disc section it may take any value in the
range [0.0,1.0]; being zero on the axis and unity on the boundary of the
disc (i.e. at the triangular surfaces of the hexcone). Saturation = 0.0 refers
to the central axis and so covers the range of greys; in particular, with
value = 0.0 yields black and = 1.0 gives white – intermediate values are
shades of grey. For the case of zero saturation hue is undefined. At all
other values of saturation the hue must be specified.   The relationship
to the artists perception of colour may be illustrated by considering a
fully saturated hue at maximum value. For example, (0.0,1.0,1.0) [hue =
0.0, saturation = 1.0, value = 1.0] is pure red. The addition of white,
thereby creating tints, is accomplished by reducing the saturation (holding
the value constant). Equivalently, shades may be created by adding black,
and is effected by decreasing the value (the saturation being unchanged).
Decreasing both saturation and value together results in tones.

The HSV single hexcone colour model may be shown
diagrammatically:

```
                              ^ V = value
                              |
              green           |         yellow
                 -------|-------
              -  .      |      . -
            -          |          -
          -            |            -
        -        .     |    .        -
      -                |                -
 cyan *            1.0 * white          * red
   \-       .     |     .            -/
    \  -          |          -  /
     \    -       |       -    /
      \     -  .     .    -    /
       \    -----------------    /
        \  blue \    |    /magenta /
         \       .    .       /
          \      \   |   /      /
           \              /
            \    \  . | .  /    /
             \              /
              \   \   |  /    /
               \    . .    /
                \   \ | /   /
                 \        /
```

```
    \   \ | /   /
     \    \     / -
      \  \|/  /          -
       \    /       -.
        \ | /    -        .H = hue
         \ /-              .
          *-------------------->
     0.0   black   S = saturation
```
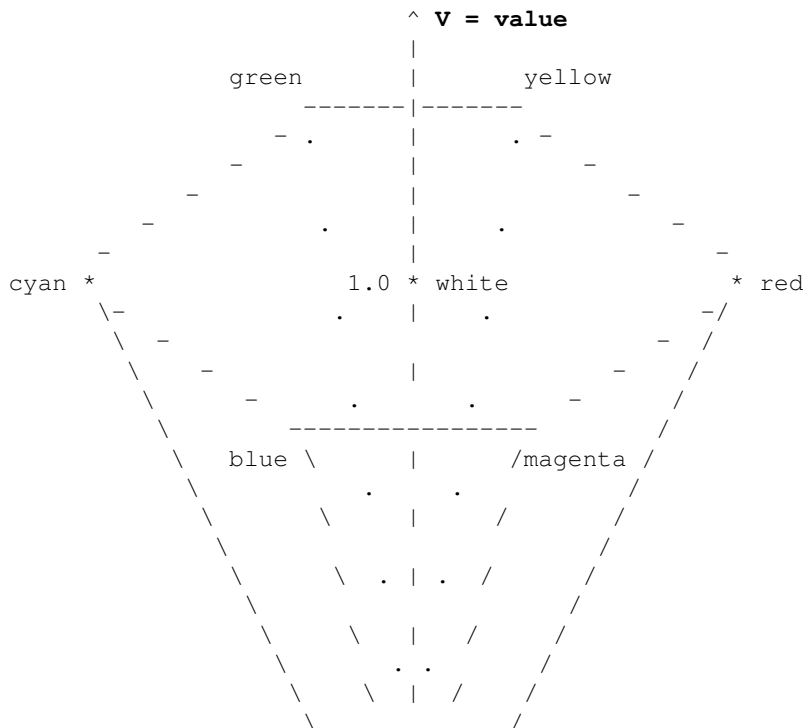
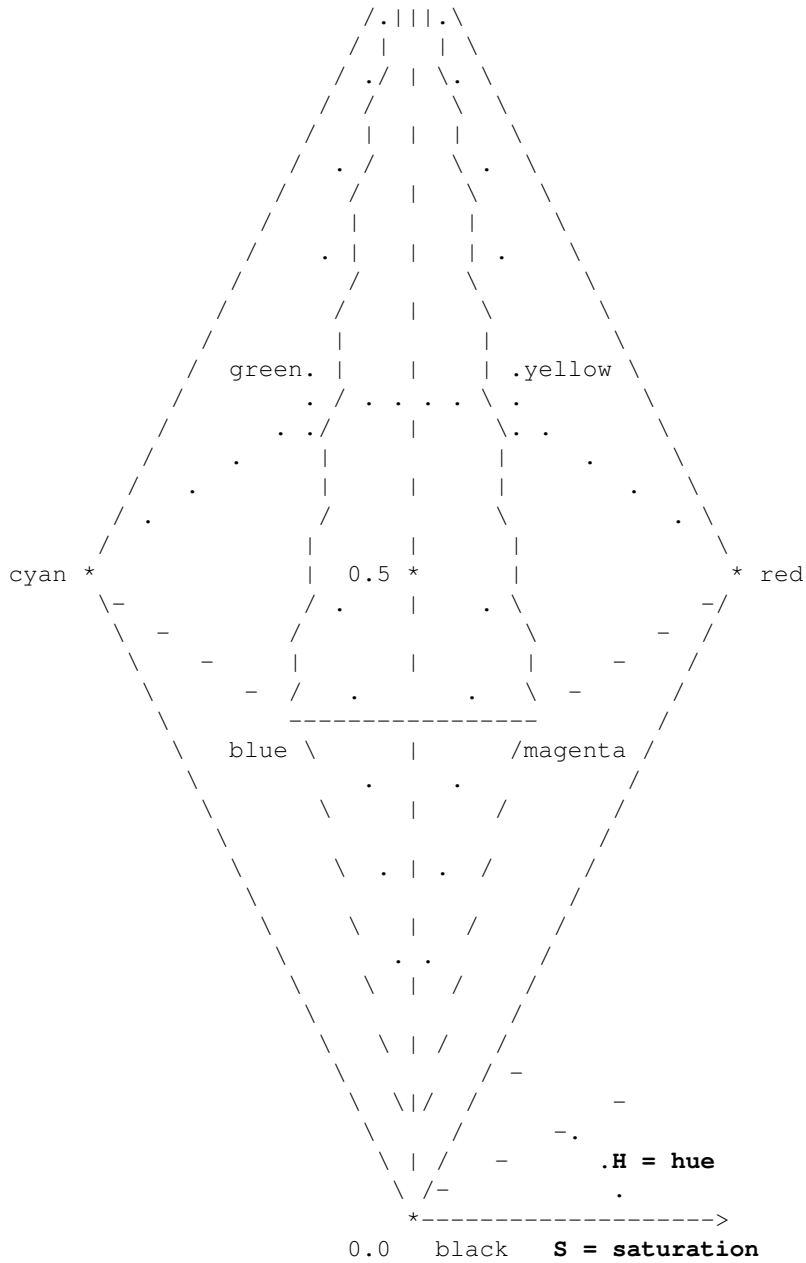   [Note: obscured edges shown dotted.]

1.8.7 – HLS Colour Model

 An alternative to the previously described HSV colour model is the HLS
model. The basic difference is the substitution of **lightness** for **value**;
hue remains as for the HSV model, while saturation is essentially equivalent
(though not numerically identical). However, single hexcone HSV space is
transformed into a double hexcone for the HLS model. This may be considered
to be accomplished by deforming the HSV space with the central white point
being drawn upwards. The central axis remains the achromatic range, and
this is the axis of **lightness** – again running from 0.0 (black) to 1.0
(white). It is important to appreciate that in this model maximally saturated
hues are located at mid-lightness (saturation = 1.0, lightness = 0.5). Considering
the surface of the double hexcone, above the mid-section the surface corresponds
to the upper (V = 1.0) hexagonal disc of the HSV model so **one of** R,G,B
will be maximum at 1.0 and lightness above 0.5 will be achieved by desaturating
this dominant component. Indeed, at L = 1.0 saturation and hue are irrelevant
as pure white results {rgb = (1.0,1.0,1.0)}. While the HLS model has **saturation**
 measured radially (in range (0.0,1.0), as before) the numeric value for
saturation of an <u>identical</u> colour in the HSV and HLS models will differ.
 The achromatic scale occurs when **saturation = 0.0**, when hue may be
undefined, and the degree of grey is determined by the value of **lightness**.
 In neither the HSV nor HLS model does lightness/value equate to luminance
(as, for example, used in the broadcast TV models), and so different colours
may share the same luminance and therefore be represented by the same grey
scale on monochromatic displays.

   The HLS double hexcone colour model may be illustrated:
```
            ^
            |  L = lightness
            |
       1.0 * white
          /|\
         // \\
```

```
                        /.|||.\
                       /  |   | \
                      /  ./ | \.  \
                     /  /     \  \
                    /   |  |  |    \
                   /  . /  |   \ .  \
                  /    /   |    \    \
                 /     |       |      \
                /      |       |       \
               /  green. |     |   | .yellow \
              /      . / . . . . \ .          \
             /     . ./    |      \. .         \
            /    .     |         |     .        \
           /   .       |    |    |        .      \
          / .          /         \          . \
         /            |    |    |              \
       cyan *         |  0.5 *    |            * red
         \-          / .   |    . \          -/
          \   -     /       |       \     -   /
           \    -  |   |    |    |  |  -      /
            \    - /  .    |    .  \ -       /
             \     ----------------         /
              \  blue \    |    /magenta /
               \       .    .          /
                \       \   |   /      /
                 \       |  |  |      /
                  \       \ . | . /  /
                   \       \   |   /      /
                    \       . .       /
                     \       \ | /   /
                      \         /
                       \  \ | /  /   -
                        \   / -
                         \  \|/  /       -
                          \    /     -.
                           \ | /   -     .H = hue
                            \ /-          .
                            *-------------------->
                        0.0   black   S = saturation
```

        [Note: Obscured edges shown dotted.]

   {The HLS convention adopted by Tektronix uses a double cone rather than
a double hexcone. Illustrations of the double hexcone often portray it as
a double cone by topologically deforming it. Additionally, Tektronix use
the convention of blue being located at zero hue. The algorithms incorporated
in DIMFILM for the HSV and HLS models relate to the geometry of the hexagonal
disc and require no trigonometric evaluations.}

1.9 – Colour Selection

 Whereas the preceding section described the various colour models supported
by DIMFILM, the present section is concerned with the specification and selection
of colours for the purpose of drawing.   As described previously (1.3.2),
DIMFILM supports three latent groups of colour/style which are accessed according
to the particular plot function in use. Colours may be specifically assigned
to any or all of these CSGroups. This colour specification may be either
by direct assignment of the value triplet in any of the supported colour
models or by reference to a look-up table (members of which may be set by
the user). Direct specification of a colour results in the computation of
the corresponding intensity (for monochromatic output); however, subsequent
change of the intensity (see 1.6, above) will maintain the hue (in terms
of the proportion of the R,G,B components).   Specification of colour etc.
is considered for each supported colour model separately. For each model,
the triplet value will be checked against range limits for each parameter;
out-of-range parameters will result in a diagnostic and the colour assignment
will be ignored. Internally, DIMFILM will ensure also that the value triplet
transforms to within the RGB colour cube space; the R,G,B values will – if
necessary – be forced into range, noting that this may well result in a change
of hue.   Some DIMFILM subroutines require specification of the triplet of
values defining the colour without specific identification of the colour
model to be used. (This is the case with certain raster operations and modification
of look-up-table entries.)   In these cases the <u>current</u> colour model will
be used in interpreting the triplet.

1.9.1 – Colour model

 DIMFILM is initialised with the RGB colour model as default. Any value
triplets <u>without</u> specific colour model reference will be assumed to be
in terms of RGB colour space and interpreted accordingly. The user may wish
to redefine the current colour model; indeed, it is likely that one of the
other models will be preferable for many applications. All subsequent triplet
definitions (other than those specifically related to other models) will
be interpreted under the new current model. This model may be changed at
will, but is <u>not</u> retrospective. Changing the colour model will not affect
the appearance of any previously assigned colour.

CMODEL(<model>)   sets the default colour model to be that defined by
the parameter,which must be of type character. The parameter may be a character
variable or a character string. The permitted values are:

    RGB

    CMY

    HLS

    HSV

    YUV

    YIQ

Thus, for example, a reference to CMODEL('HLS') will cause subsequent undesignated colour triplets to be interpreted according to the HLS model. However, this does not preclude references to routines such as CMY2 (see below) being used.

1.9.2 – Look Up Tables

There are very many situations when the user will prefer to make reference to look up tables rather than making colour selections on the basis of triplets of real values. Where only a finite range of colours is required (and elaborate facilities, such as fades etc. are not used) the look up table provides a simple and reliable method of selection. As described earlier, the default DIMFILM look up tables offer as the first 16 entries a range of basic colours. The size of the LUTs (256) makes them suitable for most purposes, and a range of different LUTs is provided for different applications.

COLOR(ILUT) sets the colour for all CSGroupings (1-3) to equate to the colour defined by entry ILUT in the **L**ook **U**p **T**able.

COLORn(ILUT)
sets the colour for CSGroup n (1,2 or 3) to equate to the colour defined by entry ILUT in the Look Up Table.

ILUT must be in the range [0,255]. It should be noted that the colour plotted will be that defined by the look up table entry at the time of issuing the plot instruction. Thus it is permissible to change the look up table value after the LUT assignment for any CSGroup has been made. The user may wish to modify any (or all) entries in the current look up table.

SETLUT(ILUT,V1,V2,V3) will set entry ILUT (which must be in the range [0,255]) of the LUT to the colour defined by triplet (V1,V2,V3) in the current colour model.

The user may make such changes at any time; they will affect all subsequent plotting that refers to the LUT. The colour will be evaluated under the colour model current at the time of definition; however, there is no restriction that all entries in a LUT shall be defined in the same model. The user may change the colour model and enter different elements of the LUT in those models. The user may, at any time, reinitialise the current LUT, either by reloading the default LUT or any of the others available in DIMFILM. Any changes made to the LUT will be lost.

LDLUT0 will reinitialise the active LUT with the DIMFILM default look up table.

1.9.3 – RGB colour model

RGB(R,G,B)   sets the colour for all CSGroupings (1-3) to equate to
the colour defined by the value triplet (R,G,B) in the RGB colour model.

RGBn(R,G,B) **– n = 1,2,3**   sets the colour for CSGroup n (1,2 or 3)
to be (R,G,B) in terms of the RGB colour model.

  R,G,B are each in the range [0.0,1.0]. {Note: Greys may be achieved by
setting R=G=B=v, for v in [0.0,1.0]; with black the special case v=0.0, and
white v= 1.0.}

1.9.4 – CMY colour model

CMY(C,M,Y)   sets the colour for all CSGroupings (1-3) to equate to
the colour defined by the value triplet (C,M,Y) in the CMY colour model.

CMYn(C,M,Y)   **– n = 1,2,3**   sets the colour for CSGroup n (1,2 or 3)
to be (C,M,Y) in terms of the CMY colour model.

  C,M,Y are each in the range [0.0,1.0]. {Note: Greys may be achieved by
setting C=M=Y=v, for v in [0.0,1.0]; with black the special case v=1.0, and
white v= 0.0.}

1.9.5 – PAL colour model (YUV)

YUV(Y,U,V)   sets the colour for all CSGroupings (1-3) to equate to

the colour defined by the value triplet (Y,U,V) in the YUV colour model
(i.e. that used in PAL broadcast TV).

YUVn(Y,U,V)   – n = 1,2,3   sets
the colour for CSGroup n (1,2 or 3) to be (Y,U,V) in terms of the YUV (PAL)
colour model.

  Y is in the range [0.0,1.0]. U is checked against the range [-0.439,0.439],
and V against [-5.007,5.007]; noting that this space is a superset of RGB
space and so certain combinations of YUV may not map into the RGB unit cube.
Care should therefore be exercised to ensure mapping onto RGB is possible.
Where the mapping onto RGB space results in a colour beyond the unit cube
truncation will be applied to the out-of-range R,G,B component – and hue
is not retained.

1.9.6 – NTSC colour model (YIQ)

YIQ(Y,I,Q)   – I is real   sets the colour for all CSGroupings (1-3)
to equate to the colour defined by the value triplet (Y,I,Q) in the YIQ
colour model (i.e. that used in NTSC broadcsat TV).

YIQn(Y,I,Q)   – n
= 1,2,3   I is real   sets the colour for CSGroup n (1,2 or 3) to be (Y,I,Q)
in terms of the YIQ (NTSC) colour model.

  Y is in the range [0.0,1.0]. I (which is a REAL variable) is checked
against the range [-0.60,0.60], and Q against [-0.52,0.52]; noting that
this space is a superset of RGB space and so certain combinations of YIQ
may not map into the RGB unit cube. Care should therefore be exercised to
ensure mapping onto RGB is possible. Where the mapping onto RGB space results
in a colour beyond the unit cube truncation will be applied to the out-of-range
R,G,B component – and hue is not retained.

1.9.7 – HSV colour model

HSV(H,S,V)   sets the colour for all CSGroupings (1-3) to equate to
the colour defined by the value triplet (H,S,V) in the HSV colour model.

HSVn(H,S,V)   – n = 1,2,3   sets the colour for CSGroup n (1,2 or 3)

to be (H,S,V) in terms of the HSV colour model.

  H – the hue – is in the range [0.0,360.0) [although it will be forced
modulus 360.0], while S and V are in the range [0.0,1.0]. {Note: Greys may
be achieved with S = 0.0, and V in [0.0,1.0]; with black the special case
V=0.0, and white V= 1.0. For zero saturation hue is undefined.}

1.9.8 – HLS colour model

HLS(H,L,S)   **– L is real**   sets the colour for all CSGroupings (1–3)
to equate to the colour defined by the value triplet (H,L,S) in the HLS
colour model.

HLSn(H,L,S)   **– n = 1,2,3**   **L is real**   sets the colour
for CSGroup n (1,2 or 3) to be (H,L,S) in terms of the HLS colour model.

  H – the hue – is in the range [0.0,360.0) [although it will be forced
modulus 360.0], while L (which is a REAL variable) and L are in the range
[0.0,1.0]. {Note: Greys may be achieved with S = 0.0, and L in [0.0,1.0];
with black the special case L=0.0, and white L= 1.0. For zero saturation
hue is undefined. Fully saturated hues are attained at L = 0.5.}

1.10 – Colour/Style Groups and Types

 All plotted/drawn lines may have a number of properties associated with
them; these may govern the appearance by specifying colour (and/or intensity),
thickness or style (e.g. solid or type of broken pattern). For the purpose
of assigning colour/style attributes, DIMFILM recognises a number of classifications
of plot operations. These classes are referred to as colour/style types
(CSType). The primary type-class comprises line drawing operations (CSType
1). The secondary type-class (CSType 2) is text output, while the tertiary
type-class (CSType 3) includes non-textual annotation of graphical plots.
Functional plot routines, with the exception of those for boundary plotting,
are performed as primary type operations (i.e. CSType 1). The exceptions
(primarily including OUTLIN, OWS, OWSVP, OWSWIN, OWCWIN, MARGIN, OPANE and
OBLANK) are plotted under CSType 3. Any other exception will be noted in
the routine description.   DIMFILM supports three <u>latent</u> groups of colour/style
attributes which are accessed according to the class of the particular plot
function. These colour/style groups will be subsequently referenced in the
form CSGroup n (n = 1,2,3). In number these CSGroups correspond to the
CSType classes. Essentially, each CSType has an associated CSGroup (although
the mapping need not be 1:1); i.e. at any time there will be a particular

colour/style set of attributes associated with each plot-type classification.
 For example: assuming the default mapping, the line produced by a call
to GRAPH (see later) is of CSType 1 and so would be drawn in the colour and
style currently defined for CSGroup 1; text (such as captions and labels
produced, for example, by GRDEF) is of CSType 2 and would be in the colour
and style set by CSGroup 2; the axes and framing of the graph (which may
also be performed by GRDEF) are of CSType 3 and would be in the colour and
style of CSGroup 3.   The assignment of plotting operations into CSTypes is
fixed. However, the user is free to define the attributes for each CSGroup
(via the colour, intensity and line style routines described elsewhere)
and also to assign the mapping of CSGroup to CSType.   Initially the mapping
is linear – with CSType 1 functions being plotted with CSGroup 1 colour/style
attributes, CSType 2 with CSGroup 2, and CSType 3 with CSGroup 3.   The user
may set up different colours/styles for each of the three CSGroups and the
relevant one will be used dependent on the nature of the actual plot operation
(i.e. determined by CSType of operation and mapping onto CSGroup). Alternatively,
the user may set all three groups to the same colour/style. By reassigning
the mapping between plot operation and colour/style group the user may cause
different types of operation to be assigned the same colour/style group
It is important to remember that conceptually there are three <u>latent</u> drawing
tools – the relevant one being automatically used for any task (and hence
becoming <u>active</u>). Each, independently or collectively, may have different
properties specified. DIMFILM will also offer the facility of a resident
pool of inactive "pens" from which the latent ones may be assigned.   [After
each plotting task is performed the currently active task type reverts to
1, and the currently active colour/style will be the CSGroup associated with
CSType 1.]

COLPT1(N) associates CSGroup N with CSType 1

COLPT2(N) associates
CSGroup N with CSType 2

COLPT3(N) associates CSGroup N with CSType 3
  {In the foregoing, N must be in the range [1,3]}

COLSET reset the
default association; i.e. CSType n is associated with CSGroup n.

  Referring to the earlier example with GRAPH, if it is desired that the
axes and boundary annotation have the same appearance as the text, it could
be accomplished by a reference to COLPT3(2), which would associate CSType
3 functions (non-textual annotation) with the same group as used by text
(CSType 2) – CSGroup 2. The user would then only need to set colour and style
for these two groups.

1.11 – Angular Measure

 There are various units of angular measure that may be used. There are
often different requirements as to which set of units is most appropriate
for a particular task, and it may be that the user will wish to use different
units for different tasks rather than an overall election of one unit of

angular measure.   DIMFILM offers the user three alternate types of angular
measure, namely degrees, radians, or grads. Additionally, plot parameters
with an angular specification are grouped into four Angular Groups (AGroups),
each of which may have an independent specification of unit of angular measure.
 The groupings and default units are given in tha accompanying table.

Assignment of angular arguments to groups

AGroup	1	2	3	4
Default	Degrees	Radians	Degrees	Degrees
	SYMANG	TRANGE	CIRARC	PARAL2
	ROTATE	POLAR	DTHETA	BOX
	HATCH	POLFN	ELARC1	POLYGN
	CLABHT	RADSEP		
	SHDEGR	POLTOA		
		POLPY3		
		POLPY5		
		POLFN		
		POLON		
		POLOFF		
		TLEVS		
		ELARC2 –	ELARC2 –	
		PHI1/2	PSI	

   The relevant angular group is noted alongside each primary parameter description.
 The group assignment may be changed, or all groups set to one angular unit.
A number of routines exist for these purposes.

RADIAN sets radians as the unit if angular measure for all AGroups.

DEGREE sets degrees as the unit if angular measure for all AGroups.

GRAD sets grads as the unit if angular measure for all AGroups.

RADGRP(N) set
radians as the unit of angular measure for AGroupN, for N=1 to 4.

DEGGRP(N) set
degrees as the unit of angular measure for AGroupN, for N=1 to 4.

GRAGRP(N) set
grads as the unit of angular measure for AGroupN, for N=1 to 4.

Chapter 2 – Refined Plotting

2.1 – Textual Plotting

 DIMFILM includes very versatile and comprehensive facilities for the output
of text strings. The user has full control over the characteristics of a
string and may change these from within the string being plotted. This section
gives a detailed description of text control and plotting. The user should
note that all switches (i.e. escape sequences) may be incorporated in strings
passed to the graphical annotation routines; the only limitation there being
the inability to continue text strings and the restriction to string length
being imposed by the Fortran implementation used.

2.1.1 – Text Founts

  (In accord with common computer parlance, the word font will hereafter
be used for fount – referring to a set of "printable" characters.) DIMFILM,
as standard, has several fonts permanently resident in memory for immediate
access by the user. These fonts comprise three for alphabets and one each
for special symbols and markers. Each alphabet font – numbered 1 to 3 – has
a maximum capacity of 96 characters and 24 accents (noting that these latter
may be combined), while the symbol font has a maximum capacity of 96 and
the marker font a maximum of 48 "characters". The DIMFILM system has access
to a considerable number of alternate fonts which may be loaded as the resident
fonts, either by a subroutine reference or from within a text string. Switching
between resident fonts is a more economical process and is to be preferred
where possible.   The fonts are arranged to reflect the printable set of 7–bit
ASCII, and are mapped onto this set from the host machine. The user need
not concern himself with the mapping mechanism while using character strings.
[However, the characters available to a user in a character string within
a Fortran 77 program are dependent on two factors: i) the particular Fortran
implementation, ii) the input character set supported by the host &/or telecommunications
system.] To simplify transport of DIMFILM programs it is possible to access
the whole character set through use of the standard Fortran 77 character
set.   For reasons that will become apparent, the 96 character font is considered
to be arranged in 6 columns (each of 16 rows). The layout is as shown:

DIMFILM Font Layout

	Upper Case			Lower Case		
	0	16	32	48	64	80
1		0	@	P		`	p
2	!	1	A	Q		a	q
3	"	2	B	R		b	r
4	#	3	C	S		c	s
5	$	4	D	T		d	t
6	%	5	E	U		e	u
7	&	6	F	V		f	v
8	'	7	G	W		g	w
9	(	8	H	X		h	x
10	)	9	I	Y		i	y
11	*	:	J	Z		j	z
12	+	;	K	[		k	{
13	,	<	L	\		l	|
14	–	=	M	]		m	}
15	.	>	N	^		n	^
16	/	?	O	_		o

Description of non-alphanumerics

1		space	28	;	semicolon
2	!	exclamation point	29	<	less than
3	"	quotation marks	30	=	equals
4	#	number sign	31	>	greater than
5	$	dollar	32	?	question mark
6	%	percent	33	@	commercial at
7	&	ampersand	60	[	opening bracket
8	'	apostrophe	61	\	reverse slant
9	(	opening parenthesis	62	]	closing bracket
10	)	closing parenthesis	63	^	circumflex
11	*	asterisk	64	_	underline
12	+	plus	65	`	grave accent
13	,	comma	92	{	opening brace
14	–	hyphen (minus)	93	|	vertical line
15	.	period	94	}	closing brace
16	/	slant	95	^	overline (tilde)
27	:	colon

NOTE:– The Fortran 77 characters set comprises

Alphabetic	A to Z
Numeric	0 to 9
Special	= equal
Characters	+ plus
	– minus

```
* asterisk
/ slash
( opening parenthesis
) closing parenthesis
, comma
. period (decimal point)
$ currency symbol (dollar)
' apostrophe
: colon
" quotation mark
  blank (space)
```

There is a range of fonts available to the DIMFILM user – and more will be added in the future. Traditionally, fonts are designed for reproduction at a specific size and with a standard appearance. To enable output with the wide range of variation permitted by such a system as DIMFILM compromise has to be made. In expanding a character beyond its designated point size smoothness and solidity may become degraded. According to demand, it is intended to enhance the range of fonts (by way of style and alphabet) – consequently user requirements should be made known. One particular difficulty in font design is the output attributes of mono or proportional spacing; each font may be ideally suited to only one of these modes. The present fonts are best output proportionally spaced (this is the default). It is intended to incorporate a selection of fonts that will be more suited to mono spacing.   Each font is identified to DIMFILM by an integer value. The present set of fonts (with residency denoted in braces) comprises:

Roman (English) Alphabetic Fonts

```
       sans-serif          seriffed

    11   Simplex {1}   14    Small Complex
    12   Duplex        15    Complex {2}
    111  Mono/Simplex  16    Triplex
                       24    Small Complex Italic
                       25    Complex Italic
                       26    Triplex Italic
                       31    Simplex Script
                       35    Complex Script
                       46    Gothic English Triplex
                       56    Gothic German Triplex
                       66    Gothic Italian Triplex
```

<u>Greek Alphabetic Fonts</u>

```
        sans-serif          seriffed

    1011   Simplex       1014   Small Complex
                         1015   Complex {3}
```

<u>Cyrillic Alphabetic Fonts</u>

```
        sans-serif          seriffed

                         2015   Complex
```

<u>Symbol Fonts</u>

```
                  8001      Mathematical (1) {S}
                  8002      Cartographic
                  8003      Astronomical
                  8004      Astrological
                  8005      Musical
```

<u>Marker Fonts</u>

```
                  9001      Principal Marker Set {M}
```

Notes

1: "small" in font description indicates a font that is suitable for
output at a smaller size; this size is directly controlled by the user.
Character sizes for all fonts are set by SYMHT. (Such fonts are generally
characterized by the individual character definition employing fewer coordinates;
spacing between multiple "strokes" for small, complex characters is also
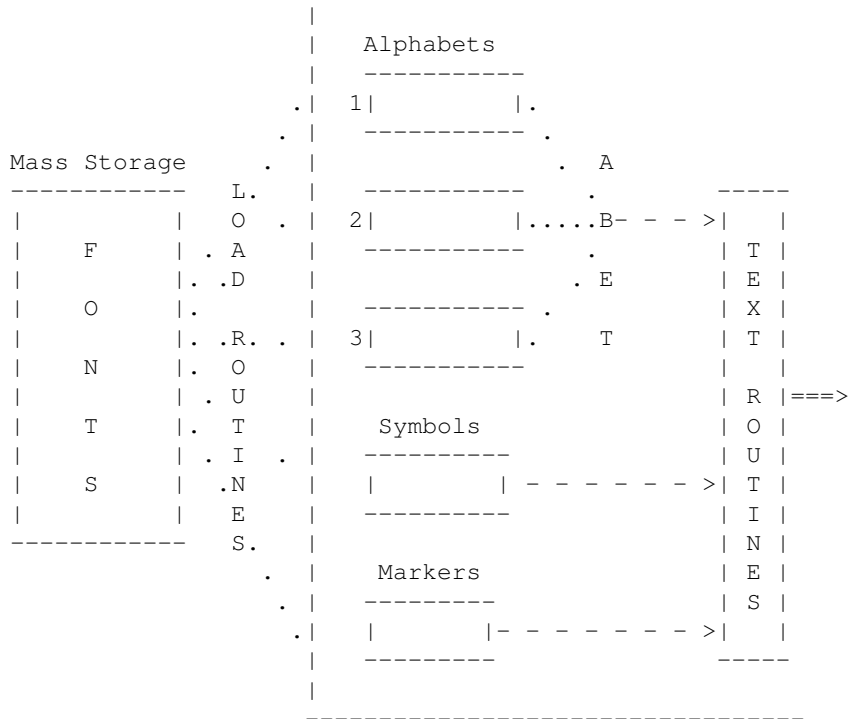more suited to output at lesser sizes.)

2: Default assignment of fonts
to alphabets is indicated, being:

```
            alphabet 1 – font 11,
            alphabet 2 – font 15,
            alphabet 3 – font 1015,
            symbols – font 8001,
            markers – font 9001.
```

   The user may load any available font into any of the memory resident set
(appropriate action is taken where selected font is incompatible with load
destination; i.e. it is possible to load symbols or markers into alphabet
fonts, which will not thereby be completely filled, etc.).   Through restricting
the number of memory resident fonts that are available to it, DIMFILM seeks
to limit its usage of central processor memory. All other fonts are held
on a mass-storage device from which they may be acquired (in place of one
of the memory resident set) as needed – this being at the users control.
This enables DIMFILM to access a very wide range of fonts, a range that
may be periodically updated. Pictorially, the process may be shown:

```
            Memory
---------------------------------
```

```
                      |
                      |   Alphabets
                      |   -----------
                    .|  1|           |.
                    . |   ----------- .
   Mass Storage      . |                    .  A
   -----------    L.  |   -----------     .        -----
   |           |  O . |  2|           |.....B- - - >|   |
   |    F      | . A  |   -----------      .        | T |
   |           ||. .D |                   . E       | E |
   |    O      ||.    |   -----------  .             | X |
   |           ||. .R . |  3|          |.    T        | T |
   |    N      |. O   |   -----------                |   |
   |           | . U  |                              | R |===>
   |    T      |. T   |   Symbols                    | O |
   |           | . I . |   ---------                 | U |
   |    S      |  .N  |   |         | - - - - - >| T |
   |           |  E   |   ---------                 | I |
   -----------   S.   |                              | N |
               . |   Markers                    | E |
              . |   ---------                    | S |
             .|   |         |- - - - - - >|    |
              |   ---------                  -----
              |
              --------------------------------
```

LDABET(IBET,NFONT)   will load font number NFONT into memory resident
alphabet IBET; where IBET = 1,2, or 3.

LDSYM(NFONT)   will load font
number NFONT into memory resident symbols.

LDMARK(NFONT)   will load
font number NFONT into memory resident markers.

2.1.2 – Initial Text Characteristics

There are a number of parameters that characterize text appearance other than size and orientation. For example , mono or proportional spacing, colour, density, italicisation. Such entities are, for the purposes of DIMFILM text output, known as the Text Characteristics. In particular, those which are active at the commencement of text form a set known as the <u>Initial Text Characteristics</u> (subsequently referred to as ITC). It is possible to change most characteristics from within a string, and the set active at any time is known as the <u>Current Text Characteristics</u> (CTC). By default, each output string will commence with the initial characteristics active; i.e. the CTC will be set to the ITC. The user may set certain characteristics of the ITC independently. The ITC comprises:

alphabet – the active alphabet from the memory resident set of alphabetic
fonts. {Default = **1**}

mapping – the mapping from the input string
may be set to one of three modes. <u>Mixed</u> identically maps the input character onto the DIMFILM font; <u>Lower</u> maps upper case input characters onto the lower case set (viz. entries in columns 3 and 4 map directly across to columns 5 and 6); <u>Upper</u> maps lower case input characters onto the upper case set (viz. entries in columns 5 and 6 map directly across to columns 3 and 4). (Columns 1 and 2 are unaffected by any mapping. Note also that use of lower case input characters is outside the scope of Standard Fortran 77.) {Default = **Mixed**}

pen – two "pens" are available for text production (each
having an associated colour/intensity). There is a <u>standard</u> colour/intensity for text and an <u>alternate</u> colour/intensity. {Default = **Standard**}

spacing – text may be produced with either <u>monotonic</u> or <u>proportional</u>
spacing. {Default = **Proportional**}

italics – text may optionally
be italicised (i.e. a slant will superimposed on each output character). Note that this characteristic is <u>additional</u> to any inherent in the font definition; if italics are set they will be superimposed on an italicised font. {Default = **italics off**}

underline – text may optionally
be <u>underlined</u>. {Default = **underline off**}

heavy – text may
be in <u>normal</u> density or <u>heavy</u> density. This affects only the intensity of the text (where supported on a device) and should not be confused with bold type which also employs line thickening. (It is possible for the intensity shift to be specified as a reduction when lightening of the density will result.) {Default = **normal**}

fonts – font assignment to each
resident alphabet, special symbol and marker set. {default = **see font descriptions**}

all other – all other text characteristics cannot
be set other than from within a string and always revert to their default state at the end of the input string. Such characteristics include superscripts/subscripts/fr to symbols and markers etc. Such states are **always** cleared at commencement/termination of each string passed via SYMTXT etc. regardless of election for text continuation.

Each string will commence at the base level (see later discussion of fractions
and super-/sub- scripts). {Default = **inactive**}

2.1.3 – **User Specification of ITC**

 The user may set the various Initial Text Characteristics and certain values
associated with them. Once changed they remain in force until reset, either
explicitly or implicitly (e.g. by total reset of all text characteristics).
The following relate to the setting of ITC values.

ABET(IBET)   selects the memory resident alphabet that is to be used.
IBET must be in range [1,3]; default is 1.

UPPER forces input characters
to be mapped onto the upper case set.

LOWER forces input characters
to be mapped onto the lower case set.

NOMAP accepts input characters
without mapping. This is the default.

ALTTYP specifies that the alternate
text pen is to be used.

DEFTYP specifies that the default (standard)
text pen is to be used. This is the default (and will normally access colour/intensity
group 2).

ALTPEN(IPEN)   sets the alternate text pen to group IPEN,
which should be in range [1,3]. The default is 3.

MONO selects mono
spacing.

PROP selects proportional spacing. This is the default.

ITALIC selects
italicisation of text fonts.

NOITAL turns off any italicisation. This
is the default.

SETITL(TANANG)   sets the italicisation to the angle
with tangent TANANG; the angle being that at which a vertical stroke would
appear measured clockwise from the upward normal to the string direction.
The default value being 0.2.

UNLIN selects underlining.

NOUNLN turns
off underlining. This is the default.

HVYTYP selects heavy type.

NRMTYP selects
normal type density. This is the default.

HVYINT(ZINT)   sets the intensity
increment to be applied to the normal type intensity as ZINT. (Note: this
may be negative, when HVYTYP would result in a lighter type.) The default
value is 0.15.

2.1.4 − Current Text Characteristics

 There are a number of parameters, in addition to the Initial Text Characteristics,
that determine the appearance of text. DIMFILM provides powerful facilities
(by way of escape sequences) whereby the user may modify the ITCs within
a string and to effect other appearance/layout changes. At any time the
set of active parameters governing text appearance is known as the Current
Text Characteristics (CTC) − this includes the ITC parameters as a subset,
but additionally incorporates a number of dynamic characteristics.   It is
reiterated that the subset of the CTC equivalent to the ITC is normally
reset identically to the ITC at the commencement of each string, while the
remaining parameters are turned off at such time (e.g. fractions are assumed
inactive at string commencement, generally being sensibly terminated at the
conclusion of each string).   The CTC comprises the ITC (each of which may
be dynamically modified) and additionally:

fractions − two levels of fraction beyond the base text level being
supported

super/sub scripts   − up to six levels of super &/or sub
scripts is supported (beyond the base level) at each fraction level (including
the base text level).

   Conveniently considered with the CTC are the escape sequences used to
access accents, special symbols, and markers, and those to perform other
transient effects (e.g. backspace, etc.). Each of these while active may
be deemed a CTC; however, their effects are terminated once executed − i.e.
they have a strictly limited scope.

2.1.5 − User Specification of CTC

 The user may set the various Current Text Characteristics, and certain
values associated with them, from within a text string. Once changed they

remain in force until reset, either explicitly or implicitly (e.g. by total
reset of CTC to ITC at end of string, or through cessation of their effective
scope). Within a text string (passed as type CHARACTER) two characters have
special significance. These are * (asterisk) and $ (dollar/currency symbol).
Generally, these are taken to indicate commencement of an "escape sequence"
[following character(s) determining escape function]. To produce either
of these characters in text it is necessary to replicate it. Thus, for
each single * required in text it is necessary to include the double **;
likewise each occurrence of $$ will yield a single output $. All other
occurrences of * or $ will either produce a CTC change or result in a diagnostic
(possibly an unrecognised escape sequence or one that is invalid at that
point). The * and $ have notionally different functions. As a rule,
escape sequences flagged by * initiate an effect, while those flagged by
$ terminate an effect (certain actions have no corresponding termination;
e.g. spacing is specified to be mono or proportional – the * sequence used
logically in each case).   The available escape sequences are given here;
all are initiated by * and so the given character syntax for the escape should
be preceded by *. If termination is applicable this is specifically stated.
Only characters in the standard Fortran 77 character set are used for escape
sequences (in particular, <u>upper case</u> letters are mandatory). In the following
specifications of escape sequences, certain escape fields are "variable"
(and should be entered as appropriate by the user) – these are denoted by
lower case strings delimited by < and >. {Some examples follow the escape
definitions.}

U – upper case mapping. Input lower case characters (i.e. those
from columns 5 or 6) are mapped onto the upper case set (i.e. columns 3 or
4, respectively). {Termination by **$U**   causes mixed (or no) mapping
to become effective.}

L – lower case mapping. Input upper case characters
(i.e. those from columns 3 or 4) are mapped onto the lower case set (i.e.
columns 5 or 6, respectively). {Termination by **$L**   causes mixed (or
no) mapping to become effective.}

1 – sets the current alphabet to
resident alphabet set 1.

2 – sets the current alphabet to resident
alphabet set 2.

3 – sets the current alphabet to resident alphabet
set 3.

+ – initiates a superscript, incrementing the super/subscript
level for the current fraction. {Termination by **$+** cancels the <u>last</u>
<u>superscript</u> (and any intermediate subscripts) for the current fraction,
decrementing the current fraction super/subscript level for each. See also
$0 below.}

– – initiates a subscript, incrementing the super/subscript
level for the current fraction. {Termination by **$–** cancels the <u>last</u>
<u>subscript</u> (and any intermediate superscripts) for the current fraction,
decrementing the current fraction super/subscript level for each.}

A –
selects the alternate type (i.e. colour/intensity). {Termination by **$A**
 selects the default colour/intensity combination.}

H – selects

heavy type. {Termination by **$H** selects the default intensity type.}

M – selects mono spaced text.

P – selects proportionally
spaced text.

= – turns underlining on in the current super/subscript
level of the current fraction. {Termination by **$=** turns underlining
off in the current super/subscript level of the current fraction.}

I –
turns italicisation on in the current super/subscript level of the current
fraction. {Termination by **$I** turns italicisation off in the current
super/subscript level of the current fraction.}

, – initiates a
fraction numerator. The fraction level is incremented and the numerator
commenced at that level. {Termination by **$,** terminates the numerator
<u>and</u> commences the denominator [this is <u>equivalent</u> to ***.**].}

. –
initiates a fraction denominator for the current fraction level; a numerator
must be active at this level. {Termination by **$.** terminates the denominator
<u>and</u> the fraction; the fraction level being decremented. See also **$Q**
 below.}

/<res>/<font>/   – loads the font designated by the unsigned
integer value <font> into the resident set <res>, which may be any of 1 2
3 S or M, denoting resident alphabets 1 to 3, special symbol set or marker
set respectively. The terminating slash is mandatory. E.g. the sequence
/3/2015/ would load the complex Cyrillic font (font 2015) into memory
resident alphabet 3. Note: the active alphabet is unchanged, and in the
example given complex Cyrillic would not become the current alphabet unless
resident alphabet 3 was currently active.

   The following characteristics have a limited range of definition; that
is, they specify an action that is to take place immediately and then revert
to the CTC that was previously active (e.g. the escape to output a special
symbol will cause output of a single special symbol and then revert to the
previously active alphabet).

:<num> – the special symbol having value <num> will be output. The
value, denoted by <num>, should be a two digit unsigned integer in the range
[1,96]. (A single digit is permissible if unambiguous in context – i.e.
it is not followed by a digit. A single digit value may be unambiguously
given as two digits, the first being a zero or space.)

::<num> –
the marker having value <num> will be output. The value, denoted by <num>,
should be a two digit unsigned integer in the range [1,48]. (A single digit
is permissible if unambiguous in context – i.e. it is not followed by a digit.
A single digit value may be unambiguously given as two digits, the first
being a zero or space.)

V<num> – the character in the currently
active alphabet having the value <num> will be output. The value, denoted

by <num>, should be a two digit unsigned integer in the range [1,96]. (A
single digit is permissible if unambiguous in context – i.e. it is not followed
by a digit. A single digit value may be unambiguously given as two digits,
the first being a zero or space.)

B – the next output character
will be backspaced over the previously output character in this string.
The backspacing will be performed so that the medians of the characters are
coincident.

N – a "null" character is output. This is essentially
used to reset the reference point for subsequent super/sub-scripts, in that
a super/sub-script is normally placed relative to the current text level
so a string may have both super and sub script suffices aligned with continuation
text clearing both script levels. The null is necessary if, for example,
a prefixed subscript is to follow a super (or sub) scripted entity.

O –
subsequent super- and sub- scripts will be taken to clear the maximum and
minimum height extents of the previous character. This is particularly useful
for limits of integrals, summations etc. in mathematical expressions (when
a backspace <u>before</u> the first limit – i.e. "super" or "sub" script will
force the limits to be directly above/below the function symbol). For all
other purposes the "limit" script will be taken as a level of super/sub scripting.
{Termination by **$O**   will cause reversion to normal super/sub script
notation.}

D – diacritics (i.e. accents) will be produced. Each
DIMFILM font has an associated set of accents that may be applied to any
character in that font (for convenience, a cross out character may be included
as an accent). There may be up to twenty-four diacritics associated with
any font, and these are referred to by the upper case letters A to X. Single,
or multiple, diacritics may be associated with any output character. The
character to which the diacritic(s) is to be applied must immediately follow
the D escape. If a single diacritic is to be produced then its associated
reference (A to X) should itself follow the base character. Where more than
one diacritic is to be applied the associated reference letters should be
enclosed in parentheses following the base character. For example, to output
a with an acute accent (reference B) the sequence ***DaB** would be used.
If the user wished to output a crossed out (reference X) c cedilla (reference
A) the sequence ***Dc(AX)** would be used.

(– this (and the next)
escape is specifically to enable access of any character in the DIMFILM font
layout from the restricted set of characters of standard Fortran 77. The
next output character will be derived from the <u>next</u> non-escape input character
by <u>first</u> mapping it <u>left</u> one column in the font layout (i.e. its ASCII
value will be decreased by 16). {<u>Note:</u> this mapping is applied <u>before</u>
any active case mapping (either upper or lower) is applied.} For example,
to output < (less than) the sequence ***(L** could be used.

) –
this (and the previous) escape is specifically to enable access of any character
in the DIMFILM font layout from the restricted set of characters of standard
Fortran 77. The next output character will be derived from the <u>next</u> non-escape
input character by <u>first</u> mapping it <u>right</u> one column in the font layout
(i.e. its ASCII value will be increased by 16). {<u>Note:</u> this mapping is
applied <u>before</u> any active case mapping (either upper or lower) is applied.}
For example, to output \ (reverse slant) the sequence ***)L** could be
used; however, if <u>lower</u> mapping were active the output character which
would result is | (vertical line).

The remaining escapes are different in that they <u>terminate</u> effects only. That is, they <u>must</u> be preceded by **$**.

0 – (zero) is used as **$0** to clear all super/sub-scripts to the base level in the <u>current</u> fraction. (Termination of a fraction also has this effect.)

Q – used as **$0** to quit all super/sub-scripts <u>and</u> fractions to the base text level. (Termination of a character string input also has this effect.)

The concepts embodied in the above discussion will become clear if the following examples are considered.

```
      CALL SYMTXT('This is DIMFILM')
              or
      CALL SYMTXT('T*LHIS IS$L DIMFILM')
              or
      CALL SYMTXT('*Ut$Uhis is *Udimfilm')

      will all plot:  This is DIMFILM

      CALL SYMTXT('This is *=underlined *Iitalic$I$= text')

      will plot: This is underlined italic text
          (with the word italic italicised)

      CALL SYMTXT('z = *,x*+2$+ + 1*.*,1*.2y$.$.')

      will plot the equation z equals x-squared plus one
      all divided by one upon two y.
```

Certain parameters associated with text layouts may be specified by the user. For example, size and positioning of super and sub scripts may be set to "personalise" the output appearance.

SETCEX(CEXP)   – will set the character expansion factor to CEXP (CEXP > 0.0). Each character will be stretched/squeezed in its width relative to the nominal height. {Default = **1.0**}

SETCSP(CSPACE)   – each character will be followed by a space of CSPACE times the current text height. A negative value may be used to close up text (caution should be exercised). {Default = **0.0**}

SETFR(FHT,FOFF,FGAP)   – controls the appearance of fractions.

Heights of each numerator and denominator will be set to FHT times current text height, while the fraction bar will be set at a position FOFF times current text height above the present text baseline. A gap will be forced between the fraction bar and both numerator base and maximum height extent of denominator; this gap will be FGAP times the fraction numerator/ denominator height. (FHT,FGAP > 0.0) {Default = **(0.6,3./7.,0.13)**}

SETSUP(SHT,SUPY)
– sets the height of superscripts as SHT (> 0.0) times the current text height at a position SUPY times the current text height above the height of the preceding character (either relative to the actual maximum height of that character or as an absolute position according to the current height – see SUPABS/SUPREL). A negative SUPY will position the superscript below the preceding character relative/absolute height. {Default = **(0.4,0.1) relative**}

SETSUB(SHT,SUBY) – sets the height of subscripts as SHT (>0.0) times the current text height at a position SUBY times the current text height above the text baseline at the current level. A negative SUBY will position the subscript below the text baseline. {Default = **(0.4,0.0)**}

SETOVR(OVHT,OVY,UNY) – when super/sub-scripts are set as over/under (via escape ***O**) their height will be OVHT times the current text height, while "over" items will have a gap of OVY times current text height between their base and the reference character, and "under" items will have a gap of UNY times the current text height between their nominal top and base of reference character. (OVHT > 0.0; OVY,UNY **not** < 0.0) {Default = **(0.4,0.1,0.1)**}

SETUN(UNGAP) – underlining will be performed at UNGAP times the current text height below the lowest extent of the underlined section of text (this will clear fractions, subscripts etc., but may not clear individual character descenders). (UNGAP **not** < 0.0) {Default = **0.1**}

SUPABS – superscripts
will be positioned relative to the absolute (i.e. currently active) height of the preceding character.

SUPREL – superscripts will be positioned
relative to the actual character height maximum of the preceding character (although it will be forced to clear any subscript of a single level applied to the same character). Thus superscripts will, for example, be differently positioned for "x" and "X". {Default = **relative**} TEMPORARY NOTE – superscripts are currently set as absolute at –0.4 (see SETSUP); the declared default state will be applied shortly.

The default values may cause fractions and superscripts to exceed the basic text extent (i.e. baseline to nominal height). Judicious redefinition of certain values can avoid this (provided over/under items are avoided); it is suggested that FHT=0.4, FGAP=0.1, FOFF=0.5, may help in this regard. However, DIMFILM is designed to accommodate variations in line spacing and the "length" text routines may be used in this respect. Full allowance for such text layout will be made by graph annotation routines, etc.

RESTXT – will reset **all** text string parameters to their default
state (**except** height and string position).

2.1.6 – String Continuation

 It is often the case that the user may require a text string to be output
positioned immediately following the last output string. This can be very
useful when values (output via INUM/RNUM) are required within a complex string.

SYMCON will result in all following direct text output to be positioned
as a continuation of the previously plotted string.Note that <u>all</u> subsequent
strings will be contiguous until this option is negated.

ENSYMC will
cause all following direct text output to commence at the current plot position.

  Each string, including those produced as positionally continuous through
the use of SYMCON, will (by default) commence with the set **I**nitial **T**ext
Characteristics. For blocks of text, whether or not these be positionally
continuous, it is often appropriate for each string to continue with the
Current **T**ext **C**haracteristics that applied at the conclusion
of the previous string.

TXTCON will result in all <u>subsequent</u> direct text output commencing
with the CTC that was active at the end of the previous string. I.e. each
string does not revert to the ITC at its commencement. Thus TXTCON should
be referenced <u>before</u> the first continuation string (and after the string
which is required to commence with the ITC). It should also be noted that
any reference to any of the ITC routines will force the <u>next</u> string reference
to commence with <u>all</u> ITCs active, after which subsequent strings will have
continuous CTCs. (<u>Any</u> reference to other DIMFILM modules, such as graphing
or contouring, will force the next direct string reference to commence with
the ITC – TEXTCON will, however, remain active for subsequent references.)

ENTXTC will cancel the effect of TXTCON; all strings will be plotted
commencing with the ITC.

  The foregoing is applicable to direct string references (i.e. SYMTXT,
RNUM, INUM etc.) – indirect string output via (for example) graph titling,
axis labelling is not affected.   It is reiterated that SYMCON and TXTCON
are independent, but may be simultaneously active if required.

2.1.7 – Text String Length

 It is often necessary, for layout purposes, for the actual output length
of a text string to be known (before it is physically plotted).

STRING(TEXT)   is a REAL FUNCTION, returning as its value the length
(in users world coordinate units) that would be occupied were the text string
in TEXT (of type CHARACTER) to be passed to SYMTXT and plotted at <u>unit</u>
height.

SRNUM(RNUMB,FMT)   is a REAL FUNCTION, returning as its value
the length (in users world coordinate units) that would be occupied were
the parameters RNUMB (of type REAL) and FMT (of type CHARACTER) to be passed
to RNUM and the resultant string plotted at <u>unit</u> height.

SINUM(INUMB,FMT)
  is a REAL FUNCTION, returning as its value the length (in users world
coordinate units) that would be occupied were the parameters INUMB (of type
INTEGER) and FMT (of type CHARACTER) to be passed to INUM and the resultant
string plotted at <u>unit</u> height.

  As an example, suppose the user requires to exactly fit a string into
17.5 of his units. He may pass the string to STRING and obtain the length
it would occupy at unit height and select a height accordingly:

```
    WIDTH = STRING('This occupies 17.5 units')
    CALL SYMHT(17.5/WIDTH)
    CALL SYMTXT('This occupies 17.5 units')
```

  There is also the case where the user's string may conclude with one or
more blanks (possibly when a variable of fixed length is used for several
strings input during program execution). In this case:

STRAIL(TEXT,NA,SLONG)   is a subroutine reference, returning in SLONG
(REAL) the length (in users world coordinate units) that would be occupied
were the text string in TEXT (of type CHARACTER) to have trailing blanks
removed and the remainder passed to SYMTXT and plotted at <u>unit</u> height.
The number of characters in TEXT, after trailing blanks are stripped, is
returned in NA (INTEGER).

  All the string length functions/subroutine take full account of embedded
escape sequences within the input strings.

Chapter 3 – Raster Plotting

 DIMFILM offers the user a wide range of raster plotting facilities. These
may be freely intermixed with all other plotting operations. However, not
all devices can support the (full) range of raster functions. Where possible
non-raster devices will perform <u>some</u> simulation of the raster commands,
but – depending on the type of device – this may be minimal. The Appendices
on device types and specific devices supported should be consulted for details
on how the raster operations will be performed. This section details the
raster operations as they would function on a full-feature device.   In calligraphic
(or vector) plotting, images are created through the drawing of line segments.
This is analogous to the draughtsman's pencil. Raster plotting produces
fully shaded (possibly multi-coloured) pictures, and is more akin to the
artist's paintbrush. Digital computers and electronic displays require an
ordered approach to picture formation, and so the raster image is composed
of a mesh of points. These are generally arranged at regular intervals on
a rectangular grid. The smallest point that may be displayed, and to which
a unique colour may be assigned, is termed a **pixel**. These points are
so arranged that under normal viewing conditions they cannot be discriminated
(within an area of constant colour); that is, they appear to touch and the
coloured area is visually continuous. On sophisticated devices the number
of pixels is large and they cannot be discriminated even across colour transitions
– the image then appears continuous and, given sufficient colour range, will
be capable of a high degree of realism. Not all applications of raster graphics
require realism. The shading of areas can greatly enhance graphic presentations
even when relatively few colours are used.   Different devices may have different
numbers of pixels in each orthogonal direction. The usual expression of
the raster capability is as 'resolution of H x V pixels' (in the Horizontal
and Vertical directions, sometimes extended by 'x B' to define number of
data, i.e. colour, bits per pixel). An alternative is to express the number
of pixels per linear unit of measure (or even per unit of area). The consequence
is that to drive a raster device the user (or software) must know the device
resolution. DIMFILM is a general purpose graphics library and, as such,
is independent (in user terms) of the graphics output device addressed.
In consequence, the user defines a raster image in terms of minimum elements
(known as **cells**) and the low level interfaces (the device drivers) convert
these into device pixels. There is no need, then, for the general user to
be concerned with the device raster specification; beyond the obvious necessity
of ensuring the chosen device can adequately represent the generated pictures
(in terms of raster resolution, colour representation, etc.). The sophisticated,
or more demanding, user will find that careful definition of his coordinate
space and his cell mapping will enable precise addressing of a particular
device's raster pixels (and still enabling an approximate depiction on secondary
devices – possibly lower resolution preview displays).   In practice, the
user defines a **cell map**, comprising his chosen raster (i.e. cell) resolutions
in both horizontal and vertical directions and a mapping of this rectangular
area onto his coordinate space. He may then assign a different colour to
each cell of this raster grid,thereby forming a **cell array** . DIMFILM
will map this **cell array** onto each device such that every <u>device pixel</u>
contained within any user cell will be assigned the colour of that cell.
Clearly, the general case will be that many device pixels are located within
a particular cell and so are coloured identically yielding a cell of the
designated colour.

This may be illustrated thus:

```
     .     .     .     .     .     .     .     .     .     .
          ------------------------------------
     .    | *   |*    *    *|   *  |   *  |  *   |.     .
          |     |     |     |      |      |      |
     .    | *   |*    *    *|   *  |   *  |  *   |.     .
          |     |     |     |      |      |      |
     .    --*----*----*----*----*----*----*----.     .
          |     |     |     |      |      |      |
     .    | *   |*    *    *|   *  |   *  |  *   |.     .
          |     |     |     |      |      |      |
     .    | *   |*    *    *|   *  |   *  |  *   |.     .
          ------------------------------------
     .    | *   |*    *    *|   *  |   *  |  *   |.     .
          |     |     |     |      |      |      |
     .    | *   |*    *    *|   *  |   *  |  *   |.     .
          |     |     |     |      |      |      |
     .    --*----*----*----*----*----*----*----.     .
          |     |     |     |      |      |      |
     .    | *   |*    *    *|   *  |   *  |  *   |.     .
          |     |     |     |      |      |      |
     .    | *   |*    *    *|   *  |   *  |  *   |.     .
          ------------------------------------
     .     .     .     .     .     .     .     .     .     .
```

Notes: i) the cells are shown boxed
      ii) device pixels are shown as . or *
     iii) pixels contained in cells = *
      iv) pixels outside cell map = .

The foregoing illustration gives some pictorial idea of how device pixels
may be mapped into cells (assignment of cell boundaries to interiors is subject
to an implementation convention). It is (generally) redundant for the user
cell resolution to exceed that of the chosen output device for then not every
cell can be represented (in this case each pixel will correspond to the one
cell containing it, and some cells will be omitted as they will not contain
any pixel). Direct device pixel addressing will be ensured when the cell
resolution equates to the pixel resolution when the coordinate space is mapped
onto the device. The user of raster plotting through DIMFILM must determine
the resolution at which he wishes to work (the cell resolution) and may
then communicate with DIMFILM routines in terms of these cells. This may
either be defined in terms of the whole **cell array** or by defining successive
single rows of this array – **cell scan**s. Considerations of computer
storage may well dictate the chosen mode, as may the algorithm used to generate
colour values for each cell.   Whichever mode is utilised, the user may assign
colour (or intensity) either by reference to look up tables or specific value
triplets. Where scans are utilised there is the option of providing data
either for each cell in the scan or for incorporating some compression in
the input.

3.1 – Cell Mapping

 As depicted above, a **cell map** consists of a rectangular grid of cells
which is mapped onto user coordinate space (i.e. world coordinates). The
definition of this cell mapping is common to all subsequent discussion of
cell-raster output. The user is required to define the **N**umber of **C**ells
in the **X**-direction (i.e. horizontally) and the **N**umber of **C**ells
in the **Y**-direction (i.e. vertically). This gives the cell resolution.
Also required is the mapping of this onto the user's coordinate space.
For this it is necessary to define the **X**-coordinate of the **START**
of the first cell scan and the **X**-coordinate of the **END** of that scan,
and also the **Y**-coordinate of the **START** of the first scan and the
Y-coordinate of the **END** of the last scan. Pictorially this may
be shown:

```
        XSTART                                    XEND
  YSTART ────────────────..────────────────
        │    │    │    │    │    │    │    │
        │  1 │  2 │  3 │    │    │NCX │
        │    │    │    │    │    │    │    │
        ────────────────..────────────────
        │    │    │    │    │    │    │    │
        │  2 │    │    │    │    │    │    │
        │    │    │    │    │    │    │    │
        ────────────────..────────────────
        │    │    │    │    │    │    │    │
        │  3 │    │    │    │    │    │    │
        │    │    │    │    │    │    │    │
        ────────────────..────────────────
        .    .    .    .    .    .    .    .
        .    .    .    .    .    .    .    .
        .    .    .    .    .    .    .    .
        ────────────────..────────────────
        │    │    │    │    │    │    │    │
        │    │    │    │    │    │    │    │
        │    │    │    │    │    │    │    │
        ────────────────..────────────────
        │    │    │    │    │    │    │    │
        │    │    │    │    │    │    │    │
        │    │    │    │    │    │    │    │
        ────────────────..────────────────
        │    │    │    │    │    │    │    │
        │    │    │    │    │    │    │    │
        │NCY │    │    │    │    │    │    │
    YEND ────────────────..────────────────
```

This shows the layout of the cell mapped area. Thus each scan comprises
NCX cells, and **NCY** rows of raster cell scans map onto the area
defined by **XSTART** to **XEND** in the X-direction (i.e. the 'horizontal'
scan direction) and **YSTART**   to **YEND** orthogonally in the Y-direction
(i.e. the vertical, or direction in which successive scans are produced).
It is this area that is mapped from the users world coordinates onto each
device display area. This concept is important to all following discussion
in this chapter.

3.2 – Cell Arrays

 In the simplest form, the user may pass data relating to the whole of the
cell mapped area in a two-dimensional data array. In this case the user
will have determined the cell values (i.e. colours) for each cell in the
NCX x NCY cell grid. In general terms this will be in an array of dimension
(NCX,NCY). That is, the value in element (i,j) of this array will be that
of the i-th cell in the j-th scan. With cell arrays, the user defines the
cell mapped area <u>and</u> passes data for the <u>complete</u> grid of cells.

 Pictorially, this may be shown:

```
    (1,1)     (2,1)  . . . . . (NCX-1,1)     (NCX,1)

    (1,2)     (2,2)  . . . . . (NCX-1,2)     (NCX,2)

    (1,3)     (2,3)  . . . . . (NCX-1,3)     (NCX,3)

     .         .                  .            .
     .         .                  .            .
     .         .                  .            .
     .         .                  .            .

    (1,NCY-1) (2,NCY-1)  . . . (NCX-1,NCY-1) (NCX,NCY)

    (1,NCY)   (2,NCY)  . . . . (NCX-1,NCY)   (NCX,NCY)
```

 The user may specify these values as either look-up-table references or
actual value triplets (in the current colour model).

3.2.1 – Cell Arrays via LUT

 One of the simplest ways to generate a raster is through the definition
of a set of **CE**ll **L**ook up table references in an **AR**ray. In this
case the cell mapped area must be defined and an array of the LUT values
provided.

CELAR(XSTART,XEND,YSTART,YEND,NCX,NCY,LUTARR)   will result in the cell
array defined in LUTARR being plotted on the cell mapped area using look
up table entries.

   XSTART is the X-coordinate of the start of the first
scan

   XEND is the X-coordinate of the end of the first scan

   YSTART is
the Y-coordinate of the start of the first scan

   YEND is the Y-coordinate
of the last scan

 (all coordinates being in the user's world coordinate
space)

   NCX is the number of cells in each scan

   NCY is the
number of scans

   LUTARR is an integer array of dimension (NCX,NCY)
with the value of the (i,j)th element being a pointer into the look up table
for the i-th cell in the j-th scan. The values in this array should be
constrained to [0,255].

   There may be occasions when the storage array used for the LUT pointers
may not be dimensioned as required (for instance, when only as subset of
the storage array is to be plotted). There is, for such purposes, an alternative
subroutine CELAR1, which requires additional information concerning the array
and subarray dimensions. (The need and use of this will be apparent only
to those familiar with Fortran array usage.)

CELAR1(XSTART,XEND,YSTART,YEND,NAX,NAY,NCS,NRS,NCX,NCY,LUTARR)   will
result in the cell array defined in LUTARR being plotted on the cell mapped
area using look up table entries. The parameters XSTART, XEND, YSTART, YEND,
NCX, and NCY are as for CELAR, above. However, the dimensioning of LUTARR
is different, and defined by other parameters:

   XSTART is the X-coordinate
of the start of the first scan

   XEND is the X-coordinate of the end
of the first scan

   YSTART is the Y-coordinate of the start of the first
scan

       YEND is the Y-coordinate of the last scan

 (all coordinates
being in the user's world coordinate space)

    NAX is the first dimension
of storage array LUTARR

    NAY is the second dimension of storage array
LUTARR

    NCS is the starting column within LUTARR at which the cell
array commences

    NRS is the starting row within LUTARR at which the
cell array commences

    NCX is the number of cells in each scan

    NCY is
the number of scans

    LUTARR is an integer array of dimensions (NAX,NAY)
which holds the pointers to the look up table for the NCX x NCY cell array
as a subset. Each scan commences at column NCS of LUTARR and the first
scan is at row NRS of LUTARR. {For this case, the value of the (i,j)th element
of the cell mapped array to be plotted will be found at (NCS-1+i,NRS-1+j)
of LUTARR. CELAR is the specific case of CELAR1 with NAX = NCX, NAY = NCY,
and NCS = NRS = 1.}

 This may be shown schematically as:

```
             1    NCS                          NAX
            ------------..-------------..-------
            |    |    |   |     |    |    |    |
        1 | |    |    |   |     |    |    |    |
            |    |    |   |     |    |    |    |
            ------------..-------------..-------
            |    |    |   |     |    |    |    |
            |    |    |   |     |    |    |    |
            |    | 1 |   | NCX|    |    |    |
            ----=======..=======------..-------
            |    I    |   |    I    |    |    |
        NRS |   1I ** |   |    I    |    |    |
            |    I    |   |    I    |    |    |
            -----I------..------I------..-------
            .    .    .   .    .    .    .    .
            .    .    .   .    .    .    .    .
            .    .    .   .    .    .    .    .
            ------------..-------------..-------
            |    I    |   |    I    |    |    |
            | NCYI    |   |    I    |    |    |
            |    I    |   |    I    |    |    |
```

```
       -----=======..=======------..-------
       |    |    |    |    |    |    |    |
       |    |    |    |    |    |    |    |
       |    |    |    |    |    |    |    |
       -----------..------------..-------
       .    .    .    .    .    .    .    .
       .    .    .    .    .    .    .    .
       .    .    .    .    .    .    .    .
       -----------..------------..-------
       |    |    |    |    |    |    |    |
  NAY  |    |    |    |    |    |    |    |
       |    |    |    |    |    |    |    |
       -----------..------------..-------
```

 The cell mapped array is the subarray of dimension NCXxNCY, with its first
element (**) located at (NCS,NRS) within the main storage array of size (NAX,NAY).

3.2.2 – Cell Arrays via Triplet

 Not all usage will be satisfied by look up table references. Where more
control (and possibly greater colour range) is needed it is necessary to
generate **CE**ll **T**riplet values by **AR**ray. The cell mapped area must
be defined, and three arrays must be provided which together define the three
colour values for each cell. The meaning of the three arrays depends on
which colour model is currently active (e.g. RGB, HLS, HSV, etc.).

CETAR(XSTART,XEND,YSTART,YEND,NCX,NCY,V1,V2,V3)   will result in the
cell array defined by (V1,V2,V3) being plotted on the cell mapped area using
direct triplet values interpreted according to the currently active colour
model.

   XSTART is the X-coordinate of the start of the first scan

   XEND is the X-coordinate of the end of the first scan

   YSTART is
the Y-coordinate of the start of the first scan

   YEND is the Y-coordinate
of the last scan

 (all coordinates being in the user's world coordinate
space)

   NCX is the number of cells in each scan

   NCY is the
number of scans

```
   V1,V2,V3 are real arrays, each of dimension (NCX,NCY),
defining the colour triplet value for each cell;  the i-th cell in the j-th
scan having triplet value  {V1(i,j),V2(i,j),V3(i,j)}.  The values must be
within the ranges  required for the currently active colour model which will
be used  for their interpretation.
```

3.3 – Cell Scans

```
 The alternative to the cell array method of plotting a  raster/cell image
is for each scan line of cells to be  passed separately to DIMFILM.  This
may well be the preferred  method when large numbers of cells and scans are
involved  (especially when computer memory is limited).   In this case the
user must first define his cell mapped area  (there is no default).  Each
scan line is then passed to DIMFILM,  which will interpret successive scan
lines as part of a cell  mapped image and ensure correct location on the cell
mapped  area. The first line will be positioned at (XSTART,YSTART)  and subsequent
lines advanced automatically toward YEND.   A pointer onto the cell mapped
area is maintained by  DIMFILM to position the next scan.  This will be  initialised
to the cell mapped area X/Y-START at any DIMFILM  frame advance, whenever
the cell mapped area is redefined,  and when a new coordinate transformation
becomes active, or  when specifically initialised (via RESCMP).  The pointer
will  not be initialised if the actual number of scan lines  generated exceeds
the predefined cell map limit (in this case,  excess lines are ignored).
It is permissible for other  DIMFILM functions to be performed between consecutive
cell  scans; the raster image should appear continuous in the  cell mapped
area (although some devices may show some  minor positional discontinuity
or colour shift).  Equally, it  is perfectly acceptable for different scans
to be generated by  different cell scan techniques (e.g. run length compressed
scans may be interspersed with full lines).
```

3.3.1 – Cell Mapped Area

```
 Prior to any cell scan reference a valid cell mapped area must  be defined.
Once defined, within a DIMFILM execution,  this cell mapped area will remain
defined for all subsequent  cell scans.  It is independent of the cell mapped
area as  specifically defined via any cell array reference, and so the  various
types of cell scan operation may be freely intermixed.
```

CELMAP(XSTART,XEND,YSTART,YEND,NCX,NCY)   will define a cell mapped
```
area for use with all subsequent  cell scan operations and initialise the
cell scan pointer  to (XSTART,YSTART).

   XSTART will be the X-coordinate
of the start of  the first scan
```

```
   XEND will be the X-coordinate of the
end of  the first scan

   YSTART will be the Y-coordinate of the start
of the first scan

   YEND will be the Y-coordinate of the last scan
```

 (all coordinates being in the user's world coordinate space)

```
   NCX will
be the number of cells in each (full) scan

   NCY will be the number
of scans for the complete image
```

  It is of note that individual scan lines will be processed up to a <u>maximum</u>
of NCX cells, while short scans are permitted.

RESCMP may be used to re-initialise the cell scan line pointer into
the mapped area (to X/Y-START). When this is referenced the following cell
scans will overlay any previously generated ones on the same mapped area
on the current frame. It is included for those cases where such overlays
are required. The user is cautioned that the effects on the specific device
in use should be considered with care. (An implicit RESCMP is performed
on any DIMFILM frame advance.)

3.3.2 – Cell Scans via LUT

 Two routines are provided for **CE**ll **L**ook up table reference **SC**ans,
with the option of **R**un length encoding (compression). The routines should
be referenced once for each scan line required on the cell mapped area.

CELSC(LSCAN,NC)   will cause the NC look up table values in integer
array LSCAN (dimension NC) to be applied as the values for NC cells in the
next cell scan. These values should be in the range [0,255].

CELSCR(LRSCAN,NC)
  is for run length encoded scans. The integer array LRSCAN with dimension
(2,NC) holds the definition for the next cell scan. For each pair of elements,
LRSCAN(1,i) is the number of consecutive cells to have look up table value
LRSCAN(2,i); i=1,NC. Thus the length of the scan defined by this reference
is the summation of LRSCAN(1,i) for i=1 to NC. (Negative or zero runs are
ignored.) Each LUT value should be in the range [0,255].

3.3.3 – Cell Scans via Triplet

 As with cell arrays, it is also possible for the user to produce **CE**lls
of **T**riplet values in **SC**an lines; again these may optionally include
Run length compression. Each triplet will be evaluated according to
the current colour model. The routines should be referenced once for each
scan line required on the cell mapped area.

CETSC(V1,V2,V3,NC)   will cause the NC triplet values defined by real
arrays V1,V2,V3 (each of dimension NC) to be applied as the values for NC
cells in the next cell scan line. The value of the i–th cell will be evaluated
as the triplet {V1(i),V2(i),V3(i)} according to the current colour model.
(The triplet values should be within the relevant range.)

CETSCR(V1,V2,V3,IRUN,NC)
  is for run length encoded scans. The integer array IRUN (of dimension
NC) holds the run length corresponding to each of the NC triplets defined
by the real arrays V1,V2,V3 (each of dimension NC). The i–th triplet {V1(i),V2(i),V3(i)}
will be evaluated according to the current colour model and applied to IRUN(i)
consecutive cells in the scan. Thus the length of the scan defined by this
reference is the summation of IRUN(1,i) for i=1 to NC. (Negative or zero
runs ignored.) Each value triplet should be within the range appropriate
to the current colour model.

Part 3 – GRAPHING

Chapter 1 – Basic Graphing

1.1 – Introduction

 DIMFILM includes a collection of subprograms to facilitate the production
of graph plots, while retaining the advantages of the overall system. The
graphing routines may be freely intermixed with the general plotting routines.
The system is designed to permit the user considerable freedom, and although
this must inevitably introduce some pitfalls for the unwary, a large amount
of checking is performed, and warning messages are included. These are controlled
as are those in the general plotting routines, through the check level.
Users interested primarily in graphing output should first familiarise themselves
with the earlier sections on basic plotting.   Graphing may be freely intermixed
with calls to all other plot routines. DIMFILM attempts to ensure that on
exit from any graphing routine, all internal parameters are set to their
values on entry. The beam is also positioned so that general plotting may
be resumed easily. Such properties as intensity, orientation and height of
symbol strings, broken lines etc., will be maintained. It is possible therefore,
to interrupt a diagrammatic plot to perform graphical plotting and then to
resume. A few limitations are, however, necessary. Primarily, on return,
a broken line pattern will commence at the beginning of its sequence (the
consequence of a beam off movement). Symbol strings will commence with <u>all</u>
their initial features active. If TXTCON has previously been called, it will
be operable on all but the first subsequent calls to SYMTXT.   Except where
otherwise indicated, all graphical drawing routines are performed as CSType
1 operations. All graph textual annotation will be treated as CSType 2, while
non-textual annotation is a CSType 3 operation. All graph titling and labelling
(user supplied) will be   subject to any escape sequences as described in Section
3. Graphs will also be subject to interruption by markers, when the routine
INTRPT has been called. This applies only to plots where a set of data values
has been supplied; when interpolation is used the interrupt only applies
to the supplied data points ie. a call of INTRPT(N,CHAR) would cause each
Nth data point to be marked with character CHAR. The height used would be
that currently used outside the graph plot.   The use of intermixed general
and graph plotting will be illustrated in a later section. It is emphasised
that this mixing may be performed freely within a frame, and at the simplest
level it may mean only the changing of intensity for axes etc.

1.2 – Graph Location

 The overall size of a graph is taken to be identical to the current pre-clipping
rectangle. That is, the user by making a call to PANE (or through the default
pre-clipping rectangle being the same as the overall bounds) has set the
area to be occupied by the graph. This area includes a margin for captions
and values, and the proportions used are preset in DIMFILM. Of specific
interest to the user, is the area occupied by the actual graph plot, excluding
margins etc. This is 83% of the pre-clipping rectangle in each direction.
The graph labels and values will be in the margins between this rectangle
and the actual graph plot. Each margin in the X(or Y)-direction will be
8.5% of the corresponding X(or Y)-dimension of the rectangle.   This enables
the user to position the graph plot anywhere within the bounds and graphs
in equal area will appear identical, whilst those of similar areas would
be reduced or enlarged accordingly. This ensures continuity over a large
number of plots and assists the user in making more than one plot in an area,
but changing axis scales as required. The actual co-ordinates used in calls
to PANE etc. are not referred to by the user during the course of graphical
plotting All scaling of graphs is accomplished internally. The X-axis of
a graph will always by parallel to the X-direction of the frame and the
Y-axis parallel to the Y-direction. Hence, the orientation of a plot is determined
by the current mode of the device.

1.3 – Axis determination

 In any graphical work the user is obliged to determine the axis types required
before commencing a plot. There are few conditions imposed on the user performing
graph plotting with DIMFILM, and these are summarised below. The axis type
need only be declared if a change is to be made, and clearly this must be
done before plotting or assigning values to the axis etc. is performed. By
default, it is assumed that both X-and Y- axes are linear. A number of routines
are available for changing or resetting axis types, at present the user
will only be concerned with:

LINX to set a linear X-axis.

LINY to set a linear Y-axis.

LINXY to
set linear X- and Y-axis.

LOGX to set a logarithmic (base 10) X-axis.

LOGY to set a logarithmic (base 10) Y-axis.

LOGXY to set logarithmic
(base 10) X- and Y-axis.

 Any combination of these calls is permissible, the last declared type for
each axis being accepted.   For a basic graph, the user will have a set of
(X,Y) co-ordinates that are to be plotted. It may be that the plot is to

be made along axes that have ranges automatically determined by the X– and
Y– ranges of the data, or alternatively along axes having predefined ranges.
The default is for automatic range determination along both axes to be performed
as in the former requirement. For the latter requirement, the user may supply
a range along either or both axes by the following.

XAXIS(XL,XR)

YAXIS(YL,YR)   where XL is the left bound and XR
the right bound for the X–axis and YL is the low bound and YU the upper bound
for the Y–axis. In each case the bounds are given in terms of the user's
graph units.

 Should both bounds be equal in a call to one of these routines, then DIMFILM
will go into auto mode for that axis.   Axis ranges set in this way remain
set until such time as they are either set to some other value or auto mode
is restored using the following routines:

AUTOX to set auto range determination on the X–axis.

AUTOY to set
auto range determination on the Y–axis.

AUTOXY to set auto range determination
on both axes.

1.4 – Drawing the Graph

 Having determined axis types and ranges as required, a single call to the
routine

GRAPH(X,Y,N)   will accomplish the plotting. X and Y are REAL arrays
(of dimension at least N) holding the N X– and Y– coordinates respectively.
This will join consecutive co-ordinates by straight line segments. The co-ordinates
should be supplied monotonically in X. N must be at least 2.

 To complete the graph with labels, values and to frame it, the following
routine may be called:

GRDEF(TITLE,XLAB,YLAB,IXY)   where TITLE,XLAB and YLAB are all <u>strings</u>
of type CHARACTER with an appropriate length:

   TITLE is the title
to appear at the base of the graph.

XLAB is the label to appear below
the X edge of the plot.

YLAB is the label to appear to the left of
the Y edge of the plot.

IXY is used to control the presence or absence
of axes, according to its value – .NL =0 :– neither axis will appear,
.NL =1 :– the X-axis only will appear and be ticked, .NL =2 :– the Y-axis
only will appear and be ticked, .NL =any other value :– both axes will appear
and be ticked.

  In addition to these titles, labels and axes, the plot will be framed
(i.e. the actual graph will be boxed), all edges will be ticked and values
will appear adjacent to the ticks.   (It should be noted that any axis can
only be plotted if it lies within the range of the orthogonal co-ordinate.
For example the X-axis at y=0 can only be plotted if zero is included in
the Y-range.)   Essentially, this routine is present to offer a default set
for the annotation of the plot. It is not expected that all users will wish
for this set of annotations and consequently a large number of routines are
included to perform the individual parts. It is suggested that each user
constructs his own standard set of calls in a single subroutine and references
this at the end of each plot. Of course, he may wish to annotate each plot
differently and in this case, the full range of routines may be accessed.
These routines will now be described.   A simple convention is used to refer
to the sides of any line. For lines in the X-direction there is a lower (L)
and upper (U) side, referring to the normal Y-values which are orthogonal
to this line.   Similarly, for a line in the Y-direction there is a left (L)
and right (R) side. The graph itself may, or may not have the axes included
in it and is bounded by four edges, surrounding which are the margins; the
whole being contained exactly within the pane.   For all the following routines
that perform titling or labelling, the symbol strings are passed through
the SYMTXT routine and all the facilities described in section 3.1 are available
(although it should be noted that text continuation will not operate, each
string commencing with the selected features). Text strings will be appropriately
centred.   Titles may be placed below the lower edge or above the upper edge
by the following.

LTITLE(TITLE)   for titling below the lower edge

UTITLE(TITLE)
 for titling above the upper edge

 TITLE is a <u>string</u> of type CHARACTER of the appropriate length, which
contains the text to be passed to SYMTXT.   Labels may be placed against any
edge or either axis, although for the latter option it is necessary for the
axis to be within the orthogonal axis range.

LXLAB(LABEL)   labels the lower edge

UXLAB(LABEL)   labels theupper edge

LYLAB(LABEL)   labels the left edge

RYLAB(LABEL)   labels the right edge

LXALAB(LABEL)   labels below the X-axis

UXALAB(LABEL)   labels above the X-axis

LYALAB(LABEL)   labels to the left of the
Y-axis

RYALAB(LABEL)   labels to the right of the Y-axis

  In each of these calls LABEL is a <u>string</u> of type CHARACTER of the appropriate
length, which contains the text to be used as the label.   The edges and
axes may be ticked as required, the ticks corresponding to steps in the range
in the appropriate direction and taking account of the axis type. The spacing
of ticks etc., will be considered below.

LXTIK will cause ticking of the lower edge

UXTIK will cause ticking
of the upper edge

LUXTIK will cause ticking of both the lower and upper
edges

LYTIK will cause ticking of the left edge

RYTIK will cause
ticking of the right edge

LRYTIK will cause ticking of both the left
and right edges

  In each of these cases the ticking would be directed inwards.

LXATIK will cause ticking on the lower side of the X-axis

UXATIK will
cause ticking on the upper side of the X-axis

LYATIK will cause ticking
on the left of the Y-axis

RYATIK will cause ticking on the right of the
Y-axis

  Values corresponding to the ticking of edges and axes may be plotted.

LXVAL plots values below the lower edge

UXVAL plots values above
the upper edge

LUXVAL plots values below the lower edge and above the
upper edge

LYVAL plots values to the left of the left edge

RYVAL plots
values to the right of the right edge

LRYVAL plots values to the left
of the left edge and to the right of the right edge

  Similarly for the axes,

LXAVAL plots values below the X-axis

UXAVAL plots values above the
X-axis

LYAVAL plots values to the left of the Y-axis

RYAVAL plots
values to the right of the Y-axis

  The axes themselves may be drawn by the following

DRAWXA to plot the X-axis

DRAWYA to plot the Y-axis

 As an alternative to ticking the edges etc., the user may grid the plot
entirely using the following

XGRID will plot grid lines normal to the X-axis (i.e. parallel to the
Y-axis)

YGRID will plot grid lines normal to the Y-axis

XYGRID will
completely grid the plot in both directions

 It is clear that when ticking (or gridding) and valuing are applied, or
labelling an axis, it is essential that the ranges for each axis are defined
first. They may either be explicitly defined by calls to XAXIS and YAXIS,
or if auto mode is in operation on either axis,they will be known internally
<u>after</u> the plotting of the graph (i.e. as described in this section through
a call to GRAPH). Clearly then a sequence for these calls has been formulated.
This may be summarised in priority, after the definition of the pre-clipping
area:

  1) axis type declarations

  2) axis range specifications

  3)
ticking, valuing, gridding or labelling axes

   Titling and edge labelling may occur at any point. If automatic mode is
in operation on either axis then the graph call which fulfills the function
of range determination must occur at step 2, otherwise it may appear anywhere
after stage 1. Steps 1 and 2 will be retained from any previous graph plot
and so do not need redefinition for subsequent plots. Of course, if auto
mode is active then the <u>previously</u> determined range(s) will hold until
reset by a call to a graph plot routine.   Although the user has considerable
freedom, it is suggested that the safest course is :

  1) axis type declarations (if changes to be made)

  2) axis range
specifications (if non–auto, or to be reset)

  3) graph plot

  4)
all annotations (in any order)

   The edges of the graph area may be plotted by

GRFRAM

 This will be required in nearly all cases. The frame is drawn as CSType
3.   To outline the current pane, call

OPANE (as described in section 2 – CSType 3.)

 The determination of value intervals (also used for ticking and gridding)
is performed automatically (although see section 4.2 later) depending on
the axis ranges.   For linear axes with a range of length d x 10**i with
d greater or equal to 1.0 and less than 10.0 (i.e. d in [1.0,10.0)) an interval
between values will be chosen according to the value of d as shown below:

  d in [1.0,3.0)   interval =   2.0 x 10**(i–1)

  d in [3.0,7.0)

```
 interval =   5.0 x 10**(i-1)

  d in [7.0,10.0)       interval =   1.0 x 10**i
```

   The values marked will be integral multiples of the interval, thus ensuring
that 0.0 would be a value point. The choice of these intervals maintains
a reasonable number of divisions without overcrowding (always more than 4
and less than 16). The values themselves will be plotted in a format to maintain
the maximum accuracy without excessive digits appearing. Where possible
a common power of ten factor will be removed (and indicated adjacent to the
edge or axis being 'valued').   The same value points are used for the tick
spacing, the tick lengths being one hundredth of the size of the shorter
pane side.   For logarithmic axes, to ensure clarity, if there are less than
two cycles along the axis values of the <u>logarithm</u> (base 10) the values
will be plotted and determined as for the linear axes. If at least two cycles
are contained along the axes then the values as integer powers of 10 will
be appropriately placed at cycle boundaries.   Ticking will occur at these
boundary points, and if fewer than five cycles are contained on the axis,
then intermediate half-size, faint ticks will be plotted (these correspond
to the normal intermediate divisions of a log cycle).   When gridding is
specified the spacing is the same as the corresponding ticking as described
above.

Chapter 2 – Refined Graphing

The features described in the previous section enable the user to plot straight
line graphs from a set of co-ordinate pairs, against linear or logarithmic
axes. The immediate extension of this is for some form of interpolation
to be performed between consecutive data points to give a smooth curve for
the finished plot. DIMFILM provides the choice of two polynomial interpolations
– third or fifth degree. To effect either of these the call to GRAPH should
be replaced by

POLY3(X,Y,N)   for third degree polynomial interpolation

POLY5(X,Y,N)
  for fifth degree polynomial interpolation

 X and Y are REAL arrays (of dimension at least N) holding the N pairs of
co-ordinates to be graphed. There is a proviso that N must exceed the degree
of the polynomial interpolation being performed. Plotting of an interpolated
graph will terminate should a singularity be found during the interpolation
process.   By default the number of steps between consecutive abscissae at
which interpolated values will be found is 5. The larger this value, the
smoother the resultant curve, although execution time for the program will

increase correspondingly. The number of steps for this interpolation may be altered by

INTERP(NSTEPS)   where NSTEPS is the number of intermediate steps at which the interpolation polynomial will be evaluated.

 An alternative output that might be required from a set of co-ordinate pairs might be a simple point plot, when no join is made between the points. To accomplish this,

PTPLOT(X,Y,N,NCHAR)   where X,Y and N are as before, and NCHAR is the number of the required marker symbol at each point (from the current Marker Font), NCHAR is an integer in the inclusive range 1 to 48 . As for the normal call to GRAPH, N must be at least 2.

  When discrete co-ordinate pairs are provided, it may be desired to produce a histogram bar chart. The bars will in each case be in the Y– direction, and correspond to each abscissa input. The appropriate instruction is

HISTGR(X,Y,N,BAR)   X,Y and N being as before.

 The fourth parameter, BAR, controls the type of histogram that will be produced.

BAR = –1. – auto-width histograms will be produced, so that adjacent bars touch.

    = 0. – histogram lines are produced (i.e. a special case of fixed width bars).

    > 0. – histogram bars width BAR (in user <u>axis</u> units) will be plotted.

   Several variations of histogram bar charts are possible.   A step histogram may be plotted by

STEPGR(X,Y,N)   the three arguments being the same as the first three arguments for HISTGR.

 The effect is to produce an auto-width histogram with common verticals

removed. This is the common staircase type of histogram. It is possible
to shade the histogram bars (when bars of a finite width are to be produced).
This is performed by

SHDEGR(X,Y,N,BAR,THETA,GAP)   The first four arguments are those for
HISTGR, with the limitation that for BAR = 0.0 histogram lines (i.e. bars
of zero width) only will be produced. The last two arguments are used to
define the type of shading, and correspond to the arguments for HATCH (section
*.x)

   THETA is the angle (AGroup1, default degrees – see I.1.9) the
shading lines make with the positive X-direction, and should be in the range
(–90.0,90.0)[degrees].

   GAP is the perpendicular distance, in user
overall units, between consecutive shading lines.

  The shading of each bar is performed separately, and is defined under
HATCH. The user may force two histogram bar charts, each with different
shading types, to be plotted with bars of the two data sets adjacent for
comparison by the appropriate positive specification for BAR, and a shift
in X range (by BAR) between the calls to SHDEGR for the two data sets.
With all the histograms, it is the case that a vertical line may coincide
with an edge of the graphical area. DIMFILM has no knowledge of whether
the user will draw the edges, and consequently the default is to plot such
verticals. However, when the Y edges are plotted, such verticals will be
superimposed, which may not be desired.

HREP causes the suppression of histogram verticals that coincide with
the Y-edges, while

HPER restores the default, whereby the plotting of
such verticals is permitted.

 It will often be the case that a graphical plot of a known function is
to be made. In this case the appropriate call would be

GRAPHF(FN)   where FN is a REAL FUNCTION, supplied by the user, and
which should be declared EXTERNAL in the calling program. The function should
have two REAL arguments and be referenced as   Y = FN(X,DX)   where the
value returned in the function is the function value (i.e. Y co-ordinate)
at X. The argument DX on return should be the increment in the X direction
to be made for the next plotted point. These values are all in terms of
the axis co-ordinate scales.

 GRAPHF starts at the minimum X-axis value, and proceeds until the maximum
is reached or , in the case of a curve which becomes double valued in X ,
the minimum is passed. That is, plotting commences until the graph leaves
the X-axis range. Alternatively, the graphing will cease when two consecutive
zero values for DX are returned. The general comments under Part I, section
4 relating to FNPLOT apply, with references to X and Y ranges being, in the
case of GRAPHF, references to the axis ranges.   With GRAPHF it is obvious

that automatic axis scaling cannot be applied; hence both axis ranges must
have previously been determined. If either axis is in auto mode at the initiation
of a call to GRAPHF, then the current axis range for that axis will hold,
ie. it will be the range used for the immediately preceding graph plot (possibly
determined under auto-mode). Should no previously defined range exist then
an error condition will occur and no plot will result.   All the above graph
plotting routines take full account of the various axis types.   A further
axis type is available, and is of particular use for commercial plots. This
is the month axis.

MONTHX(SMONTH)   will set the X-axis to months,

MONTHY(SMONTH)
 will set the Y-axis to months.

 In either case, the REAL argument SMONTH is used to indicate the starting
month along the appropriate axis. The months January to December running
from 1. to 12. respectively, a fractional value is permitted to denote the
plot is to commence at a fraction of a month. The left limit will be the
1st of this month.   E.g. February 7th would be 2.25 for SMONTH, and the axis
would commence at February 1st with the first plot occurring at February
7th.   The value of SMONTH must be in the range 1.0 to 13.0 . If an illegal
value is requested, 1.0 will be substituted.   A subsequent call to GRAPH,
POLY3, POLY5, PTPLOT, or HISTGR with arguments (X,Y,N) (and BAR in the case
of HISTGR, and NCHAR in the case of PTPLOT) will cause plotting with the
appropriate month axis from the starting point SMONTH at monthly intervals
for the N values along the orthogonal axis. This is the auto-mode month
axis, the argument N being the number of months to be as plotted. The orthogonal
axis values should be held in the REAL array X or Y as appropriate, the
other array being a scratch array that will be used during the call, and
its contents lost.   This situation will be clarified by example:

LOGY Y-axis set as a log axis

MONTHX(3.0) X-axis a month axis, commencing
at March 1st

GRAPH(X,Y,N)

 This will plot the N values from REAL array Y at monthly intervals starting
at March 1st, i.e. Y(1) will be the Y value plotted against March 1st, Y(2)
against April 1st, etc. REAL array X must be available for use by GRAPH.
The X axis will be divided into months.   Similarly,

LINX X-axis set as linear axis

MONTHY(3.0) Y-axis a month axis,
starting at March 1st

GRAPH(X,Y,N)

 will plot N values from REAL array X at monthly intervals commencing with
March 1st, i.e. X(1) will be the X-value plotted against March 1st, X(2)
against April 1st, etc... REAL array Y must be available for use by GRAPH.
The Y axis will be divided into months.   After a call of MONTHX or MONTHY
the relevant axis will remain a month axis until reset (for example, by a
reference to LINx or LOGx). It is not permissible for both axes to be of
type month for any plot.   Ticking and valuing will apply to month axes with
each month being ticked when 12 or fewer months are plotted along the axis.
For greater numbers of months not every month will be marked, as indicated:

Number of months	Month interval for ticking, etc.
≤12	1
(12,24]	2
(24,36]	4
(36,60]	6
>60	12

   The above methods have shown the use of automatic month plotting, when
each successive data point is exactly a month from the previous point.
It may be necessary to plot at intervals other than exact months, in which
case it is necessary to define the number of months that are to be plotted
along the month axis.

MONSX(NMNTHS)   sets the number of months along the X-axis

MONSY(NMNTHS)
  sets the number of months along the Y-axis

   NMNTHS is the number
(INTEGER) of months to be plotted along the appropriate axis. These calls
do not replace the calls to MONTHX or MONTHY, which should also be made.
Any subsequent call of AUTOX, AUTOY or AUTOXY will reset the relevant month
axis to auto mode.

 When this mode of fixing the month range is in operation the previously
scratch array in the X,Y arguments to GRAPH etc., is used differently. It
should now hold the offset, in months, from the beginning of the starting
month (as set by MONTHX etc). An example will illustrate this:

LOGY sets logarithmic Y-axis

MONTHX(2.5) sets X-axis to months,
beginning February 14th for auto mode month plots.

MONSX(10) indicates
fixed number (10) months along X-axis

GRAPH(X,Y,N)

 N points will be plotted, Y(I) will hold the Y-value of the Ith point (I=1,N)
and X(I) holds the offset, in months, from the 1st February (this being the
start month), although the axis will start at the 14th (as set).   Suppose
X(1)=1.25 , X(2)=1.50 , X(3)=2.00 , ...   Y(1) would be plotted as March
7th (1.25 months after Feb 1st) Y(2) would be plotted as March 14th (1.50
months after Feb 1st)   Y(3) would be plotted as April 1st (2.00 months
after Feb 1st) etc.   When discrete data is provided to a graphical routine
(excepting, therefore, GRAPHF) the user may wish symmetric or asymmetric
error bars to appear for each data point in the X and/or Y directions.
Separate calls exist for this purpose, and are listed below:

GRAPHE(X,Y,N,MX,MY)   which corresponds to GRAPH, producing straight
line joins between data points.

PTPLTE(X,Y,N,NCHAR,MX,MY)   corresponding
to PTPLOT, each data point being plotted with the marker designated by NCHAR

POLY3E(X,Y,N,MX,MY)

POLY5E(X,Y,N,MX,MY)   corresponding to POLY3,POLY5
respectively with polynomial interpolation between data points

 The arguments are of somewhat different form than the corresponding routines.
 X and Y are real arrays of dimension X(N,MX), Y(N,MY), where N is the number
of data points.   MX and MY may each have the value 1, 2 or 3.   Considering
the array Z(N,MZ); Z=X,Y, then in each case Z(I,1) holds the Ith Z data value
which will be plotted.   For MZ=1 no error bars will be produced in the Z-direction.
 MZ=2 implies symmetric error bars should appear in the Z-direction, and
for the Ith data point, the Z error bar will be +/- Z(I,2) in magnitude.
 MZ=3 implies asymmetric error bars in the Z direction, in which case the
Ith data point will have Z error bar given by +ABS(Z(I,2)), -ABS(Z(I,3)).
 The error bars are plotted with the same line type as the graph, although
the tick heads to each bar will be solid. It should be noted that MX need
not equal MY, and for error bars in one direction only the other MZ should
be set to 1. Tick heads will have the same size as other graphical ticks.
 To force the error bars to be of different line type it is necessary to
make a standard call to GRAPH etc. followed by the required line type change
and then a call to PTPLTE with NCHAR set to 0 (unless some character is required
at the data points additional to the error bars).   It is possible for axes
to be drawn at any particular X- or Y-value, i.e. at other than the normal
zero lines. To set this

AXCUT(XO,YO)   where (XO,YO) is the co-ordinate pair at axis intersection.

By default this will be at (0.,0.). The axis intersection point will then
define axis positions for all subsequent axis drawing, ticking labelling
etc., until reset.

DIVUNI(DIVX,DIVY)   (with DIVX,DIVY 0.0) will cause all subsequent linear
axis ticking etc., to be at intervals of DIV- in the axis units for the relevant
axis. If this value is found not to be sensible automatic ticking etc.,
will be performed. The condition for DIV- to be used is that it is less
than one-half and not less than one-hundredth of the relevant axis range.
It may be specified for only one axis by setting the other DIV- non-positive.

 To return to automatic division determination on both axes

AUTD

 There is also provision for controlling the intensity step for the intermediate
logarithmic axis ticks.

LOGTIK(RSTP)   will result in a step down in intensity of RSTP for the
intermediate ticks. Note that RSTP is a real value in the range 0.0 to 1.0
. The default value is 0.2

 The user may also control the criteria used for determining whether log
cycles should be ticked. Section 1 (of Chapter 4) indicated that if less
than two cycles appeared along an axis, then values of the logarithm would
be plotted, while if five or more cycles were contained on an axis, no intermediate
ticks would appear. To change these values:

LOGLIM(ZMIN,ZMAX)   where ZMIN and ZMAX are the real values of these
criteria. ZMIN being the minimum number of cycles for cycle values etc.,
to appear, ZMAX the number of cycles at and above which intermediate ticking
will not appear.

 When automatic mode is selected, each graph has individually determined
axis ranges (on either or both as specified AUTOX, AUTOY etc.). It is often
the case that while automatic range determination may be used for the first
graph, it may be wished to plot subsequent graphs against the same axis ranges.
This may be forced by:

SAMEX which fixes the X-range as that currently held   e.g. after a call

to a graphical range that X range is retained (as though the relevant call
to XAXIS had been made) for all following graphs, and automatic X-range determination
is turned off.

SAMEY acting on the Y range

SAMEXY acting on both
X and Y ranges

 By default axis values are plotted in a real format (i.e. including the
decimal point). It is possible, when the range permits, for integer values
to be plotted by linear axes or edges.

IXAX specifies that future valuing of linear X-axes should, when possible,
be in integer format

IYAX acts similarly for subsequent linear Y-axes

 Real values may be restored for all subsequent valuing by:

RXAX for real values on X-axes

RYAX for real values on Y-axes

 Common scale factors may well be included for integer values axes where
necessary, and it should be appreciated that ranges such as (10.0,10.9) are
incapable of being valued integrally. Also, the user should note that the
X and Y values passed to GRAPH etc., should still be <u>real</u> numbers.

Chapter 3 – Alternative Co–ordinate Systems

 In addition to the rectangular co-ordinate systems described above, DIMFILM
provides the means to plot against polar axis systems, i.e. radial plots.
 The area occupied by a polar plot is the same portion of the window as
used for cartesian plots. In the case of polar plots, however, the total
area is not used, the scaling and positioning of the axis system being performed
to maximise the area used.   The calls are closely analogous to those previously
described for rectangular co-ordinate systems, and the priorities for ordering
of calling sequences as given in Section 1 still hold.   The radial axis
type may be specified by:

LINR to set linear radial axis (this is default)

LOGR to set logarithmic
radial axis

The units of angular measure for angles to the radial plotting routines belong to AGroup2 (therefore being radians by default). The section on Angular Measure within the part on Basic Plotting gives further information and describes ways of changing the default.   Axis ranges may be supplied by the user:

RRANGE(RMIN,RMAX)   will set the radial range to be RMIN to RMAX, where RMAX is greater than RMIN which is greater than or equal to 0.0

TRANGE(THETA1,THETA2)
  will set the angular range from THETA1 to THETA2, both angles being AGroup2 (default radians) and should each be in the range [0.0,2pi][radians].

 If either of these calls is made the appropriate axis range is fixed until reset. By default, automatic range determination is provided. This may be reselected at any time:

AUTOR sets automatic radial range determination

AUTOT sets automatic
angular range determination

AUTORT sets automatic range determination
for both axes

 Having determined axis types, ranges as required:

POLAR(R,THETA,N)   will plot the N co-ordinate pairs from R, THETA (both REAL arrays of minimum dimension N), where R holds the radial values, and THETA holds the angular values (AGroup2, default radians).

 The graph may be completed by

POLDEF(TITLE,RLAB)   where TITLE and RLAB are <u>CHARACTER</u> arrays (of appropriate length) holding captions to be passed to SYMTXT

   TITLE is
the title to appear at the base of the plot

   RLAB is the label to appear
below the plot

 The polar boundary will be included with radii at 90 degrees, each radius will be ticked and valued, and angular values will be plotted around the circumference.   Titling etc., may be performed through LTITLE, UTITLE, LXLAB, UXLAB, LYLAB, RYLAB as for rectangular plots. Due to the varied boundary

shapes and positions that may apply to radial plots it is not sensible to provide any rigid labelling format. Generally, the –XLABEL should be satisfactory for describing the radial axis, and –YLABEL for the angular values. Ticking and valuing may be independently controlled. In the following description it should be borne in mind that in the general case the extremities of a plot will not be coincident (only in the exceptional case of the angular range 2pi [=360 degrees] will this be so), and between the extremities (i.e. going from THETA1 to THETA2, the angular range limiters) one or more principal radii may be included (i.e. those at integral multiples of pi/2 radians, viz. 0, pi/2, pi, 3pi/2, 2pi [equivalent to 0., 90., 180., 270. degrees]).

POLOUT will produce the bounding outline of the polar plot area – i.e. the arcs of minimum and maximum radius, and the radii at the bounding angles will be plotted.

POLFR will produce the bounding outline (as POLOUT) and also plot the included principal radii.

RXVAL will plot the radial values at the extreme radii.

RVAL will plot the radial values at the extreme radii and at included principal radii.

RXTIK will tick the radial values on the extreme radii.

RTIK will tick the radial values on the extreme radii and on included principal radii.

 It should be noted that, generally, the tick length is taken into account when valuing radii to avoid any overlap of tick and value. The user may exercise some control over this by:

RADSEP(ANGLE)   ANGLE is specified as AGroup2 (default radians). It controls the appearance of ticks for principal radii. If a principal radius is closer than ABS(ANGLE) to an extreme it will not be ticked. For non–coincident radii of angular range extremities to be separately valued there is a minimum necessary angular separation. If ANGLE >0.0, this separation is ANGLE. For zero or negative ANGLE automatic computation taking account of value positioning and character size is made. At less than the (fixed or computed) separation angle only the radius of least angle will be valued. While for less than twice the separation angle, values will not be spaced by the tick size from the radius, i.e. values will be adjacent to radii. The default value of ANGLE is –0.1 (radians), i.e. automatic computation of separation angle is made, but included principal radii closer than 0.1 radians to an extrema are not ticked.

POLGRD will grid the polar plot.

 The gridding of a polar plot comprises the plotting of radial lines and circular arcs at certain radii. Due to the concentration of radii at the centre of a polar plot it is generally necessary to incorporate some thinning of radial lines as the centre is approached.   This is done by dividing the

maximum radius, and each division having a separate angular increment between
radial grid lines. The circular arcs that appear are dictated by the ticking
(and valuing) of the radii, consequently the division points must be adjusted
to coincide with these arcs. The number of divisions fixes the size of each
division, and the tick point along a radius following nearest to each division
is selected to bound a division. However, a spacing less than .4 of the
fixed division is ignored to prevent nonsensical thinning of radii.   By
default 3 divisions are used with angular separation of radii being pi/2,
pi/4, pi/12 (radians) [equivalent to 90, 45, 15 degrees].   These may be changed
by the user by

TLEVS(XLEV,N)   where N, INTEGER, is the number of levels to be used
for the thinning of gridding radii – a maximum of 4 is permitted $1 \leq N \leq 4$
and XLEV is a REAL array of dimension 4, with the first N elements holding
the angular intervals between radii at each division (from 1 at the centre
to N at the maximum radius). The angles belong to AGroup2 (default radians).

 To cause angular values to be plotted along the arc of maximum radius,

TVAL

 The angles so plotted will always be in degrees, and will occur at the
interval specified for the highest level set (i.e. by default at pi/12 radians).
 Ticking of the outer arc may be done at the angular intervals by

TTIK

 for linear intervals on the radii, overriding of the automatic division
control is possible by

RDIV(D)   when D, REAL, will be the interval between ticks etc., on
the radii. This is subject to the same conditions as the similar feature
for cartesian plots.

 Automatic division may be restored by:

AUTORD

 Logarithmic axis ticking is controlled as for cartesian plots, and the

call to LOGLIM (described in section 2) controls logarithmic limits for both
types of plot.   Analagous with the calls for interpolation, point and function
plotting against rectangular axes (as described in section 2) are calls
for the similar purpose with polar plots.

POLPY3(R,THETA,N)   for third degree polynomial interpolation.

POLPY5(R,THETA,N)   for fifth degree polynomial interpolation.

 R and THETA are REAL arrays (of dimension at least N) holding the N pairs
of co-ordinates to be plotted. N must exceed the degree of the interpolation
polynomial. The number of interpolating steps between consecutive THETA
values is controlled by INTERP (see section 2), and by default is 5. THETA
is of AGroup2 (default radians).   Plotting a symbol at each co-ordinate pair
with no connecting line is accomplished by

POLPTS(R,THETA,N,NCHAR)   where R, THETA and N are as before, and NCHAR
($1 \leq$ NCHAR ≤ 48) is the desired marker symbol from the Marker Font. (N>2).

 To plot a polar function,

POLFN(FN)   where FN is a REAL FUNCTION, supplied by the user, and which
should be declared EXTERNAL in the calling program. The function should
have two REAL arguments and be referenced as   R=FN(THETA,DTHETA)
 where the value returned in the function is the radial value (i.e. R)
at angle THETA, and DTHETA, on return is the increment to be made in the
angle for the next plotted point. THETA and DTHETA are both in AGroup2 (default
radians).

 The general details of this call and its operation are directly analagous
to GRAPHF, the description of which is in section 2, and to which the user
is referred.   Analagous to SAMEX, SAMEY, SAMEXY are the references

SAMER which fixes the R range at the current range

SAMET fixing
the angular range

SAMERT fixing both the radial and angular ranges

 A different form of plot concludes this section. This is the "pie chart"
plot. In this, data is apportioned a sector of a circle in direct proportion
to its fraction of the total. As with all graphical plots, the piechart
occupies a portion of the current window. Within this space the chart is
centred.

PIECHT(R,CAP,PERCNT,N,IPERS)

   R is the radius of the chart, in
current users units, and this must not exceed half of 83% (the graph area)
of the minimum pane dimension.

   N is the number of data items to be
plotted.

   CAP is a <u>CHARACTER</u> array of dimension N, containing the
N captions of the data items. When percentages are to be plotted for each
data item, this will be appended to each caption.

   IPERS indicates
the form of data input, and whether percentages are to be plotted.

 If
IPERS is negative, the actual data value of each item is input through PERCNT,
in which case there will be exactly N sectors. When IPERS is positive, the
actual percentages associated with each are input through PERCNT. In this
case, a check is made by totalling the percentages 1 to N, and if a total
in excess of 100.1 is found an error is recorded and the call ignored. Conversely,
a total less than 99.9 will result in a blank sector appearing to force
the total to 100.

 When the absolute value of IPERS is 1 no percentage
will accompany the data names on the plot, while an absolute value of 2 causes
percentages to be plotted.

   PERCNT is a REAL array of dimension N,
holds the percentages or values (as determined by IPERS) associated with
each item.

Chapter 4 – Graphing in depth

  This section contains details of a handful of routines that will probably
be accessed only infrequently by the user. They do, however, offer the means
for the more ambitious to obtain superior graphical plots. Unless specifically
stated, the following facilities refer only to plots against rectangular
axes, i.e. those discussed in sections 1 and 2.   Although DIMFILM is able
to provide ticking of the various bounds and axes, as described in II 1.4,
it has no prior knowledge of whether axes will in fact be drawn. Where axes
are included, a tick at the axis point will result in a heavier density or
a colour "change" where axis and tick coincide. Conversely, if the axis
tick were omitted <u>and</u> the axis not drawn then a seemingly nonsensical
gap in the ticking would appear. By default ticking is suppressed at the
axis as it is felt that generally an axis will be incorporated if it falls
within range. To reverse this, so that ticking occurs at the axis point the
following routine may be called :

AXTIK

  The default condition may be restored by :

NOAXT

 It is obviously the user who must consider whether or not his axis will
be contained in the plot.   The ticking of bounding edges is directed inwards.
A set of routines exist to place outward pointing ticks on the bounding edges:

LXOPT for outward ticking of the lower edge

UXOPT for outward ticking
of the upper edge

LUXOPT for outward ticking of both upper and lower
edges

LYOPT for outward ticking of the left edge

RYOPT for outward
ticking of the right edge

LRYOPT for outward ticking of both left and
right edges

 It is also possible to plot any combination of edges (rather than the complete
frame of four edges, through GRFRAM – see II.1.4). The ticking of the open
ends of the plotted edges may also be controlled.

EDGES(LKJI,ITICK)

  LKJI is a four digit INTEGER, each digit being
0 or 1. There is a 1:1 correspondence between the digits and the four edges,
and a value 1 causes the relevant edge to be plotted, while for 0 it is
not plotted.   The correspondence is: I – lowest edge   J – left hand
edge   K – uppermost edge   L – right hand edge

  ITICK controls
the ticking of the open ends of the plotted edges (i.e. at ends of the plotted
edges not meeting another plotted edge). The values ITICK may have are :
 +1 – inward ticks   –1 – outward ticks   +2 – in/out double ticks
  0 – no ticks

  For example CALL EDGES(0011,2) would produce the lower and left edges

with double ticks.   It may sometimes be more convenient to know the actual
plotting co-ordinates (in the users overall frame system) of a point known
in terms of currently used graphing axes. This is necessary when one wishes
to overplot text on a graph, or, possibly, to blank out a section of the
graph (where it is known that no relevant data will occur) for future use
as a legend box.   To obtain the actual plot co-ordinates (XA,YA) of a point
with current graph co-ordinates (XG,YG) , call the following:

GRTOAC(XG,YG,XA,YA)

   Note: the desired coordinates (XA,YA) are actual coordinates in terms
of the users (world) coordinates (e.g. as determined through a reference
to BOUNDS). They take <u>no</u> account of any rotation or temporary origin.
As such they are suitable for use as blanking/clipping parameters, and should
only be used for plotting when rotation and temporary origin are both zero.
 There is provision for movement from the current position to a point referred
to by its graph coordinate (XG,YG); this may be either a positional or drawing
move.

GRON(XG,YG)   is equivalent to ON2; drawing will occur from the current
plot position to the <u>actual</u> point with graph coordinates (XG,YG) at the
current colour/intensity and in the current style (i.e. CSType 1).

GROFF(XG,YG)
  is equivalent to OFF2; the current plot position will be repositioned
to the <u>actual</u> point with graph coordinates (XG,YG) without drawing any
line.

  Note: in these cases the movement is to the <u>actual</u> point with graph
coordinates (XG,YG). This point will become the current plot position.
Effectively rotations and temporary origin are accounted for (although there
is <u>no</u> rotation within the graphing area).

 The sequence
     CALL GRTOAC(XG,YG,XA,YA)
     CALL ON2(XA,YA)
will not generally have the same result as
     CALL GRON(XG,YG)

 When there is no rotation and no temporary origin the results will be identical,
otherwise the coordinates (XA,YA), which are actual world coordinates, will
be transformed by ON2 according to the existing rotation and temporary origin
before the move is made. It is emphasised that (XA,YA) are the actual coordinates
relative to the user's bounds; they are not relative to the current plotting
axes.   Directly analogous routines exist for polar plots. These are

POLTOA(R,THETA,XA,YA)   which will convert the given coordinates (R,THETA)
on the polar plot to actual coordinates (XA,YA).

POLON(R,THETA)   to
draw a line (CSType 1) from the current plot position to polar point (R,THETA).

POLOFF(R,THETA)   to reposition the current plot position to polar point
(R,THETA).

   The notes pertaining to the equivalent cartesian routines apply equally
to those for polar coordinates.

Part 4 – CONTOURING

Chapter 1 – Basic Contouring

1.1 – Introduction

 DIMFILM includes a collection of subprograms for the production of basic
contour plots. These may be used in conjunction with all other facilities
of the overall system. As with the graphing routines, the user has considerable
freedom. Error checking is performed and warning messages included, these
being controlled through the check level. Users interested primarily in
contour plots should first familiarise themselves with the earlier sections
on basic plotting.   Contour plots may be freely intermixed with calls to
the general plotting routines. DIMFILM attempts to ensure that on exit from
any contour routine, all internal parameters are set to their values on
entry. The beam is also positioned so that diagrammatic plotting may be resumed
easily. Such properties as intensity, orientation and height of symbol strings,
broken lines etc., will be maintained. It is possible therefore, to interrupt
a diagrammatic plot to perform graphical plotting and then to resume. A few
limitations are, however, necessary. Primarily, on return, a broken line
pattern will commence at the beginning of its sequence (the consequence of
a beam off movement). Symbol strings will commence with <u>all</u> their initial
features active. If TXTCON has previously been called, it will be operable
on all but the first subsequent calls to SYMTXT.   Except where otherwise
indicated, all contour plotting routines are performed as CSType 1 operations
and all textual annotation will be treated as CSType 2. Additional annotations
and framing may be incorporated through use of the graphical annotation routines
described in Part II.   At present, all DIMFILM contouring routines require
regularly spaced input data. However, it is intended that fuller contouring
facilities will be made available in due course.

1.2 – Contour Location

 The contour plot is positioned in precisely the same way as graphing output,
thereby permitting the use of the graphical annotation routines. Thus, the
contoured data will be fitted into the current pre-clipping rectangle with
appropriate margins allowed for titles. By default the sub-area will be
centred as 83% of the pre-clipping rectangle in either direction. The contour
grid will be scaled to fit this area. This enables the user to position the
contour anywhere within the current bounds, and to generate multiple contour
plots per frame (through respecification of PANE). Orientation of the contour

```
plot is determined by the current mode  of the device.
```

1.3 – Contouring Regularly Spaced Data

```
 Contours are constructed based on a number of datum points.  This  section
is concerned with contouring regularly spaced data.  In this  case a rectangular
base mesh is provided with datum points at each  intersection.  The numeric
data is passed via a two-dimensional  real array GRID defined as:
```

REAL GRID(NX,NY)

```
  This defines a mesh of NX by NY elements, whereby the value of the  (i,j)th
element of the array GRID is the value of the mapping  variable (e.g. for
a physical relief map this would be the height)  at the (i,j)th intersection
of  the base defining mesh.  Pictorially, this is shown as:
```

```
 (1,NY)                                              (NX,NY)
   *-------*------  . . . . .  -----*-------*-------*
    |       |                       |       |       |
    |       |                       |       |       |
 (1,NY-1)--*-------*- . . . . . ---*-------*-------*
    |       |       |               |       |       |
    |       |       .               .       |       |
    .       .       .               .       .       .
    .       .       .               .       .       .
    .       .       .               .       .       |
    |       |       .               |       |       |
 (1,3)-----*-------*--- . . . . . -*-------*-------*
    |       |       |               |       |       |
    |       |       |               |       |       |
 (1,2)---(2,2)-----*---- . . . . --*-------*----(NX,2)
    |       |       |               |       |       |
    |       |       |               |       |       |
   *-------*-------*------  . . . ---*-------*-------*
 (1,1)   (2,1)   (3,1)               (NX-1,1)  (NX,1)
```

```
  The rectangular mesh defined by the regularly spaced data will be  mapped
onto the contouring (i.e. graphing) sub-area of the  current pre-clipping
rectangle.  That is,  the (NX-1) intervals in the X-direction are exactly
scaled to the  contouring area X-dimension, while the (NY-1) intervals in
```

the Y-direction are scaled to exactly fit the contouring area Y-dimension. Consequently, unless the proportion of the contouring area X:Y dimensions equals (NX-1):(NY-1) grid dimensions there will not be identical X- and Y-scales. If, for example, measurements have been conducted at different intervals in the X and Y directions it will be possible to reflect this in the definition of the pre-clipping rectangle (via PANE) to force equal scaling in the output plot. For example, where the data is defined via GRID(6,3), and sampling was at equal intervals in both the X and Y directions, the pre-clipping rectangle should be redefined with its X-dimension a factor of 5/2 [=(6-1)/(3-1)] of its Y-dimension. {This assumes equal annotation margins – the default.} Alternatively, GRID(7,4) with X sample interval half the Y sample interval reflects data from a square base mesh (6/3, but each X interval is half each Y interval), requiring a square contouring area; i.e. a pre-clipping rectangle with equal X and Y dimensions. {Same assumption.} <u>Note</u>: DIMFILM routines are written to ensure the integrity of all arguments (except, obviously, the case of arguments provided for the return of information to the calling program; for example, inquiry routines). However, in the case of contouring from a grid of data the values passed in array GRID <u>may</u> be modified in their **least significant bit** <u>only</u>. This obviates the provision of scratch space to the routines, but should be borne in mind if subsequent calculations are to be performed on the data (the effect is likely to be minimal). This effect is <u>not</u> cumulative between successive calls to contouring routines.

1.3.1 – Basic Contouring

At the basic level, the user will have his regularly spaced data stored in an array and will know the values at which contour lines are required. The most basic routine will compute contour intersections with the mesh lines (via linear interpolation between mesh points) and generate specified contours. This routine is intended as a quick preview facility and produces contours that are straight lines between these points. The production of smoothed contours is described in the next section (Refined Contouring). By default, contour values (where possible) will be automatically included across the lines.

CONTR(GRID,NX,NY,NC,CVALS)

GRID is a two dimensional REAL array of dimensions (NX,NY). This holds the datum values that define the rectangular mesh with NX columns and NY rows.

NX is the number of columns of data (NX>1).

NY is the number of rows of data (NY>1).

NC is the number of contour levels that it is desired to plot (NC>0).

CVALS is a singly dimensioned REAL array of dimension (NC) which holds the NC values at which contours are to be plotted.

The contoured area may be boxed in the same manner as for   graphical
routines; that is, by a reference to GRFRAM. Other   titling or labelling
may be applied as for the graphical   routines.

1.3.2 – Refined Contouring

  The user is able to control many defaults such as contour annotation and
interpolation. This latter is used to generate smooth contours. Once a
default is changed it will remain in effect until reset by the user.   The
cubic interpolation of contours may be simply controlled by the user via:

CINTER(ISTEPS)   which will turn on the cubic interpolation of contours,
where ISTEPS is the number of subintervals of the basic mesh that are to
be used in evaluating the interpolating curve. (ISTEPS>0)

NOCINT to
turn off this interpolation.

  The contours plotted will be generated using a cubic interpolation. It
should be noted that on some data forms unacceptable oscillations and inaccuracies
may be introduced. This should be borne in mind when interpreting results.
The number of steps at which the interpolating function is evaluated will
affect the smoothness of the plots – but will also be reflected in the degree
of computation involved. Clearly the size of the data grid will be a relevant
criterion when setting this value, as will the physical size of the contour
area in the frame.   The user has considerable freedom in specifying the way
in which contour values may be plotted across the generated contours. The
default is for such labelling to be produced wherever possible (proximity
to an edge, or next contour, may prohibit production of these labels).

NOCLAB will turn off the plotting of contour values.

CLABEL will
turn on the plotting of contour values.

CLABHT(HT,ANGLE)

  The user may specify the character height at which contour values will
be plotted and the orientation. The default character size is computed at
1/50 of the least dimension of the current pre-clipping rectangle (and may
be further scaled if CQUAD or CxHALF is active), and values to be generated
horizontally with respect to the frame.

CLABHT(HT,ANGLE)   will set the character height and orientation for
contour values used in evaluating the interpolating curve.

   HT is
the character height in overall user units. (HT>0.0 sets height, other values
will cause the default height to be effective)

   ANGLE is the angle
at which values will be produced, and is measured counter-clockwise from
the positive X-direction of the frame (AGroup 1, default degrees).

  The contour values that are output are, by default, produced according
to the Fortran format specification G9.3. This format may be changed by
the user. Additionally, in deriving the plotted contour value a scale factor
may be applied to the actual contour values.

CFMT(JFMT)   will set the format for outputting contour values to that
passed in JFMT (type CHARACTER, of length <11). This parameter should hold
the required <u>real</u> format as a character string commencing with an open
parenthesis and terminating with a close parenthesis. E.g. a valid parameter
would be '(F10.3)'.

CSCALE(SCALE)   non-zero values of SCALE will be
used as a factor by which all contour values will be pre-scaled before plotted
as annotation.

  Frequently, contour data may be symmetrical (for example, plots of some
electromagnetic fields). In these cases, the user need supply only a quarter
or half of the total pattern and the various contouring routines will fill
out the whole area by a series of mirrorings.

```
A----------------S----------------B
|                |                |
|                |                |
|                |                |
|                |                |
|                |                |
|                |                |
P----------------Q----------------R
|                |                |
|                |                |
|                |                |
|                |                |
|                |                |
|                |                |
D----------------T----------------C
```

Part 4 – CONTOURING

  The total area that is to be contoured is given, in every case, by ABCD.
The user may provide a data grid for either the upper-right quadrant, the
upper half or the right half and have the complete area symmetrically contoured
by appropriate reflection(s).

CQUAD will assume subsequent GRID data corresponds to the upper-right
quadrant SBCQ, and reflect about both axes in turn.

CVHALF will assume
subsequent GRID data corresponds to the vertical half SBCT, and reflect about
ST.

CHHALF will assume subsequent GRID data corresponds to the horizontal
half ABRP, and reflect about PR.

CWHOLE will restore subsequent GRIDs
to accommodate the whole area ABCD (this is the default).

  All the foregoing default modifiers should be referenced before any contour
for which they are to take effect.   There are alternate forms of contouring
routines, applicable to regularly spaced data, that may be used. These permit
contours at equal intervals to be plotted or for closer control of the line
colour/intensity for different contours.

CONTR2(GRID,NX,NY,NC,CSTART,CSTEP)   will generate contours at regular
intervals.

   GRID is the two dimensional REAL array of dimensions (NX,NY).
This holds the datum values that define the rectangular mesh with NX columns
and NY rows.

   NX is the number of columns of data (NX>1).

   NY is
the number of rows of data (NY>1).

   NC is the number of contour levels
that it is desired to plot (NC>0).

   CSTART is the value of the first
contour to be plotted.

   CSTEP is the value by which the contour value
will be incremented for successive contours.

CONTR1(GRID,NX,NY,NC,CVALS,ICSG)
  will generate contours at specified intervals, each being of independently
set CSGroup colour/style.

   GRID is a two dimensional REAL array of
dimensions (NX,NY). This holds the datum values that define the rectangular
mesh with NX columns and NY rows.

   NX is the number of columns of
data (NX>1).

   NY is the number of rows of data (NY>1).

NC is
the number of contour levels that it is desired to plot (NC>0).

   CVALS is
a singly dimensioned REAL array of dimension (NC) which holds the NC values
at which contours are to be plotted.

   ICSG is a singly dimensioned
INTEGER array of dimension (NC). The absolute value of all elements in ICSG
must be in the range [1,3]. The absolute value of ICSG(i) will be used as
the CSGroup identifier determining the colour/style in which contour of
value CVALS(i) will be plotted. When the actual value of ICSG is positive
the corresponding contour will be labelled in the <u>same</u> colour/style.

CONTR3(GRID,NX,NY,NC,CVALS,ICLUT)   will generate contours at specified
intervals, each being plotted according to a different colour look-up-table
reference. All arguments, bar the last, are identical with those for CONTR1.

   ICLUT is a singly dimensioned INTEGER array of dimension (NC). The
absolute value of all elements in ICLUT is a reference to the colour look-up-table
entry that will be used in plotting the corresponding contour and, when the
actual value is positive, the value label(s) for that contour. Each element
must have an absolute value not exceeding 255.

  If the user requires even greater control of contour line/annotation style
all the above CONTRx routines may be repeatedly called for the plotting of
single contour lines, with change of colour/style between each. However,
no account of adjacent contours can be taken when determining where annotations
are produced.

Part 5 – DIMFILM Guidelines

 The new DIMFILM has undergone a major rewrite. While rationalising certain
routine names and converting to Fortran 77 the opportunity was taken to improve
the internal operation of the product. Thus, all internal (i.e. inaccessible
to user) routines now have the prefix DFX. A similar convention was applied
to common blocks and internally accessed files. The user is therefore cautioned
not to use the prefix DFX for any purpose associated with DIMFILM. Failure
to respect this condition may well result in unpredictable results and the
probable failure of DIMFILM.

Appendix 1 – Using DIMFILM at ULCC

 DIMFILM is currently available on both the Amdahl 5890 and the Cray X-MP
at ULCC. It is intended that the user interface as described in this document
will not be subject to change. Please advise Graphics Development Group
at ULCC of any problems with this product. We shall endeavour to fix reported
bugs at the earliest possible moment. Similarly new features will be incorporated
as they become available and will be documented by revisions to this guide.
 It should be noted that some microfilm services (notably, colour film)
may be subject to a charge. In these cases users <u>must</u> ensure that their
intended usage is covered by an official college order. This should be directed
to the Microform Section of the Operations Division (not Allocation and Control
or Administration). The order may cover a single account number, a department
or whole institution, and may indicate either a monetary limit or expiry
date. Any enquiries regarding this should be directed to Claire Steward
(extension 365).

1.1 – Amdahl 5890 – VM/CMS Service

 The DIMFILM library is produced under VS FORTRAN Version 2, a Fortran 77
compiler available under VM/CMS. DIMFILM is built with the latest version
of this compiler.   Under CMS the implementation of DIMFILM and its internal
interface to the Dicomed film recorder makes use of virtual device addresses
– these may be in the range 180-183 (inclusive), and so should be avoided
by the user.

1.1.1 – VS Fortran Version 2

 The DIMFILM library for use with VS Fortran is available on the GRAPHICS
mini-disk. Access to this must be achieved by inclusion of an appropriate
command: GIME GRAPHICS   To access the DIMFILM library it is further
necessary to incorporate a couple of GLOBAL commands. These, for regular

users of DIMFILM, might best be included in their PROFILE EXEC file:

```
GLOBAL TXTLIB DIMFILM VSF2FORT CMSLIB

GLOBAL LOADLIB VSF2LOAD
```

 A typical Fortran program would appear as:

```
PROGRAM TEST
...
CALL D35
.. <Program including calls to DIMFILM routines>
CALL DIMEND
STOP
END
```

 This should be compiled, loaded and executed in the normal manner. The
GLOBAL commands will ensure the correct satisfaction of references to DIMFILM
routines at the appropriate point.
 The DIMFILM library is large, the memory required will depend on the modules
referenced by the user.   When multiple libraries are used with DIMFILM it
is possible that the default Loader table sizes may be exceeded – to avoid
this, the default table size may be increased by including a command similar
to the following: SET LDRTBLS 14

  Note: Users who are new to the Amdahl will find general documentation
on running jobs in the Amdahl Handbook, published by ULCC. In particular
see Chapter 2.2 for details of running jobs and Chapters 3.0 and 3.2.1 for
details of using the FORTRAN compiler.

1.2 – Amdahl 5890 – MVS Service

The DIMFILM library is produced under Fortran 77 and made available for
use with both VS Fortran and Siemens Fortran 77. Separate libraries are
maintained for use with these compilers. DIMFILM is built with the latest
version of each compiler, and so all jobs should follow the JOB statement
with   /*JOBPARM P=PROC02

1.2.1 – VS Fortran

The DIMFILM library for use with VS Fortran is available on SYS2.DFILMVS.LIB,
and should be used in conjunction with release 5.0 of VS Fortran, which is
obtained by immediately following the JOB statement with   /*JOBPARM
P=PROC02   To access the DIMFILM library LLIB1='SYS2.DFILMVS.LIB' may be
used in the FVS procedure. A simple job to produce output on 35mm black
and white film might be :

```
  JCL

     //<jobname>   JOB  '<name and delivery>',REGION=2000K
     //*PASSWORD <MVS password>
     /*JOBPARM P=PROC02
     //        EXEC    FVSCLG,LLIB1='SYS2.DFILMVS.LIB'
     //C.SYSIN  DD  *
          PROGRAM TEST
          ...
          CALL D35
          .. <Program including calls to DIMFILM routines>
          CALL DIMEND
          STOP
          END
     //
```

```
  Phoenix 3

     //<jobname>   JOB  '<name and delivery>',REGION=2000K
     //*PASSWORD <MVS password>
     //          EXEC   PHOENIX3
     FVSCLG PROGRAM=%H+ LIBRARY=SYS2.DFILMVS.LIB
          PROGRAM TEST
          ...
```

```
        CALL D35
        .. <Program including calls to DIMFILM routines>
        CALL DIMEND
        STOP
        END
  +
  //
```

The DIMFILM library is large and it will almost certainly be necessary
to request more than the default memory allocation using the REGION parameter
on the JOB statement. The default is 600K and the maximum is over 7000K,
although no more than 6500K should normally be requested. (Numbers are decimal
and specify units of 1K (1000) bytes).   When multiple libraries are used
with DIMFILM it is possible that the default Linkage Editor/Loader table
sizes may be exceeded giving an MVS IEW0484 TABLE OVERFLOW error. To avoid
this, the default table sizes may be increased by including the following
parameters in the appropriate JCL EXEC statement:

```
   Using the Linkage Editor:
        LPARM='SIZE=(2000K,64K)',LREGN=2000K
   or using the Loader:
        LPARM='SIZE=2000K',LREGN=2000K
```

1.2.2 – Siemens Fortran

The DIMFILM library for use with Siemens Fortran is available on SYS2.DFILM77.LIB.
This may be accessed by utilising the LLIB1 parameter in the F77 procedure.
It is recommended that Version 4 of Siemens Fortran be used – considerable
savings in CP time and memory can be achieved. A simple job to produce output
on 35mm black and white film might be:

```
    //<jobname>   JOB  '<name and delivery>',REGION=2000K
    //*PASSWORD <MVS password>
    /*JOBPARM P=PROC02
```

```
//      EXEC    F77CLG,LLIB1='SYS2.DFILM77.LIB'
//C.SYSIN  DD  *
       PROGRAM TEST
       ...
       CALL D35
       .. <Program including calls to DIMFILM routines>
       CALL DIMEND
       STOP
       END
   //
```

```
 Phoenix 3

   //<jobname>   JOB  '<name and delivery>',REGION=2000K
   //*PASSWORD <MVS password>
   //            EXEC    PHOENIX3
   F77CLG PROGRAM=%H+ LIBRARY=SYS2.DFILM77.LIB
       PROGRAM TEST
       ...
       CALL D35
       .. <Program including calls to DIMFILM routines>
       CALL DIMEND
       STOP
       END
   +
   //
```

 The DIMFILM library is large and it will almost certainly be necessary
to request more than the default memory allocation using the REGION parameter
on the JOB statement. The default is 600K and the maximum is over 7000K,
although no more than 6500K should normally be requested. (Numbers are decimal
and specify units of 1K (1000) bytes).   When multiple libraries are used
with DIMFILM it is possible that the default Linkage Editor/Loader table
sizes may be exceeded giving an MVS IEW0484 TABLE OVERFLOW error. To avoid
this, the default table sizes may be increased by including the following
parameters in the appropriate JCL EXEC statement:

```
   Using the Linkage Editor:
       LPARM='SIZE=(2000K,64K)',LREGN=2000K
   or using the Loader:
       LPARM='SIZE=2000K',LREGN=2000K
```

Note: Users who are new to the Amdahl will find general documentation
on running jobs in the Amdahl Handbook, published by ULCC. In particular
see Chapter 2.2 for details of running jobs and Chapters 3.0 and 3.2.1 for
details of using the FORTRAN compilers.

1.3 – Cray X–MP – UNICOS Service

The DIMFILM routines are available as a library, compiled with the CFT77
compiler and accessible on /lib as libdimfilm.a. The normal usage will involve
only the declaration of dimfilm as a library; either to segldr or cf77, via
the -l option. There is no need to use any other graphics library in order
to produce output on film. A simple job to produce output on 35mm black and
white film might be:

```
cat test.f <</EOF
      PROGRAM TEST
      ...
      CALL D35
      .. <Program including calls to DIMFILM routines>
      CALL DIMEND
      STOP
      END
/EOF
cf77 -l dimfilm test.f
```

alternatively, the last command may be replaced by:

```
cft77 test.f
segldr -l dimfilm test.o
```

The DIMFILM library is large; the required memory will depend on the
modules acccessed by the user program.   Where plot data is staged via the
front end, an indication of file size and frames generated will be given,
but confirmation of successful staging is a function of the front end.

   Note: Users who are new to the Cray will find general documentation on
running jobs in the Cray Handbook published by ULCC. In particular see Chapter
2.2 on running jobs and Chapter 3.1 on the use of the CFT77 FORTRAN compiler.

Appendix 2 – Graphics Devices

 There are many types of graphic output devices for the display of computer
generated images. These may differ significantly in the way they generate
graphical output, and this may have considerable effect upon the visual
appearance of pictures (which may look very different when directed to alternative
devices). This can be very confusing. An understanding of the basic principles
of the operation of a device is essential in order to be able to see why
it produces the results it does. There may also be significant differences
in the functioning of devices that affect their suitability to certain applications,
or that dictate alternate modes of operation for efficient utilisation.

2.1 – Image Appearance

 Of paramount concern to the user of any graphics output device is the actual
appearance of the picture – and how it may differ from that expected. This
section deals with the makeup of the image in its final perceived form and
how different devices may present the same data. A generalised description
of picture production is followed by consideration of different device classes.

2.1.1 – Colour Generation

 [Note: Terminology in this section should not be interpreted rigorously
or pedantically; the intention is to introduce the inexperienced user to
the complexities of colour mixing, etc. Consequently, terms such as colour
and hue are used here in a general sense, and are interchangeable.]   There
are, essentially, two ways in which a colour image may exist; it may be permanent
(i.e. hardcopy) or transient (e.g. electronic image on a cathode ray tube
– **CRT**). Additionally, a permanent image can be created directly or
derived from a transient image (or from a combination of several). All these
methods of image formation result in subtle (and not so subtle) differences
in the final product.   Transient images are typically produced through emission
of (coloured) light from a luminous source. Indeed, the CRT relies on the
eye/brain physiological interpretation of colour whereby <u>any</u> hue may be
created through the admixture of appropriate proportions of the three primary
colours; these being, red, green, and blue. Equal combinations of all three
(saturated) yields white, while varying the proportions of each independently
will create any other colour. Red, green, and blue, then, are known as the
additive primaries (or simply as the primary colours/ primaries).

Mixing the saturated primaries in equal proportions gives the following results:

```
       none              ==> Black
       Red + Green       ==> Yellow
       Green + Blue      ==> Cyan
       Blue + Red        ==> Magenta
       Red + Green + Blue  ==> White
```

 It is the application of luminous sources of the additive primaries that
is fundamental to the colour CRT (and television). On the colour CRT, red,
green, and blue potential light sources are located in close proximity to
each other, where they may be independently excited in varying proportions.
Each such cluster (or triad) of red/green/blue will then be interpreted
by the viewer (at sufficient distance) as a single source of the hue associated
with that combination. Such spatial integration will be referred to again.
{This is why colour graphic displays are often referred to as RGB. They
accept separate signals for each of the **R**ed, **G**reen, and **B**lue
components. The standard broadcast television transmission technique relies
on a different convention, which is briefly treated in the discussions on
colour models in the main part of the User Guide.}   As the three primaries
are independent (i.e. none of them can be formed of any combination of the
other two), it will be seen that the hues resulting from the combination
of any two primaries have the property that they contain <u>no</u> component of
the third primary. Thus, for example, cyan (which is the combination of green
and blue) has no trace of red. That is, a cyan pigment or dye will filter
white light (the combination of red + green + blue) of all red. For this
reason cyan, magenta and yellow are known as the **subtractive primaries**
 (or, alternatively, as the   secondary colours/secondaries). The value
of this will become evident in the later discussion.   The relationship between
the additive and subtractive primaries can be clearly shown by a diagram
of the hue spectrum:

```
                       *
                    _     _
       green  120.0  -           -  60.0  yellow
                 *                 *

              _                        _

            _                          _


          _                          _
   cyan  180.0 *        hue spectrum        * 0.0  red
            _                        _
```

```
             _                           _


         _                           _


             *                   *
   blue  240.0  –            –  300.0  magenta
                     _     _
                  *
```

 Permanent images may either be reflective or transmissive. The former
encompasses all forms of hardcopy on an opaque base (e.g. paper) while the
latter is the province of transparencies requiring projection. However,
if one considers an overhead transparency this may be viewed either with
back-lighting (i.e. analgous to projection) or it may be placed on a reflective
(e.g. white) background and viewed under incident lighting. This is a good
analogy to the effect of transmissive ink on the printed page; the coloured
transparency functions by filtering the incident light and transmitting
only a particular colour through the clear base medium, this is reflected
(in this case) by the white base and retransmitted through the transparency
with (ideally) no further filtration. It is reasonable, then, to direct attention
first to hardcopy output and to look at the ways in which coloured inks are
used by the printer (these techniques are relevant to all present forms
of colour hardcopy device).   The printed colour on a page is the result of
either reflective or transmissive dyes/pigments and the way they are combined
on the page. The first category relies on an essentially opaque pigmented
ink (or paint) which relies on reflection of the required colour and absorption
of all others. Mixtures of such inks will, when applied to the paper, yield
a colour that corresponds to the proportional mix of those yielded by each
ink individually (perfect mixing is assumed throughout this discussion).
If two such inks are applied <u>sequentially</u> to the paper (the first allowed
to dry thereby precluding mixing) the resultant colour will be that of the
last used ink (in reality the opaqueness of a paint or ink can result in
some modification of the perceived hue). In essence, for such inks the colour
is dependent on the particular properties of a single ink. This is the
method of colour creation used in oil painting – the pigments are mixed additively
and then applied to the canvas. However, it is well worth noting that the
artist has a different colour requirement than the physicist or photographer.
These two are intent on capturing the real world (theoretically or photographiclly,
respectively). For them the hue spectrum based on the red, green, blue model
(and, thus, cyan, magenta, yellow) is an adequate representation of light
(i.e. the combination of various monochromatic frequencies). The artist
is dependent on displaying colour through the removal of various combinations
of monochromatic frequencies from the incident light through selective absorption
by the pigments in his paints. In fact, the standard artist's palette comprises
red, <u>yellow</u>, and blue as his primaries with corresponding secondaries orange,
green, and violet. White is now the absence of any colour, and a separate
pigment is specifically needed for black. The artist also relies on a white
pigment to lighten any of the hues. It would not be possible to use the
additive primaries as the mixture of any two would (theoretically) remove
<u>all</u> colour and yield black. In practice, red and blue paints result in
violet due to the pigments used.

```
none              ==> White (canvas)
Red + Yellow      ==> Orange
Yellow + Blue     ==> Green
Blue + Red        ==> Violet
Red + Yellow + Blue ==> Olive-Brown
Black             ==> Black
```

 Turning to transmissive dyes/inks, a single such ink will transmit its
own colour (absorbing others) and this may then be reflected from the base
material and re-transmitted. Such inks are translucent. If two (or more)
are overlaid, the first will transmit its own colour which will then pass
into the second (lower) ink, which will in turn transmit <u>only</u> its own colour.
Thus only that subset of colours that is common to <u>both</u> inks can be transmitted
through to the base, where it will be reflected and re-transmitted. Each
ink has subtracted (i.e. removed, or filtered) particular colour components
from the incident light. From the earlier consideration of the relationship
between the primary colours and their secondaries, it should be fairly obvious
that these latter colours are ideal candidates for the printer's palette.
Clearly, the additive primaries would be no use – any combination would
bar all <u>transmission</u> (each will pass only its own component). However,
the subtractive primaries will be seen to be ideal as each will remove <u>only</u>
its complementary primary and transmit the other two.

```
none                     ==> White
Yellow + Cyan            ==> Green
Cyan + Magenta           ==> Blue
Magenta + Yellow         ==> Red
Yellow + Cyan + Magenta  ==> Black
```

  Note: it must be remembered that these are
        transmissive, <u>not</u> reflective.

 In actuality, most printing processes based on the subtractive primaries
employ a fourth ink for perfect representation of black. {Theoretically,
yellow + cyan + magenta yield black, whereas the mixing of printer's inks
tends to produce a rather dark brown.}   The foregoing description of hardcopy
colour has been on the basis that each point in the picture is coloured by
application of one or more inks to that point. As the standard printing
process is not suited to delivery of varying quantities of each ink to a
point (rather any point either has none or a fixed amount of each ink) the
subtractive process would be capable of printing only eight discrete colours
(being, yellow, cyan, magenta, green, blue, red, black and white). Clearly,
much printed material requires rather more than eight hues. This may be
achieved through different techniques.   If a standard newsprint picture
is studied closely it will be seen that this is composed of a vast number

of discrete points. For a black and white picture the points are dispersed
with varying size. Thus in a given area there will be a greater density
of ink where the dots are largest, while those with decreasing size will
have less ink density over an area. When viewed from a normal reading distance
this variation in dot size, and hence ink density, will be perceived as varying
shades of grey – between black and white (assuming white paper). This techniques
is known **halftoning**. The print screen may be around 60 to 80 dots per
inch for newspapers, up to 150 dots per inch in quality magazines and books.
 For colour printing a similar technique may be applied. The master image
(e.g. film) is passed through a process known as colour separation to yield
four colour masters (i.e. yellow, cyan, magenta, black). These are screened
and colour plates are produced for each colour; the plates now consisting
of single tone images with areas of different dot size (and so different
colour density). The final step is for the final image to be printed from
each plate with its corresponding ink (preserving registration between each
pass). This is basic four- colour lithography. {Note, that as overlapping
colours are used this requires the use of translucent inks.} These techniques
are all relevant to computer generated graphics (and may be applied to transient
images, also). Here the problem is that the dot size is likely to be fixed,
as is the display resolution. Considering monochromatic images first, halftoning
may be simulated if four pixels (i.e. an array of 2x2 dots) is used to represent
each point in the image. With a bi-level device (i.e. monochromatic with
no intensity modulation – black or white, only, at any pixel/dot) the 2x2
array gives the option to set any number of the four points – from none to
all four - as black, thereby giving 5 levels of dot density. If this is
viewed from sufficient distance (i.e. the pixels are sufficiently small)
the perceived effect is that the bi-level device is now showing 5 levels.
If the device actually offered 4 levels (zero to three) then a 2x2 arrangement
could have a summed value of zero to twelve and the perceived range would
be extended to thirteen. This process – equivalent to screened half-tones
– is known as **dithering**, dots of different intensity/colour being grouped
to create the illusion of a greater range (due to spatial integration).
When the device is a colour display/printer with only three primaries available
then a 2x2 array yields 125 different colours/shades, because for <u>each</u>
primary, 5 intensities may be achieved. It is worth noting that dithering
can be applied to opaque (i.e. purely reflective) paints, although the range
will then less as each pixel can only be one of the paints (assuming no mixing).
Dithering 3 such colours over a 2x2 array yields a total of 34 different
colour combinations (compare with 125 for subtractive inks). {The present
paragraph has, up to now, assumed that the image data is of lower resolution
than the display so multiple pixels (dots) may be assigned to each data point
in the image. Where the image and display are of the same resolution the
same techniques may be applied to increase the intensity/colour range of
the device through **ordered dithering**. A decision, based on the actual
address of each point in the image and its intensity, is made as to how it
should be represented in the display. For areas of constant value the appearance
will be as for a similar halftone pattern, and the effect will only be noticed
in areas of changing colour/intensity.}

2.1.2 – Device Categories

 In general, drawing may be performed as line drawing (**calligraphic**)
or painting (**area fill**). Computer graphics uses a terminology which
reflects the technology used. Thus line drawing is referred to as **vector**
 plotting (a straight line is plotted between two endpoints), while painted
– or shaded – images are termed **raster** plots when produced on particular

output devices (which construct a picture through a matrix of discrete points).
Although specific devices may be better (or only) suited to one mode of
drawing there are usually ways in which both modes may be represented (or,
at the least, approximated).

  <u>Pen plotters</u>   These are the closest device to the draughter's pencil
and drawing board. Here it is generally the case that lines are drawn in
the same order as they are generated by the computer program. Thus overlapping
lines (or areas of colour) will be drawn as such and the effect will depend
on the colours used and the type of ink; with felt-tips, and other pens
with translucent ink, line intersections will be noted for their change in
colour as for subtractive processes (if the additive primaries are used
intersections will – in theory – be black). Pigmented, opaque inks will
leave intersections of lines reflecting the colour of the last drawn line.
{The same applies to attempts to shade areas with multiple, parallel strokes
of the pen.}

  <u>Raster CRT displays</u> are considered next, due to their relevance
in discussions concerning other graphic output devices; although they generate
transient images rather than the permanence of hardcopy output. The images
produced on such devices are arrays of points (the red/green/blue triads)
known as pixels and (normally) spaced regularly across the display screen
with regularly spaced lines of pixels. The whole image is stored in the
display memory with a triplet of values stored for each pixel to represent
the red, green, and blue intensity levels. Maximum intensity n all three
combined will yield white. Values are assigned to each pixel affected as
graphics instructions are processed – the order is crucial. Output vectors
are evaluated for their point of intersection with the scanlines, and those
points assigned the colour of that vector (a process known as rasterisation).
In assigning colour (similar arguments apply to monochromatic displays)
to any point there are two basic options: the point may be assigned the
colour of the most recently drawn line, or any value already at a point may
be merged with that of this vector. The first option is equivalent to the
use of opaque inks, while the second has other properties. Although merging
of colours is analgous to additive mixing it is <u>not</u> identical. The merging
of each component may be additive (although a maximum for each component
may not be exceeded) or by taking only the maximum contribution to each component
from the intersecting lines. Using the notation (r,g,b) to be the value
of a colour triplet some examples of merging can be given.

<u>Additive overlay</u>

      <u>line1</u>   <u>line2</u>   <u>result</u>

  i) (0.4,0.5,0.0) + (0.6,0.0,0.0) ====> (1.0,0.5,0.0)

 ii) (1.0,0.5,0.0) + (0.6,0.0,0.0) ====> (1.0,0.5,0.0)

iii) (0.7,0.6,0.1) + (0.7,0.1,0.8) ====> (1.0,0.7,0.9)

<u>Combined overlay</u>

      <u>line1</u>   <u>line2</u>   <u>result</u>

```
    i) (0.4,0.5,0.0) + (0.6,0.0,0.0) ====>  (0.6,0.5,0.0)

   ii) (1.0,0.5,0.0) + (0.6,0.0,0.0) ====>  (1.0,0.5,0.0)

  iii) (0.7,0.6,0.1) + (0.7,0.1,0.8) ====>  (0.7,0.6,0.8)
```

<u>Opaque overlay</u>

| <u>line1</u> | <u>line2</u> | <u>result</u> |

```
    i) (0.4,0.5,0.0) + (0.6,0.0,0.0) ====>  (0.6,0.0,0.0)

   ii) (1.0,0.5,0.0) + (0.6,0.0,0.0) ====>  (0.6,0.0,0.0)

  iii) (0.7,0.6,0.1) + (0.7,0.1,0.8) ====>  (0.7,0.1,0.8)
```

   It can be seen that the different techniques give very different results.
In no case can the maximum intensity of a component be exceeded so, for
example, multiple maximum reds (in additive overlays) would always result
in maximum red. {A further complication may occur when the device drives
the CRT through a look up table. In this case, the values in the device
(pixel) memory are index pointers into look up tables; consequently, it
can be possible to combine – logically OR – the overlaying LUT pointer values
(in binary arithmetic) yielding a new pointer with the output colour having
<u>no</u> relation to the colours held in the look up table for either of the
two overlapping colours.} Filled areas <u>may</u> be handled similarly, but often
a hardware fill is utilised which will colour all points of the area with
the required colour in an opaque mode. Similarly, pixel scan line data
passed in graphics data will normally overwrite any existing pixel values
(no merging). Subsequent vectors, however, will be handled in whichever
mode is currently supported/active. For most purposes the display may be
considered as erasable, with the most recently interpreted graphic instruction
taking precedence of visibility.

 Vector refresh displays rely on buffering plot commands into device
memory (as a display list) from where they may be continually scanned and
displayed on a long persistent phosphor CRT. This is <u>not</u> the line–by–line
scanning of a raster device; each directive is traced across the CRT as
a <u>vector</u> (in any direction). There is a limit to the number of directives
that may be held and cycled through in the display list before flicker becomes
noticeable. Overlapping vectors will all be traced and the effect will be
the repeated intensification of the overlapping parts of the image. (Refresh
displays are generally monochrome – although bi-colour devices are available.)

 Storage tubes, more correctly known as direct–view storage tubes,
are in the category of "additive" displays. Each plot instruction remains
visible until the whole display is erased. As an instruction is "plotted"
it is written by an electron beam on a fine dielectric mesh behind the phosphor
on the CRT; a flood gun continually transfers this image onto the CRT. The
image on the wire mesh is additive, but the range of viewable intensities
is small so overlapping plot elements show little hotspotting – the most
intense of the overlapping elements will predominate.

<u>Non-photographic, colour hardcopy devices</u> include many types of
output technology. These may include electrostatic plotters, ink jet plotters/printers,
thermal plotters/printers (e.g. wax/dye transfer), inked ribbon systems (matrix
printers), and laser printers. Each of these is characterised in the way
the picture is produced. Whether it is a single or multi-pass operation,
the image is drawn by a series of closely spaced lines composed of many points
(which may overlap). The resolution of these devices is expressed in points
per scan line unit and lines per unit measure (sometimes points per unit
area are given). Typical resolutions are between 100 and 560 points per
inch (although some ribbon printers may be lower). For all these devices
a whole image is (notionally) created as an array of all pixel points with
a colour/intensity value at each. Thus the graphics instructions are pre-processed
before plotting (that is, rasterised) to generate the appropriate array of
pixel data, which is then output line by line. (For colour applications,
it is generally the case that three – or four, if a separate black is provided
– complete passes of the array are necessary, one for each ink. This requires
careful registration of the plot medium between passes to ensure that scan
lines overlap correctly. To avoid this some plotters have a complex mechanism
to enable a single pass across the image area; some ribbon printers shift
the ribbon and overstrike each line before advancing the paper, and some
ink jet devices perform a similar one pass action switching between their
jets.) For bi-level monochrome plotters one of the dithering techniques
may be employed to give multi-level halftone pictures. For colour plotters,
most employ subtractive inks and do not need to rely on spatial integration
for the basic mixing of colours. Dithering may be used to increase the colour
range, although this will often depend on the source of data (i.e. the rasteriser).
This source may be a raster display system feeding the hardcopy RGB signals
at its display resolution and limited to a small colour range (by virtue
of the number of pixel planes); such signals must be mapped onto CMY space
for the subtractive process. In rasterising data there is one further complication
– efficient rasterising, without a large buffer, will require pre-sorting
of vectors. This may mean that vectors are scan converted in an order different
from that in which they were generated by the user. However, this problem
is less likely to occur with the large capacity rasterisers now associated
with many hardcopy units (or with the raster displays used to buffer the
data for lower resolution devices).

<u>Film recorders</u> are of two sorts. Both function by recording, on
film, the trace of a beam on a precision CRT through an optical system.
Those with a vector capability (i.e. a vector may be drawn on the CRT) function
in an additive manner. The shutter remains open and film is exposed continuously
as the graphics data is interpreted and plotted on the CRT. Such film recorders
may be considered as active. This image (on the CRT) is transient and is
<u>not</u> refreshed in the way the display of a raster graphics terminal (or
television) is. The whole image is not built up and then captured, rather
each element is recorded as it is created (and the CRT image of that element
would appear to the eye as no more than a spot, of varying intensity, traversing
the shape of the element). This leads to an additive or cumulative exposure
on film – multiple intersections of vectors, or overlapping solid fill areas,
will result in an exposure equivalent to the sum of <u>all</u> the coincident
intensities. Indeed, the nominal device maximum may be exceeded many-fold
at any point. The effect of this then becomes a function of the physics/chemistry
of film. Essentially, film can only be rendered clear (reversal stock) by
maximum exposure. However, an effect known as halation then becomes apparent
and the 'bright' point/area leeches into surrounding areas with a fuzzy appearance.
Colour film is produced on film recorders by splitting each element to be
drawn into its constituent colour components, which are then plotted separately
onto the CRT (near-white spot) and recorded through the corresponding colour
filter (normally red, green, and blue are used). Generally the whole image

is pre-streamed for each individual filter thereby minimising filter changes.
With colour, multiple exposures through any filter will be additive and
once a colour is fully saturated on film further exposure may adulterate
the image to the extent that not only halation occurs but the colour will
be (ultimately) washed to white (all colour layers clear, on reversal stock).
The additive nature will be seen dramatically where two maximum intensity
primaries overlap (thus, red and green lines crossing will show a very predominant
yellow point of intersection). This gives problems when a coloured background
is required behind a vector picture; colours must be selected so their combination
results in the required vector colour! The other class of film recorder
handle only raster data. This may be through pre-processing (i.e. rasterising)
via software on a host computer, or they may fed RGB signals from a graphics
raster display (which then acts as the rasteriser). The image is recorded
on film line by line, and it is useful to consider such film recorders to
be passive. This type then will function as a raster device, and generally
the last plotted entity will determine the colour at any point.

Appendix 3 – Device Specifications

 This Appendix describes the characteristics and capabilities of those devices
supported by DIMFILM. As additional devices are made available in DIMFILM
their descriptions will be added here. The previous Appendix should be
consulted for comparison of various device categories in so far as their
mode of operation affects the appearance of graphical output.

3.1 – DICOMED Colour Film Recorder D(1)48C

 The Dicomed D48C graphic COM system is a high precision colour film recorder.
In D148C configuration it includes a Digital Equipment Corporation controlling
mini-computer. The name D48C refers to the graphic recorder, comprising
firmware controller, precision CRT, film transports and optical assemblies.
 The D48C is capable of image recording in vector, raster and alphanumeric
modes, and may be used for either colour or monochrome recording dependent
on film/optics used. It is ideally suited to many applications, including
scientific plotting, business graphics (vector or raster), animation, engineering
and design, and alphanumeric COM (**C**omputer **O**utput **M**icroform,
e.g. microfiche output – which may incorporate other graphics). As a film
recorder the Dicomed D48C captures the image of a variable intensity spot
as it traverses a precision cathode ray tube. This CRT uses the P48
phosphor (wide spectral response), and colour is achieved by insertion of
coloured filters in the optical path. For each colour filter the corresponding
components of the image are traced – that is, recording is in an additive
manner (filters being used singly, _not_ in combination). The filter assembly
comprises seven colours (red, green, blue, yellow, magenta, cyan, and neutral)
and is used in conjunction with specially designed colour corrected optics.
(Note: colour output at ULCC is currently produced by accessing the red,
green, and blue filters alone and streaming the plot data accordingly.)
For use with black and white film the neutral filter may be utilised; however,
for the highest resolution in black and white on the aperture format a separate
narrow spectral response lens is used (without filter assembly). Three optical

assemblies are in use at ULCC for graphic COM: F300 for 16mm cine (colour
and black and white), F304 for 35mm slide (i.e. comic – colour), and F309
for 35mm aperture card format (black and white). The film transport (F333A
– manufactured by Marron Carrell), which sits atop the optical assemblies,
is able to handle 16mm (perforated) and 35mm (both perforated and unperforated)
with 400 feet film magazines. This is accomplished through changes to the
transport mechanics and use of appropriate aperture plates. In **vector
mode** , the D48C has an addressability of 32,768 x 32,768 upon the CRT (the
optics/film format will result in a [maximised] subset of this being available
for each film format). The D48C is capable of very fast vector plotting:
a minimum length vector (single point) will take 15 microseconds, while a
full-screen, maximum intensity vector can be drawn in 45 milliseconds (single
pass). The time required is roughly proportional to the vector length.
Exposure may be at any of 256 levels (for any one filter). In drawing a vector,
the D48C deflection system is used to accurately position the beam at selected
points on the 32K x 32K matrix when the exposure logic turns the beam on
for the requisite time. The points are selected so that the vector is drawn
as a continuous and uniform intensity line (through ensuring adjacent points
touch). In **raster mode** , the D48C is able to plot on a matrix of 4096
x 4096 addressable points (pixels – i.e. picture elements) on the CRT (again,
a subset is available to each film format) at any of 256 exposure values.
The difference between vector and raster addressability is due to the spot
size generated; the raster spacing ensures a continuous image through the
touching of adjacent points. The Dicomed provides the user with the means
to construct display pixels as composite points with control over point and
element spacing. This enables other raster "resolutions" to be imaged.
Point plotting in raster mode may be accomplished at speeds of 100K–250K
points per second, enabling maximum resolution raster images in full colour
(three passes) to be plotted in around 3 minutes. The maximum flyback time
for repositioning the beam to the next scan line is 100 microseconds; for
certain applications a short end of line delay of (about) 15 microseconds
may be introduced. In **point mode** , the D48C may be positioned to any
of the CRT's 32,768 x 32,768 addressable points, which may then be exposed
at any of 256 levels. The maximum time to position to any point is 100 microseconds,
while exposure may be from (about) 4 to 9 microseconds. In **character
mode** , the D48C has a high speed hardware character generator. This is
particularly relevant to alphanumeric COM applications. The character fonts
are stored in local memory, and depending on complexity a number of fonts
may be accommodated simultaneously, with less than 50 milliseconds required
for font changes. Characters may be generated at different sizes and in
four orientations, as standard. Speed depends on font complexity, but for
typical alphanumeric COM applications speeds in excess of 30,000 characters
per second may be achieved. (In practice, throughput is dependant on input
rates to the D48C.) The hardware character generator functions at only one
exposure level and differs from other drawing modes through operating as
a stroke generator. Each character is drawn at high-speed as a series of
short vectors and the CRT beam is unblanked throughout the duration of each
stroke. **Intensity/exposure** control is over an 8-bit (i.e. 256 level)
range. However, this range is selected according to the plot mode in use:
vector, raster, point, and character modes all access different CRT intensity
scales in order to achieve compatible film exposure when modes are intermixed
on a frame. [It should be noted that hardware characters may only be generated
at a single exposure level.] The D48C incorporates further internal exposure
tables that are selectable according to the optics and emulsion characteristics
to ensure consistent film exposure under the different operating conditions.
Separate table references are made according to the selected filter. These
tables are programmable so that alternative emulsions may be used and the
system calibrated accordingly. Indeed, calibration is required for each
film type and for each optical assembly with which it may be used: this
calibration is for each plotting mode and (for colour assemblies) each filter.
In vector mode the internal intensity value is dynamically compensated according

to the angle at which a vector is drawn. The exposure technique is **time
modulation**   which controls the density of individual points directly through
the duration for which the CRT is excited. This ensures precise control
and a high degree of repeatability with consistent line width.   The geometrics
 **of the D48C offer a very high precision with minimal distortion and a
high degree of repeatability. With regard to the major axis**, trapezoiding
, rectangularity , linearity , **and** line curvature **(or** pin
cushion distortion **) are all maintained to less than (plus or minus) 0.1%,
while orthogonality of the horizontal and vertical axes is within (plus or
minus) 0.15%. The** spatial repeatability **is within (plus or minus) 0.02%
over a ten-minute interval, and (plus or minus) 0.05% over a thirty-minute
interval (after a one-hour warm up).   The** photometrics **of the D48C offers
a dynamic exposure range of 256 levels. In terms of diffuse density units
(D), the range on a typical panchromatic film (e.g. Kodak Plus-X) will be
1.5D when processed to gamma 1.0 (or 2.0D for gamma 1.5). The** exposure
uniformity **over the entire plotting area has a maximum deviation of (plus
or minus) 0.15D (measured through a circular aperture of 3 millimetres diameter).
 The** film resolution **is a function of optics and film type. Measurements
are in line pairs per millimetre and are referenced to black and white film.
For the optics in use at ULCC the target resolutions are:**

Optics	resolution
	(lpm)
F300 (16mm)	90
F304 (35mm slide)	40
F309 (35mm aperture)	55

 [It is important to appreciate that these are the resolutions that may
be achieved at the film plane on high resolution film, such as DACO E. It
is possible to resolve 3000 lines across the image area on the CRT of the
D48C. The 4K raster of the D48C is designed so adjacent lines merge and
a continuous image is perceived.] **Registration accuracy** on the F333A
film transport is maintained within (plus or minus) 0.005 millimetres. This
is of significance in generation of cinematographic film and is adequate
for producing steady images.   In practice, measurements on the D48C at ULCC
(for geometrics, photometrics resolution and registration) have consistently
indicated performance well within these limits.   There are a number of aspects
concerning use of the D48C that are generally applicable to all devices.
These are dealt with in the following sections, after which each film mode
is considered separately (i.e. these being the different **devices** supported
on the Dicomed).

3.1.1 – Prime

An image is recorded from the excitation of phosphor on the CRT (through a controlled exposure of an electron beam). After the excitation – and associated luminescence – of the phosphor it must stabilise at a neutral level. If this stabilisation is not allowed to occur the next phosphor excitation may result in greater luminescence. That is, the exposure on the film may be greater than that desired. This effect <u>may</u> be observed as the "ghosting" of earlier images superimposed upon the current picture, and is known as **image retention** . It occurs most commonly on raster images (particularly in those areas of relatively low intensity) and is more usually noticed on grey scale images (on black and white film). The retained image is often a previously drawn vector picture or components from several. Usually lines drawn at high intensity will be subject to retention, although repeated tracing across the same area of the CRT [i.e. identically positioned vectors, which might be user data or the ULCC identification frame(s)] may also be retained. Consider the case of a high intensity vector (redrawn several times) followed by a low intensity uniform raster. Each point in the raster image will be subject to the same length of electron beam excitation on the CRT, and it would be expected that even exposure of the film across the raster would result. However, those areas of the CRT on which the vector was drawn may not have relaxed to the appropriate neutral level. At those points although excited for the same length of time phosphor luminescence will be at a different level than elsewhere: the result is a uniform raster with a vector superimposed. The effects on vector images are virtually unnoticeable. This is due partly to the very small overlap of vectors which might exhibit retention (intersections of line width slightly enhanced would not be perceived), but mainly to the fact that vectors are produced via higher exposures than rasters (dictated by relative visual perception) and so retention effects are proportionately lower. Similarly, with colour rasters exposures are at a higher level than for black and white due to film characteristics and spectral filtering (i.e. the filter cuts the transmission of phosphor luminescence considerably; conversely a longer exposure is required for a similar effect on film). Again, the effect may be observed on low intensity colour rasters. Similarly, high intensity raster images occasionly result in noticeable image retention on subsequent lower intensity rasters. The effect is <u>not</u> confined to adjacent images and may become apparent some considerable time later. Image retention is dependent on phosphor characteristics (in the case of the D48C this is P48). At ULCC the Dicomed is periodically **primed** ; that is, the CRT is saturated with a uniform high intensity raster. This reduces image retention to a very high degree. However, image retention may become a problem during graphic generation from stacked data (the method by which ULCC combine the plot data of many individual users into a single, continuous input to the Dicomed). To help obviate this DIMFILM users are provided with a means to prime the CRT during their plot run.

DPRIME(<device>) – this will cause the CRT to be primed for the designated device, provided it is currently active. If the display surface is non-empty it will be cleared before the prime is initiated. The frame will always be advanced after the prime. The single parameter is of type CHARACTER and should be a string comprising the device designation:

'D35' – will prime device D35 (35mm black and white aperture card format)

'D35P' – will prime device D35P (35mm black and white portrait mode format)

'D16' – will prime device D16 (16mm black and white movie)

```
   'D35C' – will prime
device D35C (35mm colour slide)

   'D16C' – will prime device D16C (16mm
colour movie)

   'ALL' – will prime all active Dicomed devices
```

   Specific references exist for each individual device if such mode is preferred
(these being detailed under the section for relating to each device).   It
should be noted that priming will expose a frame of film (the film transport
having no shutter). Thus, priming is inappropriate to cinematographic output
except preceding sequences.   Most users will not find image retention a problem.
The facility is included for users who encounter such effects. Priming
is a relatively slow process and excessive use could degrade throughput
of the Dicomed system. The length of this section is not intended to emphasise
the problem, but to explain and reassure users when it is first encountered.

3.1.2 – Frame repetition

 In certain applications it is useful to repeat (either singly or multiply)
the generation of each successive frame. In particular, for animations which
require twnety-four images for every second of screen time (twenty-five if
destined for PAL video) great economies of computation may be made by "double
framing", i.e. repeating each image. Thus for every second only twelve images
will have been computed. Much animation can be successfully created in this
way with little, or no, jitter or judder apparent to the viewer.   Most animations
will require action to be held for periods of time; it is then sensible to
instruct the system to replicate a particular frame the relevant number of
times (avoiding repeatedly re-computing the same image).   Two device specific
routines are provided in DIMFILM for these purposes.

DMULTI(<device>,MFRAM)   – this will set the multi-frame count for the
specified device (or, optionally, for all active Dicvomed devices).

   <device> –
is of type CHARACTER, and should be a string containing the device designation
(i.e. any of 'D35', 'D35P', 'D16', 'D35C',or 'D16C' for a specific device
or 'ALL' for all active Dicomed devices).

   MFRAM – is the required
multi-frame count (in range [0,5]) for all subnsequent frames.

   It must be noted that DMULTI is a <u>deferred</u> action: it does not take
effect until after the <u>next</u> frame advance is executed. Subsequent to that
advance, following frames will be replicated to generate the specified number
of copies of each.   In particular, double framing will be achieved by setting
MFRAM = 2. While for MFRAM = N a total of N frames will be generated for
each following image. MFRAM = 1 is single framing, and is equivalent to

MFRAM = 0. In each case the action will commence <u>following</u> the next reference
to FRAME (it is, therefore, sensible – as an aid to clarity – to position
the refernce to DMULTI to immediately precede a frame advanmce (i.e. reference
to FRAME): the image following that advance will then be the one for which
the DMULTI action first takes effect. It will remain active until reset.
  (Single and double framing may also be achieved through reference to the
DSET routine, described below.)

DHOLD(<device>,NFRAMS)   – this will repeat the next frame the specified
number of times, for the designated 16mm device, provided it is currently
active.

   <device> – is of type CHARACTER, and should be a string containing
the device designation (i.e. any either 'D16', or 'D16C'). This option is
only available for 16mm animations on the Dicomed film rcorder.

   NFRAMS –
is the numbers of times the following complete frame is to be generated.

   It must be noted that DHOLD is available only for 16mm film devices.
The action is <u>deferred</u> until the next frame advance is executed. The following
frame will then be generated the requisite number of times, after which framing
will revert to single or multi framing (as presently set). In particular,
NFRAMS is the number of frames in total that will be generated for the hold.
That is, if NFRAMS = 1 a single frame only will be produced (even if currently
multi-framing). In general, for NFRAMS = N the original frame will be generated
followed by (N–1) repeats, yielding a total of N frames. As DHOLD is a deferred
action it is sensible to position the reference to it to immediately precede
a frame advance (i.e. reference to FRAME): the image following that advance
will then be the one to which the DHOLD action refers.

3.1.3 – Dicomed device control

 There are other ways in which Dicomed devices can be controlled or reset.
A routine, DSET, has been provided for some such requirements. In particular,
this can be used to globally change the orientation of a device or to control
single or double framing (in this it is a simple alternative to the facility
provided in the DMULTI routine).

DSET(<device>,KEY)   – this will select particular operational modes
for the specified device (or, optionally, for all active Dicvomed devices).

   <device> – is of type CHARACTER, and should be a string containing
the device designation (i.e. any of 'D35', 'D35P', 'D16', 'D35C',or 'D16C'
for a specific device or 'ALL' for all active Dicomed devices).

   KEY –
is a character string specifying the required action. It may be one of the
following:

```
   SINGLE – turn on single framing, deferred until after the
next  frame advance.

   DOUBLE – turn on double framing, deferred until
after the next  frame advance.

   COMIC – turn on comic mode orientation.


   CINE – turn on cine mode orientation.

   LANDSCAPE – turn on landscape
mode orientation.

   PORTRAIT – turn on portrait mode orientation.
```

   For a fuller discussion of single and multi framing refer to the description of DMULTI.   Orientation of an image may be considered in two ways. Typically one refers to a picture as being in either landscape or portrait orientation. A picture is said to be landscape if its longer dimnension is horizontal, while it is portrait if its longer dimension runs vertically.   For film the equivalent term is comic or cine; for comic mode the picture base is longitudinal with respect to the (edge of the) film, while in cine mode the base of the image is transverse (i.e. across the film). For each film format there is a different default and imaging is done relative to this. Through reference to DSET it is possible to change that orientation. The effect is of rotating the image through ninety degrees before recording it on the film. The default mode for 16mm film is CINE and for 35mm is COMIC, but due to the aspect ratios of each both correspond to LANDSCAPE images. (For 16mm the longer dimension is transverse, while for 35mm it is longitudinal.) Thus the default image layout is LANDSCAPE for all the Dicomed film modes. The routine DSET enables the production of, for example, portrait mode slides or portrait mode images on 16mm (the latter would be unsuited for projection, but would be quite suitable for archive purposes).

 For each device the equivalent combinations may be shown as:

DEVICE:	D35	D16	D35P	D35C	D16C
ASPECT RATIO					
LANDSCAPE	COMIC	CINE	COMIC	COMIC	CINE <--DEFAULT
PORTRAIT	CINE	COMIC	CINE	CINE	COMIC

   It is strongly advised that changes of orientation are made <u>only</u> at the beginning of a frame; they take immediate effect and various device

transformations will be reset to reflect the selected mode.

3.1.4 – Raster elements

As intimated above, the raster element as used on the D48C may be composed
of a grouping of points with both point and element spacing being variable.
DIMFILM raster facilities do not require the user to address this topic.
All mappings from the cell mapped area to device raster elements are performed
within the device drivers. However, it is recognised that for certain serious
applications the user will require more detailed understanding of the raster
construction. Additionally the user is provided with the means to respecify
the raster characteristics (and DIMFILM will take account of the actual raster
structure when transmitting cell data to the Dicomed).   The **Raster element**
 is the unit by which the D48C plots a raster picture (i.e. it is the basic
picture element, or **pixel**). When a raster is transmitted to the D48C
a series of required pixel intensity values are transferred. Successive
raster elements are plotted with those successive intensities, the same
intensity being applied to each point in the element. When one element has
been plotted the beam is spaced to the next; when a scan line is complete
the beam is repositioned ready for the next scan. The **element spacing**
 (both horizontally and vertically) may be set. Each raster element is
a square or rectangular array of one to sixteen **points** in each of the
horizontal and vertical directions. The **point spacing** (both horizontally
and vertically) within the element may also be set. The points correspond
to the 32K grid addressability which may be used to expose a single spot
on the CRT – these points should not be confused with pixels, of which they
are a component. Only in the limiting case of elements consisting of single
points will points be equivalent to pixels. The raster elements are the pixels
of which the image is composed, and it will be seen that the D48C offers
the user the ability to specify both the size and spacing of pixels.   The
raster structure is determined by a number of parameters. The **raster element**
 consists of an array of points, arranged with a specified **number of
points horizontally** and a **number of points vertically** , these being
in the range [1,16]. Spacings are expressed in raster units. That is, in
terms of the 4K x 4K maximum raster grid of the D48C. However, the spacings
may be stepped by eighths of the raster unit which is equivalent to units
on the 32K x 32K addressable grid of the CRT. (It is these latter units
which are used in DIMFILM.) The spacing of points within this array may
be set independently in each direction via the **horizontal point spacing**
 and the **vertical point spacing** , each of these being in the range
1 to 7-7/8 raster units (i.e. in the range [8,63] CRT addressable units).
Similarly, elements may be spaced in each direction by specifying the **horizontal
element spacing** and the **vertical element spacing** , again in the range
1 to 7-7/8 raster units (i.e. [8,63] CRT addressable units). The element
spacing is measured between the lower left corners of adjacent elements (in
the appropriate direction), and should be such that adjacent elements do
not overlap (i.e. in either direction the element spacing should exceed the
product of the number of points and the point spacing).

The element structure may be shown:

```
_____
 |    |    |    |    |    |    |
 |    |    |    |    |    |    |
 |    |    |    |    |    |    |
_____
 |    |    |    |    |    |    |
 |    |    |    |    |    |    |
 |    |    |    |    |    |    |
O----O---------X----X------------
 |    |    |    |    |    |    |
 |    |    |    |    |    |    |
 |    |    |    |    |    |    |
_____
 |    |    |    |    |    |    |
 O    O    |    X    X    |    |
 |    |    |    |    |    |    |
_____
 |    |    |    |    |    |    |
 |    |    |    |    |    |    |
 |    |    |    |    |    |    |
O----O---------X----X------------
 |    |    |    |    |    |    |
 |    |    |    |    |    |    |
 |    |    |    |    |    |    |
_____
 |    |    |    |    |    |    |
 |    |    |    |    |    |    |
 |    |    |    |    |    |    |
_____
 |    |    |    |    |    |    |
 |    |    |    |    |    |    |
 |    |    |    |    |    |    |
_____
 |    |    |    |    |    |    |
 |    |    |    |    |    |    |
 |    |    |    |    |    |    |
X----X---------X----X------------
 |    |    |    |    |    |    |
 |    |    |    |    |    |    |
```

In the foregoing illustration, the lines represent a part of a 4K x 4K
grid on the CRT (i.e. raster units). The 'O's are the points comprising
one raster element, while the 'X's are points in adjacent elements. The
grid divisions (i.e. 1/4096 of image dimension) are termed **raster units**
. In the example a **raster element** consists of six **points** arranged
as 2 horizontally by 3 vertically. The **horizontal point spacing** is 1
raster unit, and the **vertical point spacing** is 1.5 raster units. The
spacing between elements is determined by the **horizontal element spacing**
 of 3 raster units, and a **vertical element spacing** of 7 raster units.
 DIMFILM provides the means whereby the user may change the specification
of the raster structure.

DRAST(<device>,HES,HPS,HPE,VES,VPS,VPE)   – this will set the raster
structure for the designated device, provided it is currently active.

```
  <device> –
is of type CHARACTER, and should be a string  containing the device designation
(i.e. any of 'D35', 'D35P', 'D16',  'D35C', 'D16C' for a specific device or
'ALL' for all  active Dicomed devices),
```

HES – is the **horizontal element spacing** in CRT addressable units (= 8 x **raster unit**), as an INTEGER in the range [8,63], but subject to exceeding HPExHPS.

HPS – is the **horizontal point spacing** in CRT addressable units (= 8 x **raster unit**), as an INTEGER in the range [8,63].

HPE – is the number of **horizontal points per element** , as an INTEGER in the range [1,16].

VES – is the **vertical element spacing** in CRT addressable units (= 8 x **raster unit**), as an INTEGER in the range [8,63], but subject to exceeding VPExVPS.

VPS – is the **vertical point spacing** in CRT addressable units (= 8 x **raster unit**), as an INTEGER in the range [8,63].

VPE – is the number of **vertical points per element** , as an INTEGER in the range [1,16].

```
  Specific references exist for each individual device if such  mode is preferred.
 On initialisation, each Dicomed device is initialised with a 2048 x  2048
raster with elements consisting of 2 x 2 unit spaced points.  Element spacing
is 2 raster units.  The equivalent parameters  to DRAST would be HES = 16,
HPS = 8, HPE = 2, VES = 16,  VPS = 8, and VPE = 2.  For most purposes this
will yield  eminently acceptable results.
```

3.1.5 – Line Width/Spot Size

```
 The Dicomed devices are able to produce thickened lines and  (hardware)
spots of various sizes.  These are controlled through the  DIMFILM routines
for line width and spot size scale factors.  The supported range is from 1
to 16 in each case – the default  being for a single width line and minimal
point.  Scale factors below  1.0 will equate to the (default) of 1.0, while
values above 1.0 will  select the nearest integral value up to a maximum of
16.   NOTE : At present multi-width lines on the Dicomed are produced
on a vector by vector basis and no account is made of any preceding  or following
vector.  Unacceptable results will be produced for  thick lines forming parts
of curves (in particular characters).  This will be investigated in due course;
until it has been resolved  users are advised to restrict their use of thickened
lines to  lower values of the scale factor (a maximum of 3.0 is suggested).
```

3.1.6 – Plot Dataset Length

 DIMFILM attempts to utilise the Amdahl queues in an economic and efficient
manner. To this end a limit is applied to the permitted size of a user's
plot dataset. At present, the Dicomed drivers will terminate when a dataset
length of 100 MBytes is exceeded. We shall review this limit in the light
of experience, and would be pleased to hear from users who encounter difficulties
as a result of this limit. It is possible that the limit may be reduced
if operational difficulties result – every effort will be made to give notice
of such change, but this cannot be guaranteed.

3.1.7 – Device Sizes

 The physical image size on film for each of the Dicomed modes is maintained,
to a close tolerance, according to various international standards. The
default coordinate system for each film device is intended to produce images
of the defined size (within the physical device limitations). While these
have been carefully determined, the right is reserved to amend them should
the need arise. There will be no effect upon users specifying their own
coordinate system through reference to BOUNDS.

3.1.8 – D35 – monochrome aperture card format (35mm)

```
 Medium: 35mm unperforated film (reversal)
 Imaging: monochrome, vector and raster
 Physical image size: 36.41x27.94 millimetres
 Default horizontal coordinates:  0.0,32236.0
 Default vertical coordinates:    0.0,24708.0
 Default image/film mode:  landscape/comic
 Default (initial) intensity: 0.8
```

3.1.9 – D35P – monochrome 35mm portrait mode

```
 Medium: 35mm unperforated film (reversal)
 Imaging: monochrome, vector and raster
 Physical image size: 27.94x36.41 millimetres
 Default horizontal coordinates:  0.0,24708.0
 Default vertical coordinates:    0.0,32236.0
 Default image/film mode:  portrait/cine
 Default (initial) intensity: 0.8
```

3.1.10 – D16 – monochrome 16mm (movie)

```
Medium: 16mm double perforated monochrome film (reversal)
Imaging: monochrome, vector and raster
Physical image size: 10.26x7.49 millimetres
Default horizontal coordinates:  0.0,32767.0
Default vertical coordinates:    0.0,23920.0
Default image/film mode:  landscape/cine
Default (initial) intensity: 0.8
```

3.1.11 – D35C – colour slide 35mm

```
Medium: 35mm double perforated, reversal colour film
Imaging: colour, vector and raster
Physical image size: 36.3x24.5 millimetres
Default horizontal coordinates:  0.0,32148.0
Default vertical coordinates:    0.0,21698.0
Default image/film mode:  landscape/comic
Default (initial) intensity: 0.8
Default (initial) colour: (0.8,0.8,0.8)
```

3.1.12 – D16C – colour 16mm (movie)

```
Medium: 16mm double perforated, reversal colour film
Imaging: colour, vector and raster
Physical image size: 10.26x7.49 millimetres
Default horizontal coordinates:  0.0,32767.0
Default vertical coordinates:    0.0,23920.0
Default image/film mode:  landscape/cine
Default (initial) intensity: 0.8
Default (initial) colour: (0.8,0.8,0.8)
```

Appendix 4 – Conversion to the new DIMFILM

```
 This chapter tabulates user accessible routines in the original  DIMFILM
and indicates, where relevant, the new name and any change  in arguments.
```

The major parameter change is associated with character strings. Implementation
under Fortran 77 has required the replacement of (hollerith string,character
count) by a single parameter of type character.

```
   Thus, for example, the old coding
        CALL SYMTEXT(6HABCDEF,6)
   becomes
        CALL SYMTXT('ABCDEF')
```

 The other general change affecting parameters is the change from an integer
intensity to a real value (normalised in the range 0. to 1.).   Other than
subroutine name changes, the major difference the user will encounter is
the replacement of the single character alternators (user assignable) by
a predetermined set of escape sequences. This has the advantage of only removing
two characters from normal usage in strings and simultaneously offering even
greater flexibility to the user in formatting and controlling his textual
output.   New features are not generally described in this Appendix unless
they directly replace an original feature. Thus, for instance, colour specification
and raster plotting is not described here – nor are the device nomination
routines appropriate to them.

4.1 – Routine name conversion

  Notes:

---

 n/a – indicates routine or equivalent is no longer available

 blank in second column if no change

 +C+ Hollerith string, count parameter pair(s) replaced by
character type parameter(s)

 +R+ integer intensity value replaced by real value

    Original   New routine name

    ACCENT   n/a – but see new text description
    AUTODIV   AUTD
    AUTOR
    AUTORTH   AUTORT

```
AUTOTH        AUTOT
AUTOX
AUTOXY
AUTOY
AUTRDIV       AUTORD
AXISCUT       AXCUT
AXTIKON       AXTIK
BLANK
BLNKFRM       OBLANK
BNDRY         INQB
BOUNDS
BOX           RECT      – see also new routine BOX
BRKSYM
CALPNT        n/a
CAM16MM       D16
CAM35MM       D35
CASEALT       n/a – but see new text description
CFMT                    +C+
CHAR          MARKER     – but see new description
CHARA         n/a
CHARL         n/a
CHECK
CHHALF
CINE          n/a
CINTERP       CINTER
CIRCLE
CIRCSEG       CIRARC
CLABEL
CLABHT
COMIC         n/a
CONTOUR       CONTR
CONTR1                  – see description re last parameter
CONTR2
COORDS        INQC
CQUAD
CSCALE
CVHALF
CWHOLE
DASH
DASHDOT       DSHDOT
DASHOFF       DSHOFF
DEFALTS       n/a
DEFINT        n/a – but see general description
              of intensities
DEFMATH       n/a – but see new text description
DEGREES       DEGREE
DEGSEP        RADSEP
DIVUNIT       DIVUNI
DOT
DRAWXAX       DRAWXA
DRAWYAX       DRAWYA
DTHETA
EDGES
ELLIPSE       ELLIPS
ELLSEG1       ELARC1
ELLSEG2       ELARC2
ENDBLNK       ENBLNK
ENDCONT       ENSYMC
ENDFILM       DIMEND
ENDWIND       ENPANE
FNPLOT
```

```
FRACTNS       n/a – but see new text description
FULL16        n/a
GRAFCON       GRCON
GRAFDEF       GRDEF     +C+
GRAPHER       GRAPHE
GRAPHFN       GRAPHF
GRAPHIC       GRAPH
GREEK         n/a – but see new text description
GRFRAME       GRFRAM
GROFF
GRON
GRSLIDE       GRCONA
GRTOAC
HATCH
HISTGRM       HISTGR
HPER
HREP
HVYALT        n/a – but see new text description
HVYINT                  +R+
HVYTYPE       HVYTYP
IBACK         n/a – but see general description
              of intensities
IBCOUNT       IBCNT
IDEF          n/a – but see general description
              of intensities
INTERP
INTREND       INTRND
INTRUPT       INTRPT
INTSTY        INTSY   +R+ – single real parameter (0.,1.),
              see description
ITALIC
ITALT         n/a – but see new text description
IVAL          INQR/INQI
IXAX
IYAX
IZVAR         INTVAR    +R+
JOIN
LIMIT         n/a
LINALT        n/a
LINESYM       n/a
LINK
LINR
LINX
LINXY
LINY
LNUM          n/a
LOGLIM
LOGR
LOGTIK                  +R+
LOGX
LOGXY
LOGY
LOWCASE       LOWER
LRYOPPT       LRYOPT
LRYTICK       LRYTIK
LRYVAL
LTITLE                  +C+
LUXOPPT       LUXOPT
LUXTICK       LUXTIK
LUXVAL
LXAXLAB       LXALAB    +C+
```

```
LXAXTIK        LXATIK
LXAXVAL        LXAVAL
LXLABEL        LXLAB      +C+
LXOPPT         LXOPT
LXTICK         LXTIK
LXVAL
LYAXLAB        LYALAB     +C+
LYAXTIK        LYATIK
LYAXVAL        LYAVAL
LYLABEL        LYLAB      +C+
LYOPPT         LYOPT
LYTICK         LYTIK
LYVAL
MARGIN
MATHEXP        n/a – but see new text description
MONTHSX        MONSX
MONTHSY        MONSY
MONTHX
MONTHY
NEWFRAM        FRAME
NEWORIG        NEWORG
NEXTONE        n/a – but see new text description
NOACCNT        n/a – but see new text description
NOAXTIK        NOAXT
NOCHECK        NOCHCK
NOCINT
NOCLAB
NOFRACT        n/a – but see new text description
NOITAL
NOIZVAR        NOIVAR
NOMATH         n/a – but see new text description
NONEXT         n/a – but see new text description
NOROT
NOSPEC         n/a – but see new text description
NOSPOS
NOSUB          n/a – but see new text description
NOSUPER        n/a – but see new text description
NOTXTCN        ENTXTC
NOUNLIN        NOUNLN
NRMTYPE        NRMTYP
NUM            RNUM/INUM  – real/integer value, format as
               character type
OFFTOXY        OFF2
OLDORIG        OLDORG
ONTOXY         ON2
OPMESS         n/a
OUTLINE        OUTLIN
OUTSLDE        n/a
OWNDASH        OWNDSH
PARALL1        PARAL1
PARALL2        PARAL2
PIECHRT        PIECHT – caption now singly dimensioned
               character array
POINT
POLRDEF        POLDEF     +C+
POLYGON        POLYGN
POLY3
POLY3ER        POLY3E
POLY5
POLY5ER        POLY5E
POSN
```

```
PTPLOT
PTPLTER         PTPLTE
RADIANS         RADIAN
RELAXES         RELAX
RELFRAM         RELFR
REPEAT          n/a – to be implemented in more versatile
                manner
ROMAN           n/a – but see new text description
ROMGRK          n/a – but see new text description
ROTATE
RRANGE
RTHDIV          RDIV
RTHETA          POLAR
RTHFN           POLFN
RTHFRAM         POLFR
RTHGRID         POLGRD
RTHLEVS         TLEVS
RTHOFF          POLOFF
RTHON           POLON
RTHOUT          POLOUT
RTHPLY3         POLPY3
RTHPLY5         POLPY5
RTHPTS          POLPTS
RTHTOAC         POLTOA
RTIKS           RTIK
RVALS           RVAL
RXAX
RXTIK
RXVAL
RYAX
RYAXLAB         RYALAB      +C+
RYAXTIK         RYATIK
RYAXVAL         RYAVAL
RYLABEL         RYLAB       +C+
RYOPPT          RYOPT
RYTICK          RYTIK
RYVAL
SAMER
SAMERTH         SAMERT
SAMETH          SAMET
SAMEX
SAMEXY
SAMEY
SERIAL          n/a
SETDASH         SETDSH
SETITAL         SETITL
SETUP           n/a
SHDEGRM         SHDEGR
SLDSYM
SLIDE           n/a
SNUM            SRNUM/SINUM – real/integer value, format
                as character type
SPECHAR         n/a – but see new text description
STEP            n/a
STEPGRM         STEPGR
STRAIL                      +C+
STRING                      +C+
SUB             n/a – but see new text description
SUBALT          n/a – but see new text description
SUPALT          n/a – but see new text description
SUPER           n/a – but see new text description
```

```
SYMANGL        SYMANG
SYMCONT        SYMCON
SYMHT
SYMINT          - to be made available (renamed)
SYMPOS
SYMTEXT        SYMTXT      +C+
TEXTCON        TXTCON
THRANGE        TRANGE
THTIKS         TTIK
THVALS         TVAL
TICKCON        TIKCON
TITLE          n/a
TMPORIG        TMPORG
UNLIN
UPCASE         UPPER
UTITLE                     +C+
UXAXLAB        UXALAB      +C+
UXAXTIK        UXATIK
UXAXVAL        UXAVAL
UXLABEL        UXLAB       +C+
UXOPPT         UXOPT
UXTICK         UXTIK
UXVAL
WINDFRM        OPANE
WINDOW         PANE
XAXIS
XGRID
XUNIT
XYGRID
YAXIS
YGRID
```